

AWS Best Practices

v1.0



Minh Tran/ NashTech

Jun/2023

Nash
Tech.

Agenda

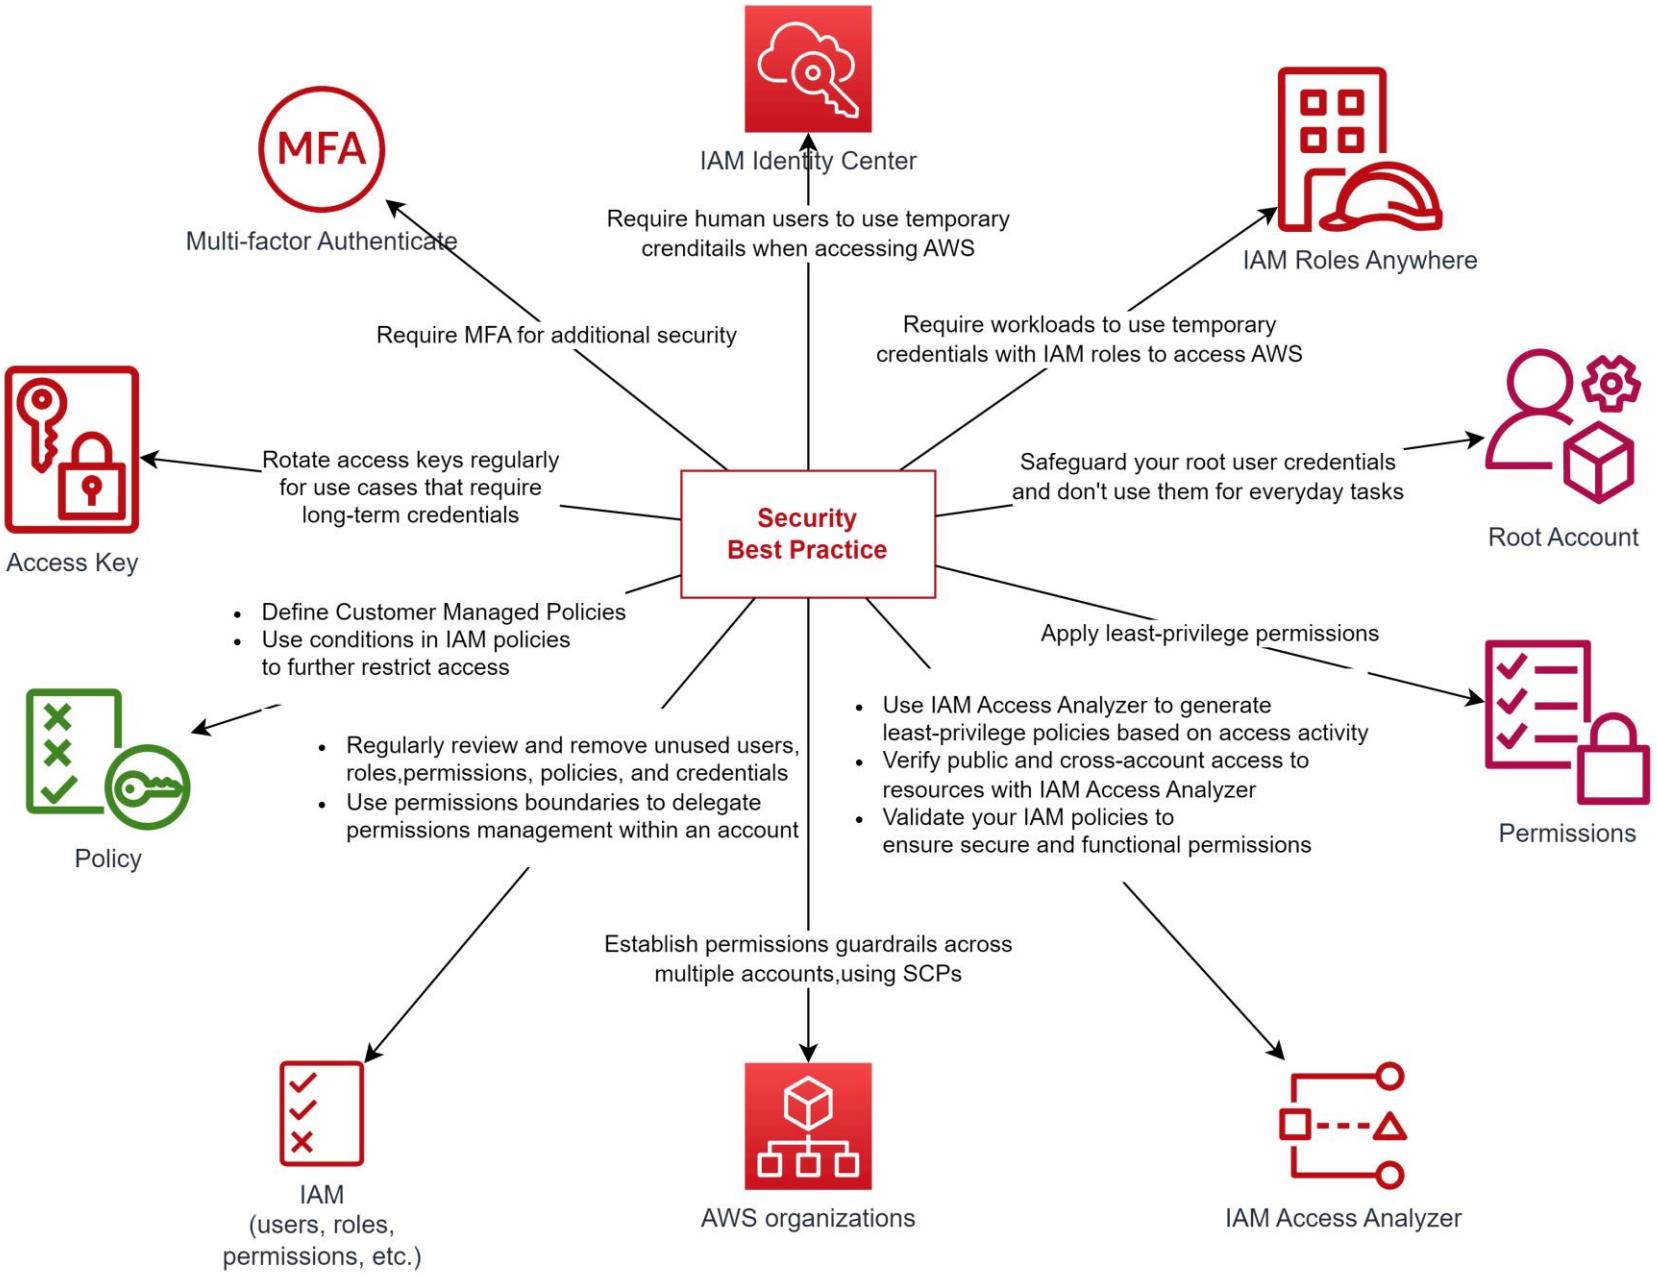
1. AWS Identity and Access Management (IAM)
2. AWS Organizations
3. Amazon EC2 (Elastic Cloud Computing)
4. Amazon RDS (Relational Database Services)
5. Amazon Aurora
6. Amazon S3 (Simple Storage Service)
7. Amazon Lambda
8. Amazon Cognito
9. Amazon SNS (Simple Notification Service)
10. Amazon VPC (Virtual Private Cloud)
11. Amazon Kinesis
12. Amazon Auto-scaling
13. Dynamo DB
14. Amazon SQS (Simple Queue Service)
15. AWS CloudFormation

AWS Identity and Access Management (IAM)

Nash
Tech.

Security Best Practice

- [Require human users to use federation with an identity provider to access AWS using temporary credentials](#)
- [Require workloads to use temporary credentials with IAM roles to access AWS](#)
- [Require multi-factor authentication \(MFA\)](#)
- [Rotate access keys regularly for use cases that require long-term credentials](#)
- [Safeguard your root user credentials and don't use them for everyday tasks](#)
- [Apply least-privilege permissions](#)
- [Get started with AWS managed policies and move toward least-privilege permissions](#)
- [Use IAM Access Analyzer to generate least-privilege policies based on access activity](#)
- [Regularly review and remove unused users, roles, permissions, policies, and credentials](#)
- [Use conditions in IAM policies to further restrict access](#)
- [Verify public and cross-account access to resources with IAM Access Analyzer](#)
- [Use IAM Access Analyzer to validate your IAM policies to ensure secure and functional permissions](#)
- [Establish permissions guardrails across multiple accounts](#)
- [Use permissions boundaries to delegate permissions management within an account](#)

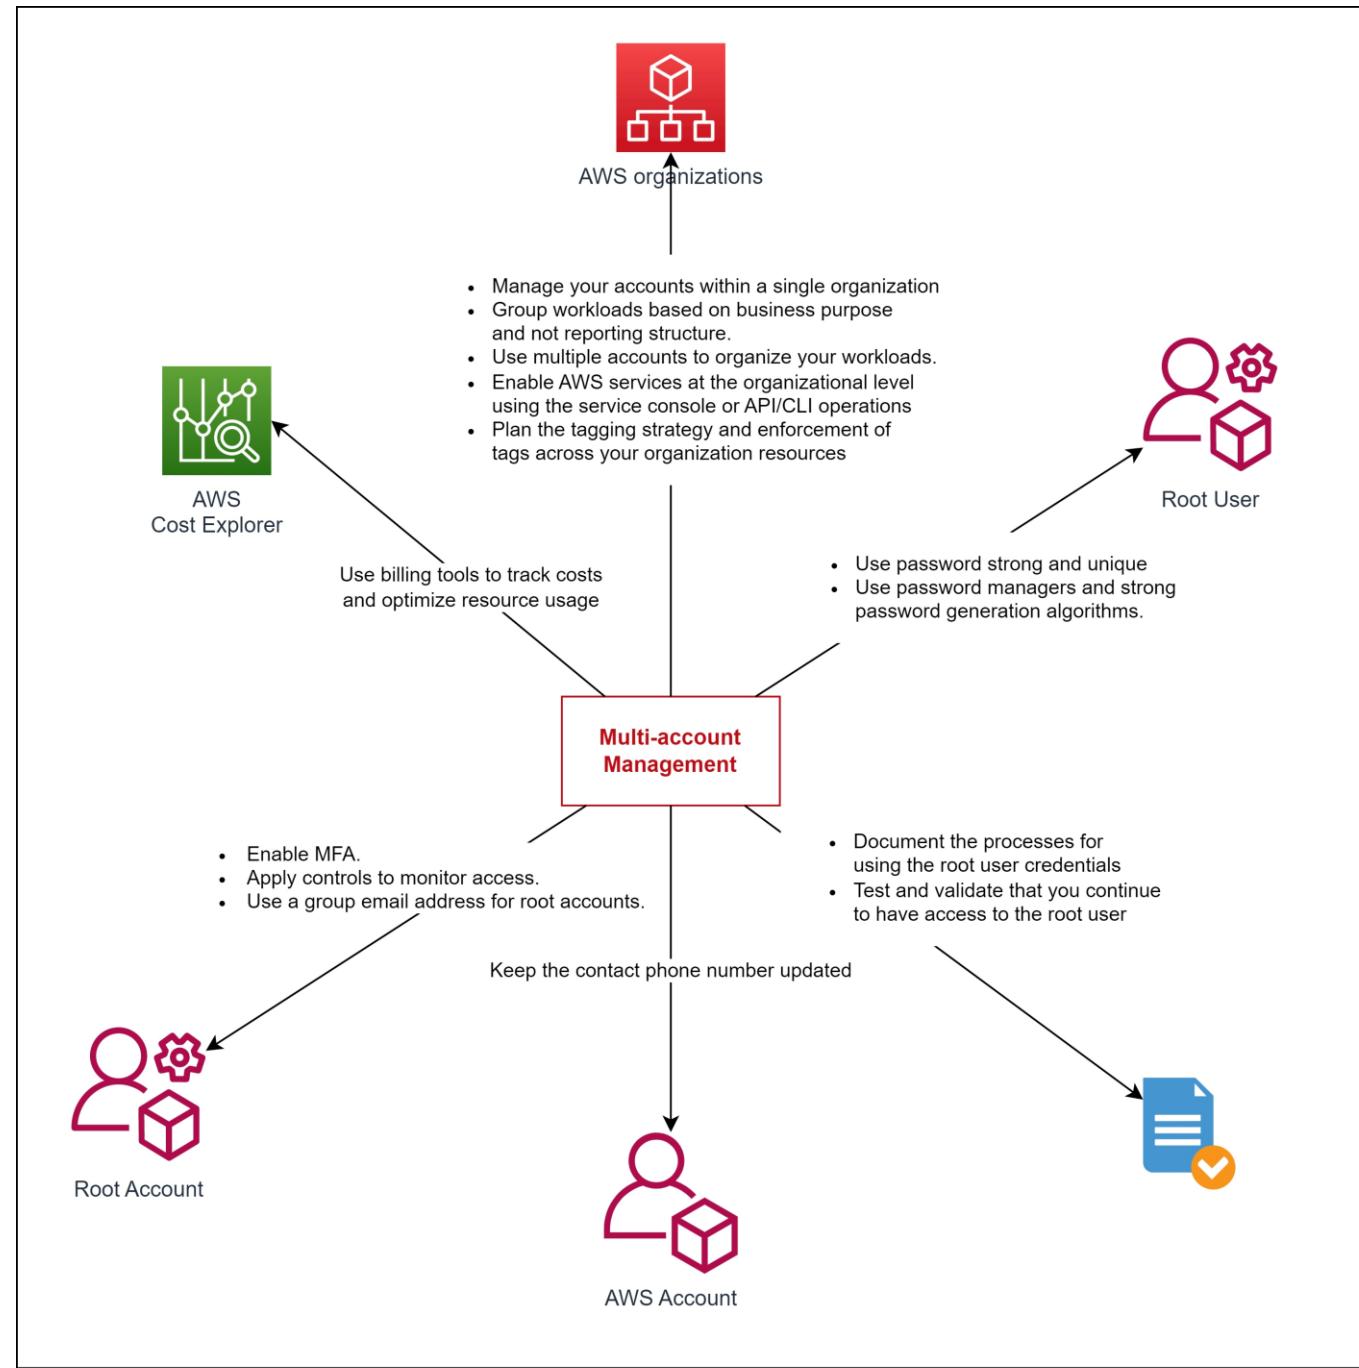


AWS Organizations

Nash
Tech.

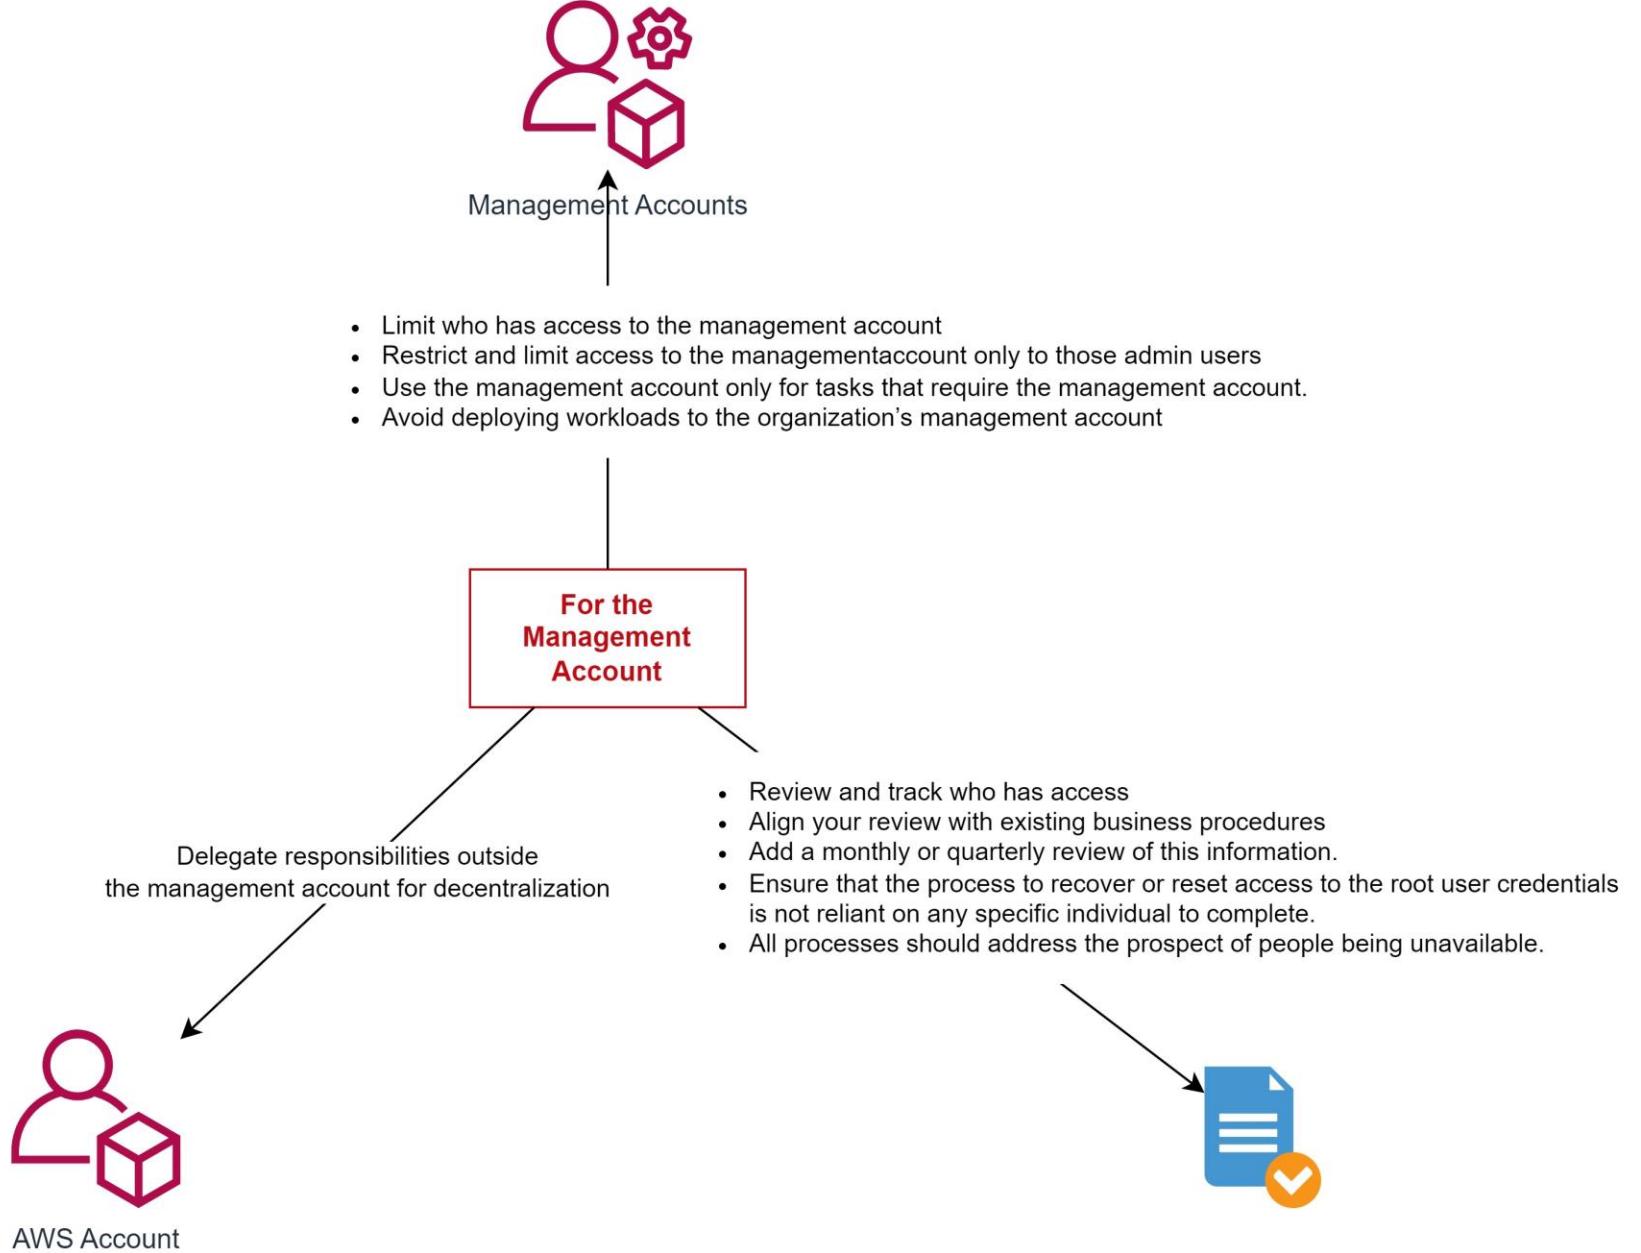
Best practices for multi-account management

- [Manage your accounts within a single organization](#)
- [Use a strong password for the root user](#)
- [Document the processes for using the root user credentials](#)
- [Enable MFA for your root user credentials](#)
- [Apply controls to monitor access to the root user credentials](#)
- [Keep the contact phone number updated](#)
- [Use a group email address for root accounts](#)
- [Group workloads based on business purpose and not reporting structure](#)
- [Use multiple accounts to organize your workloads](#)
- [Enable AWS services at the organizational level using the service console or API/CLI operations](#)
- [Use billing tools to track costs and optimize resource usage](#)
- [Plan the tagging strategy and enforcement of tags across your organization's resources](#)



Best practices for the management account

- Limit who has access to the management account
- Review and track who has access
- Use the management account only for tasks that require the management account
- Avoid deploying workloads to the organization's management account
- Delegate responsibilities outside the management account for decentralization



Best practices for member accounts

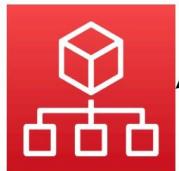
- [Define account name and attributes](#)
- [Efficiently scale your environment and account usage](#)
- [Use an SCP to restrict what the root user in your member accounts can do](#)



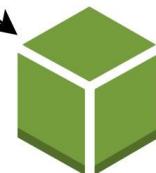
- Define account name and attributes.
- Before creating new accounts, make sure accounts for similar needs do not already exist.

For the Member Account

- Before closing accounts, note that they are subject to close account quota limits.
- Check Quotas for AWS Organizations.
- Check CloseAccountAPI limits.
- Clean up unneeded resources or workloads from the accounts
- Save the Accounts for a future use, if you're planning to reuse the accounts.



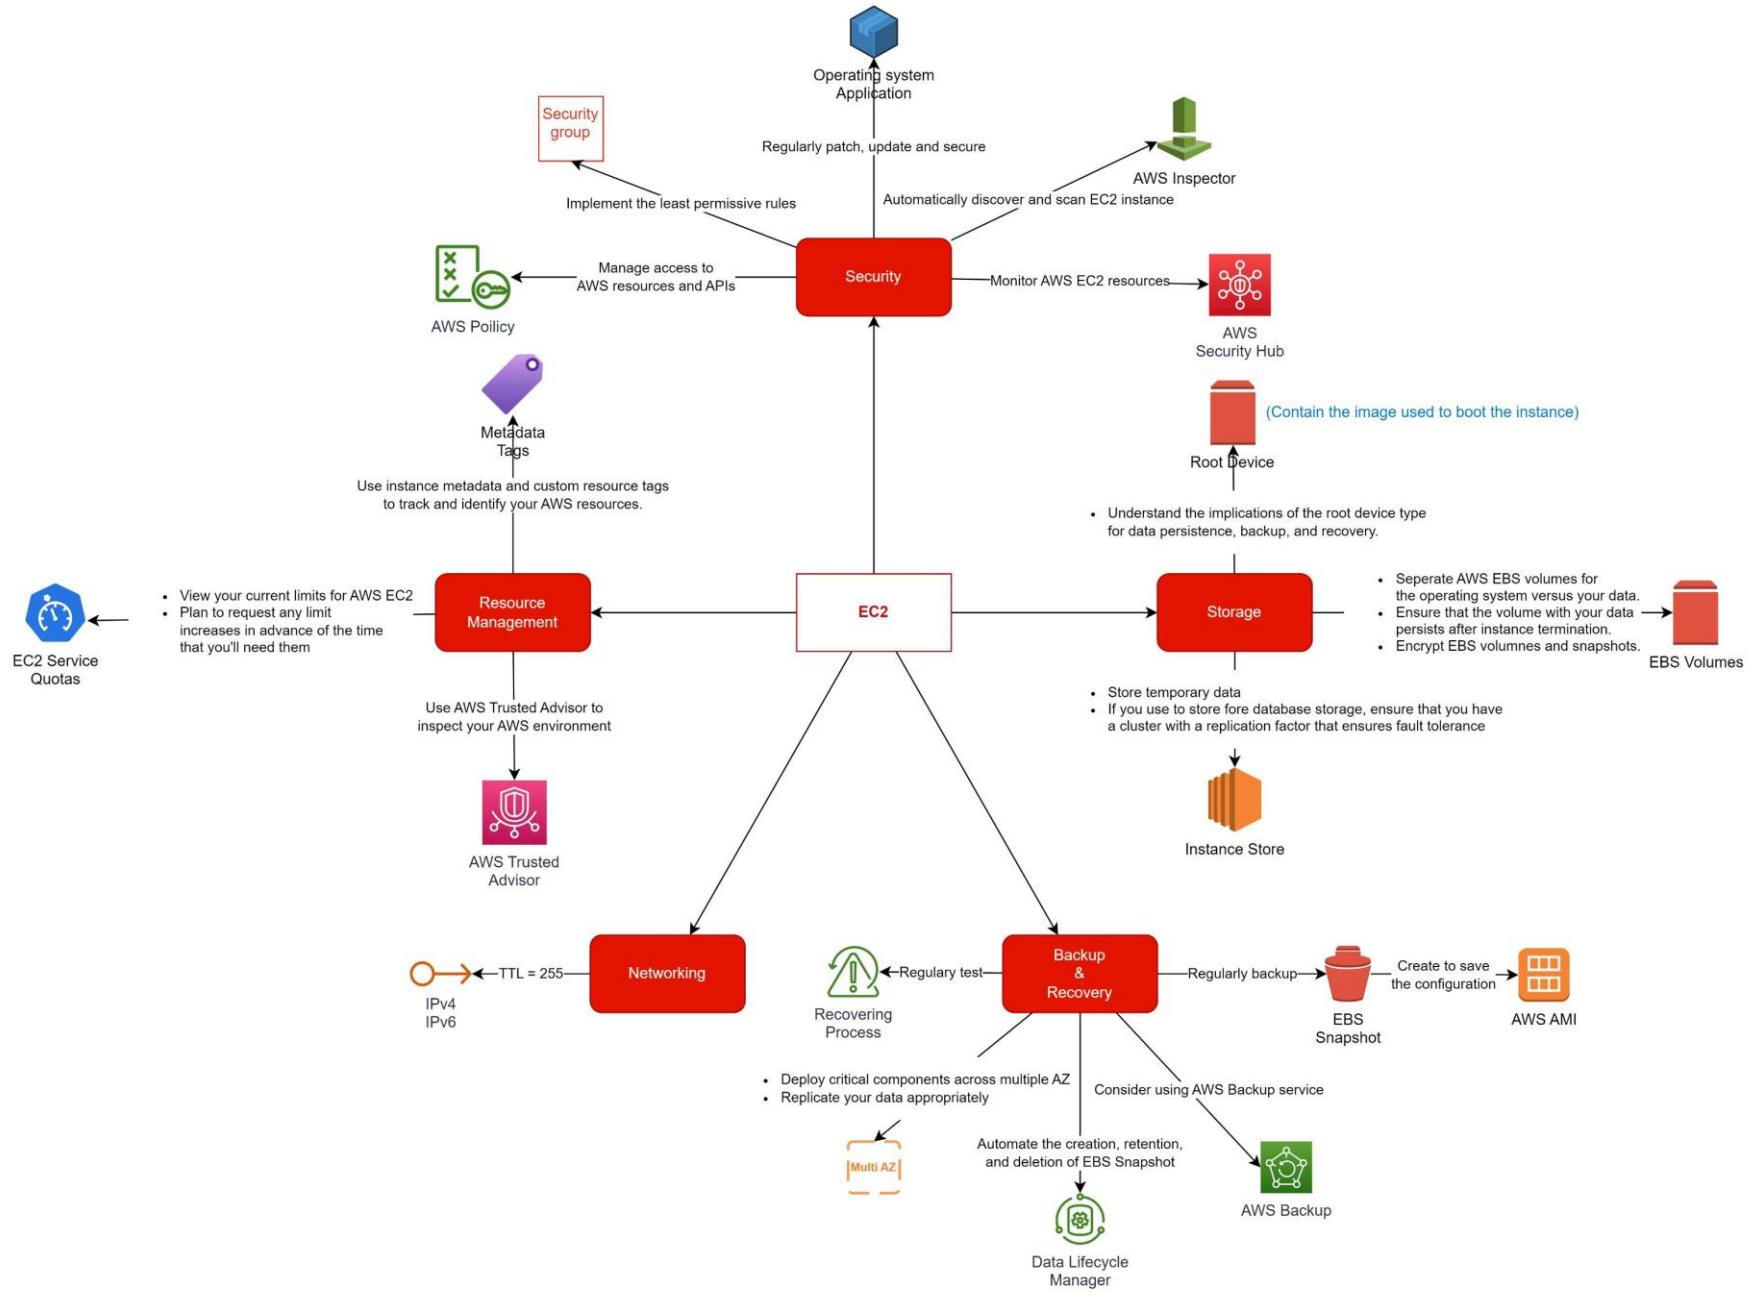
AWS organizations



AWS Resources

Amazon EC2 (Elastic Cloud Computing)

Nash
Tech.

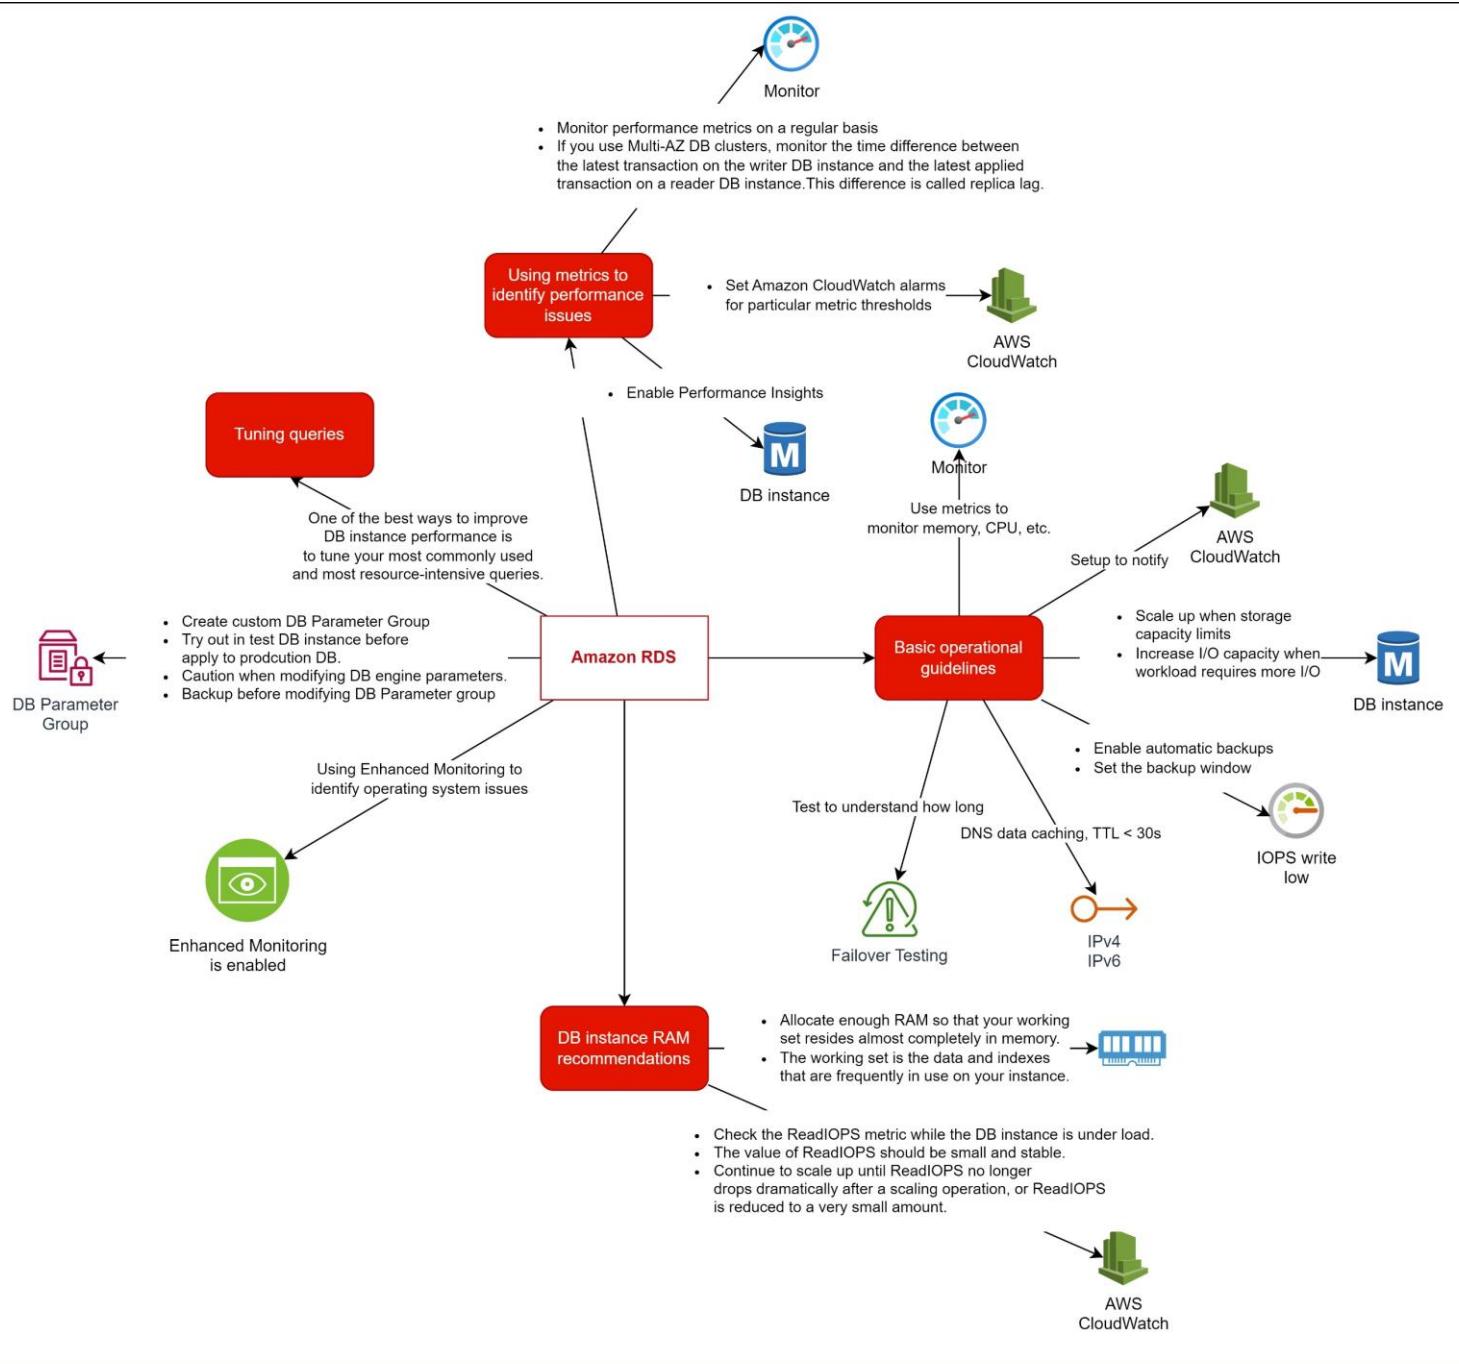


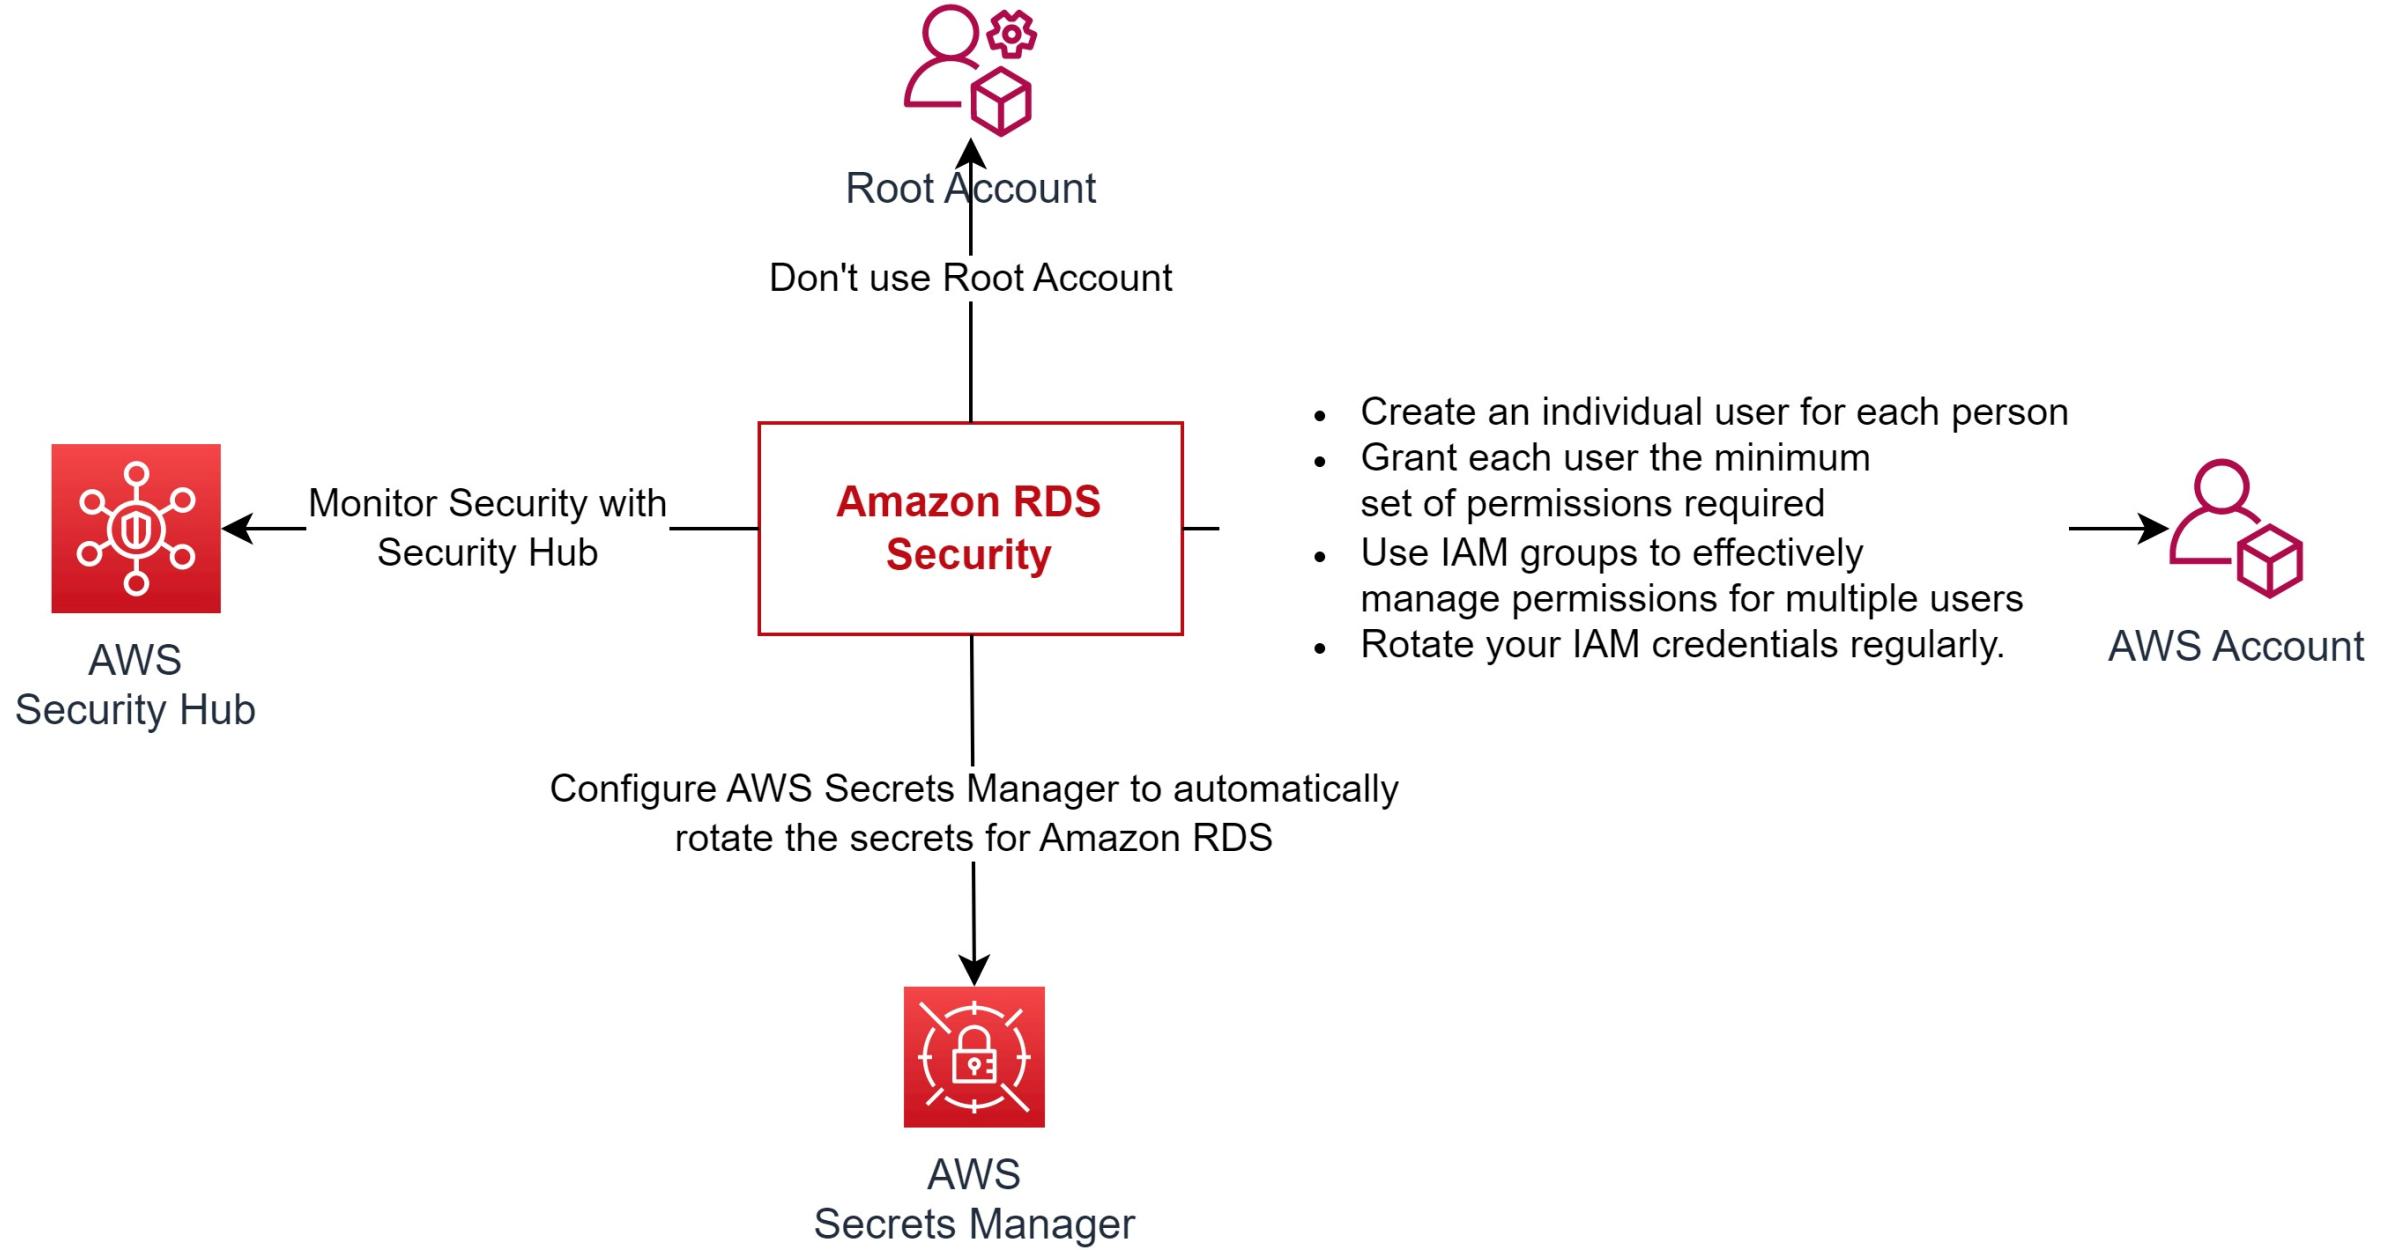
Amazon RDS (Relational Database Services)

Nash
Tech.

Best practices for Amazon RDS

- [Amazon RDS basic operational guidelines](#)
- [DB instance RAM recommendations](#)
- [Using Enhanced Monitoring to identify operating system issues](#)
- [Using metrics to identify performance issues](#)
- [Tuning queries](#)
- [Best practices for working with MySQL](#)
- [Best practices for working with MariaDB](#)
- [Best practices for working with Oracle](#)
- [Best practices for working with PostgreSQL](#)
- [Best practices for working with SQL Server](#)
- [Working with DB parameter groups](#)
- [Best practices for automating DB instance creation](#)
- [Amazon RDS new features and best practices presentation video](#)



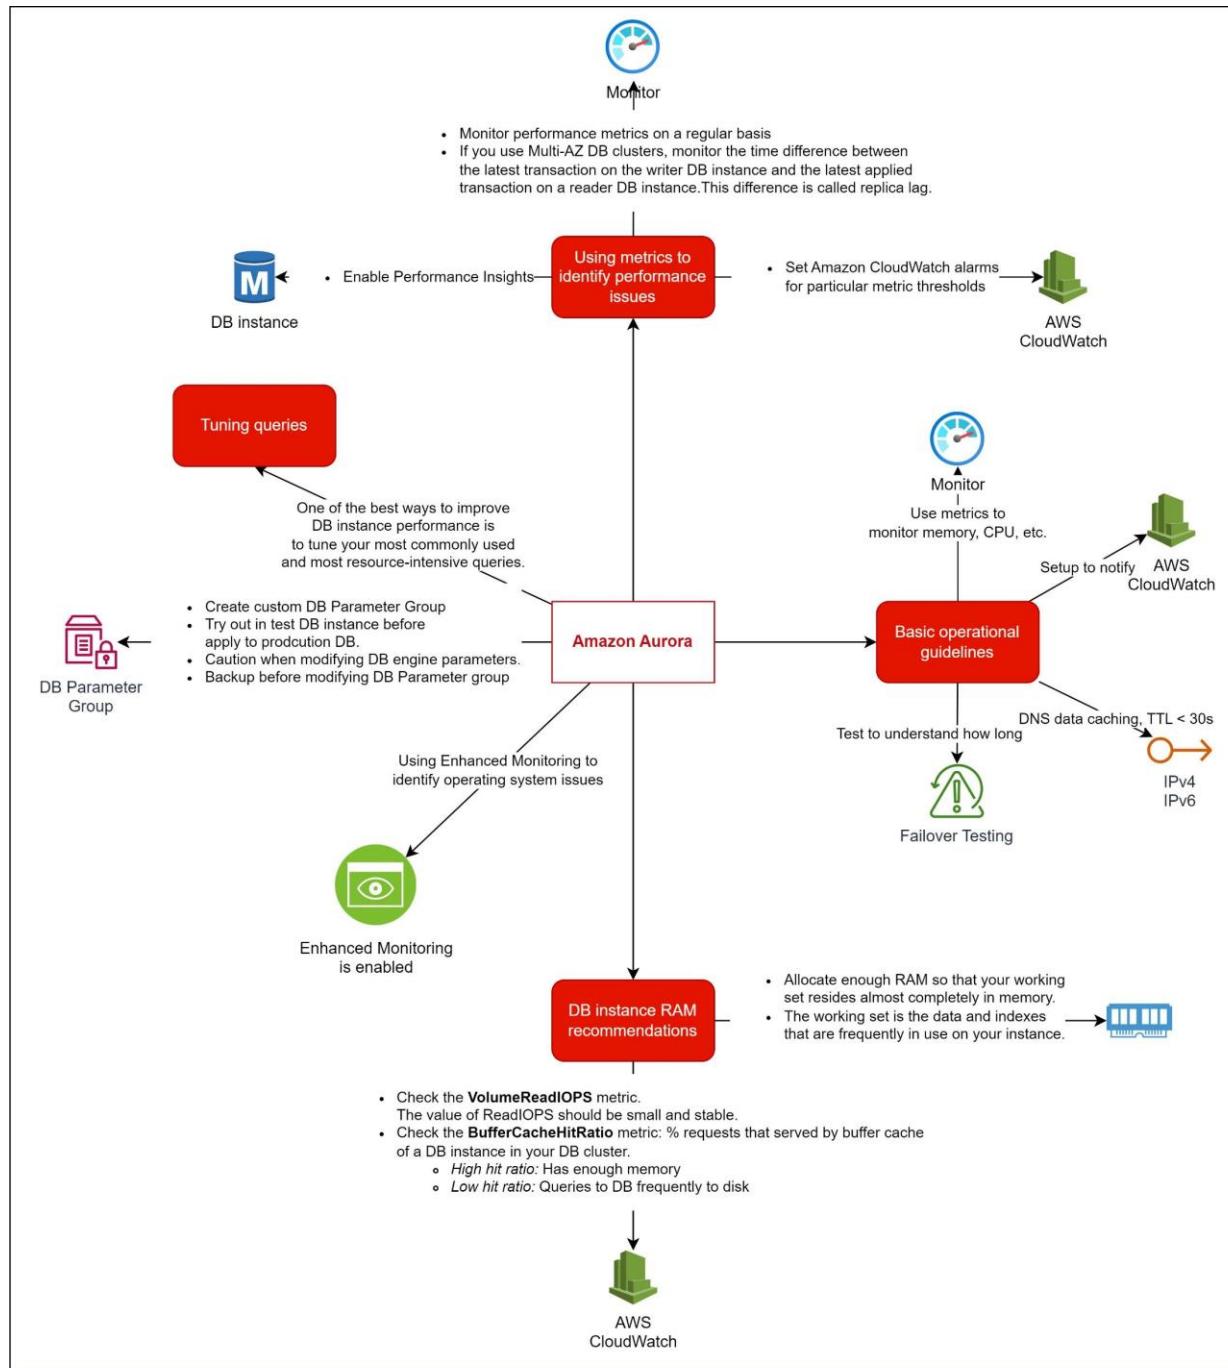


Amazon Aurora

Nash
Tech.

Best practices with Amazon Aurora

- [Basic operational guidelines for Amazon Aurora](#)
- [DB instance RAM recommendations](#)
- [Monitoring Amazon Aurora](#)
- [Working with DB parameter groups and DB cluster parameter groups](#)
- [Amazon Aurora best practices video](#)



Best practices with Amazon Aurora MySQL

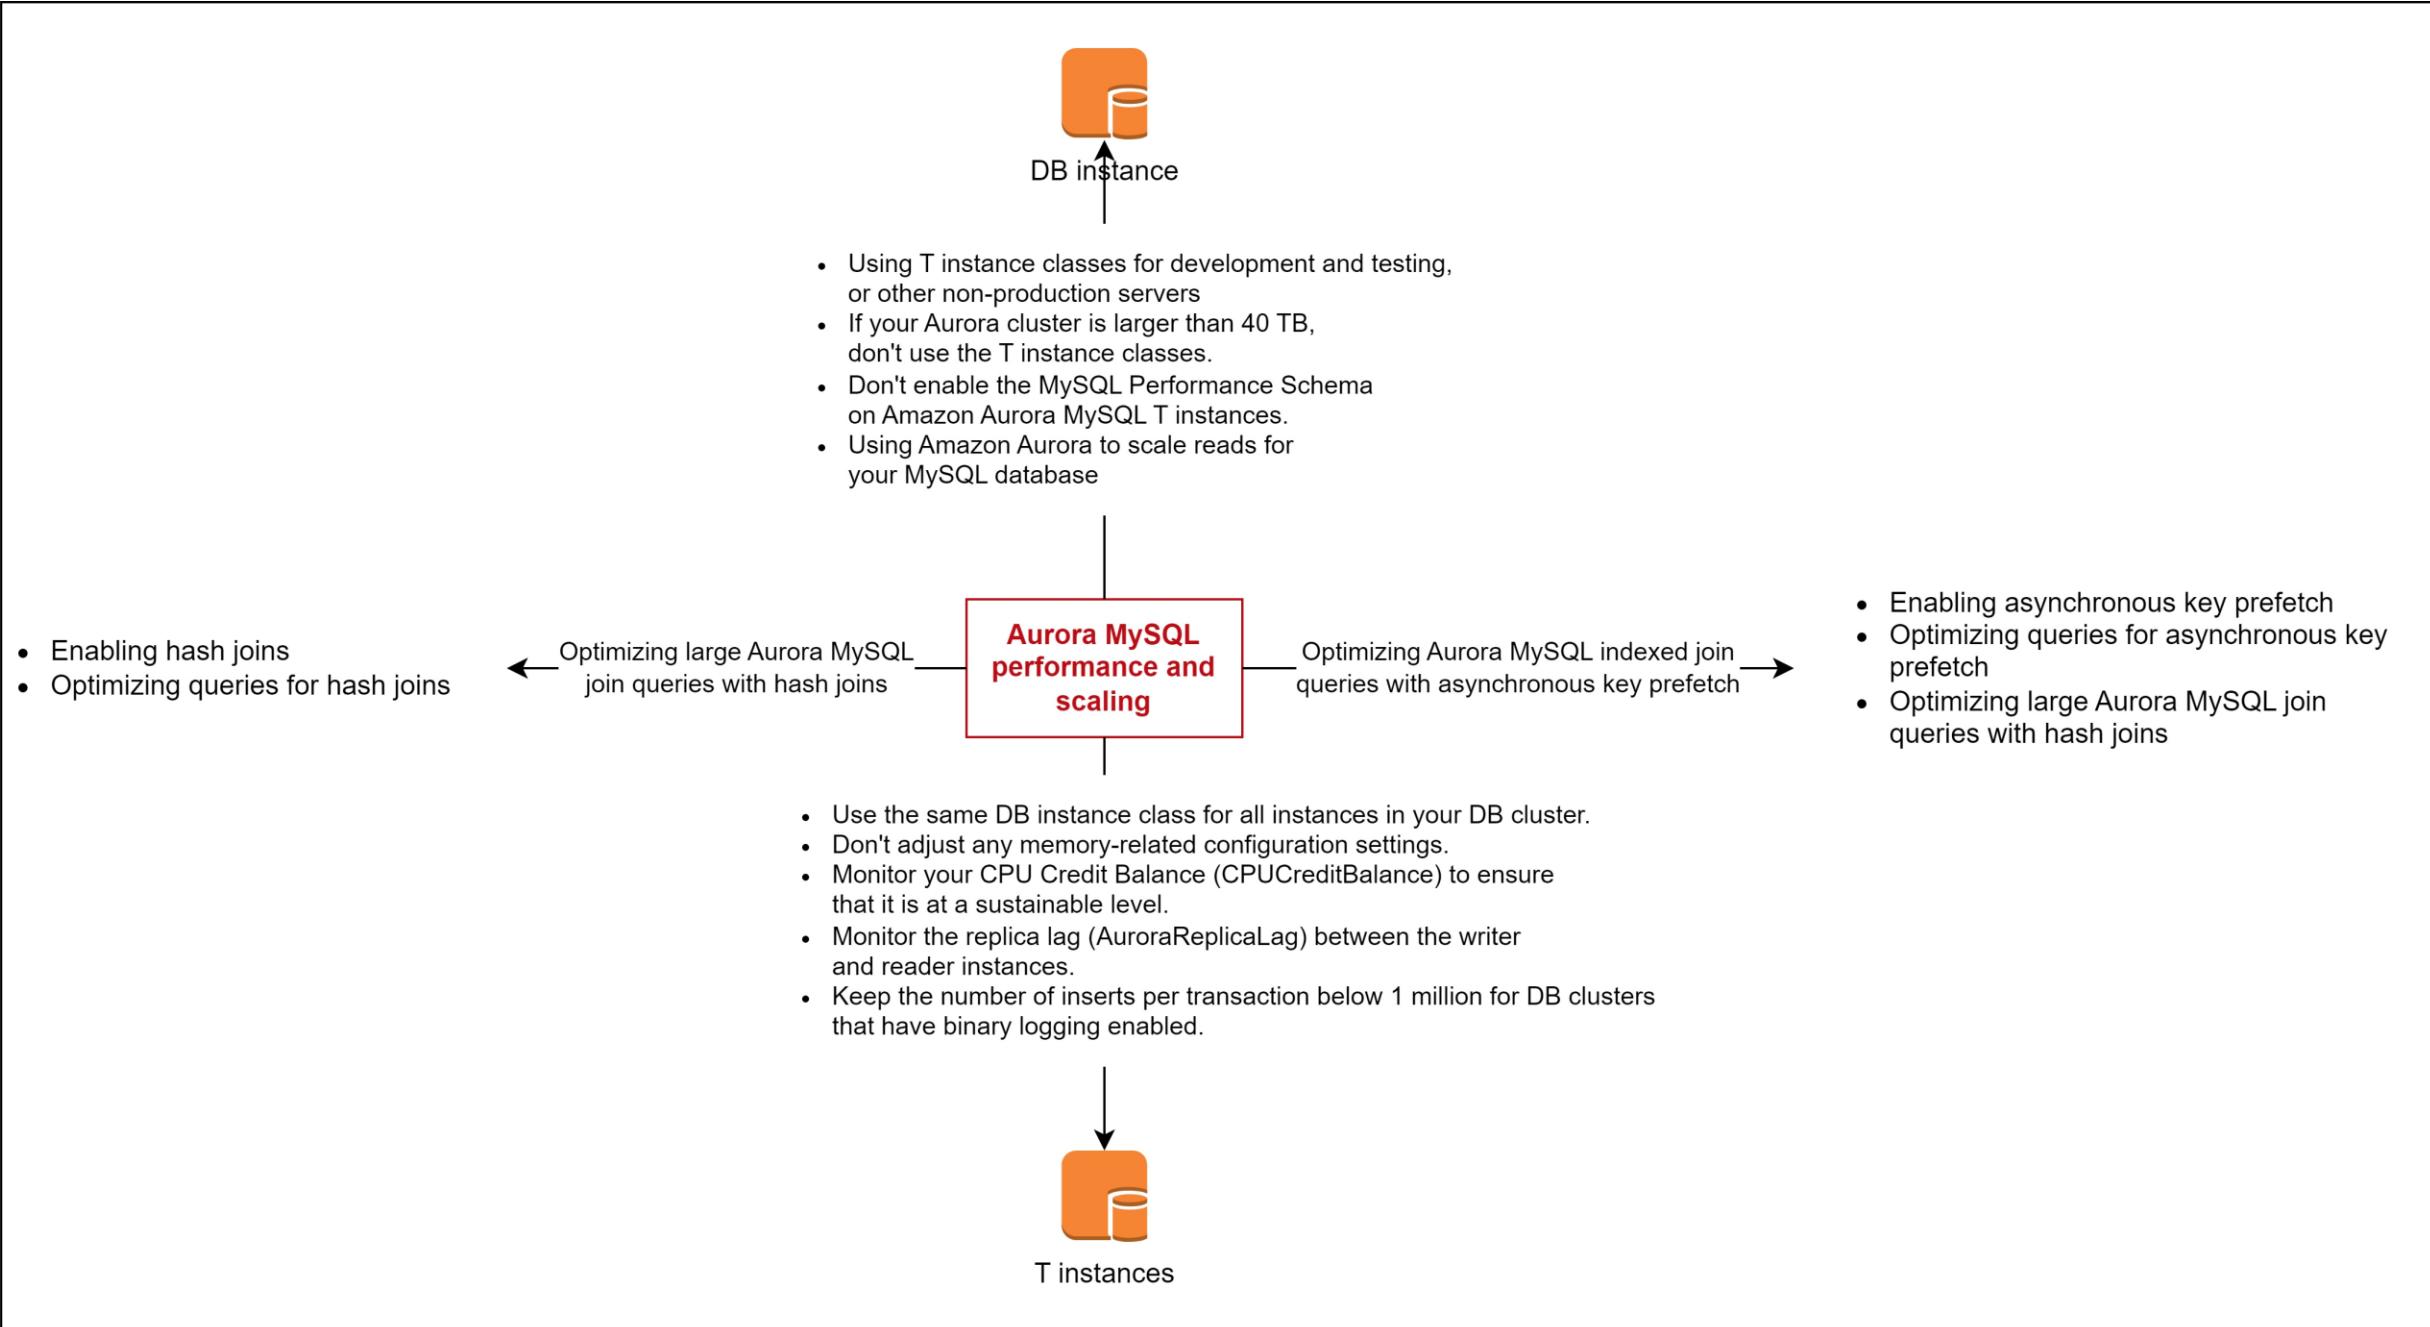
- [Determining which DB instance you are connected to](#)
- [Best practices for Aurora MySQL performance and scaling](#)
 - [Using T instance classes for development and testing](#)
 - [Optimizing Aurora MySQL indexed join queries with asynchronous key prefetch](#)
 - [Enabling asynchronous key prefetch](#)
 - [Optimizing queries for asynchronous key prefetch](#)
 - [Optimizing large Aurora MySQL join queries with hash joins](#)
 - [Enabling hash joins](#)
 - [Optimizing queries for hash joins](#)
 - [Using Amazon Aurora to scale reads for your MySQL database](#)
- [Best practices for Aurora MySQL high availability](#)
 - [Using Amazon Aurora for Disaster Recovery with your MySQL databases](#)
 - [Migrating from MySQL to Amazon Aurora MySQL with reduced downtime](#)
 - [Avoiding slow performance, automatic restart, and failover for Aurora MySQL DB instances](#)
- [Recommendations for Aurora MySQL](#)
 - [Using multithreaded replication in Aurora MySQL version 3](#)
 - [Invoking AWS Lambda functions using native MySQL functions](#)
 - [Avoiding XA transactions with Amazon Aurora MySQL](#)
 - [Keeping foreign keys turned on during DML statements](#)
 - [Configuring how frequently the log buffer is flushed](#)
 - [Minimizing and troubleshooting Aurora MySQL deadlocks](#)
 - [Minimizing InnoDB deadlocks](#)
 - [Monitoring InnoDB deadlocks](#)

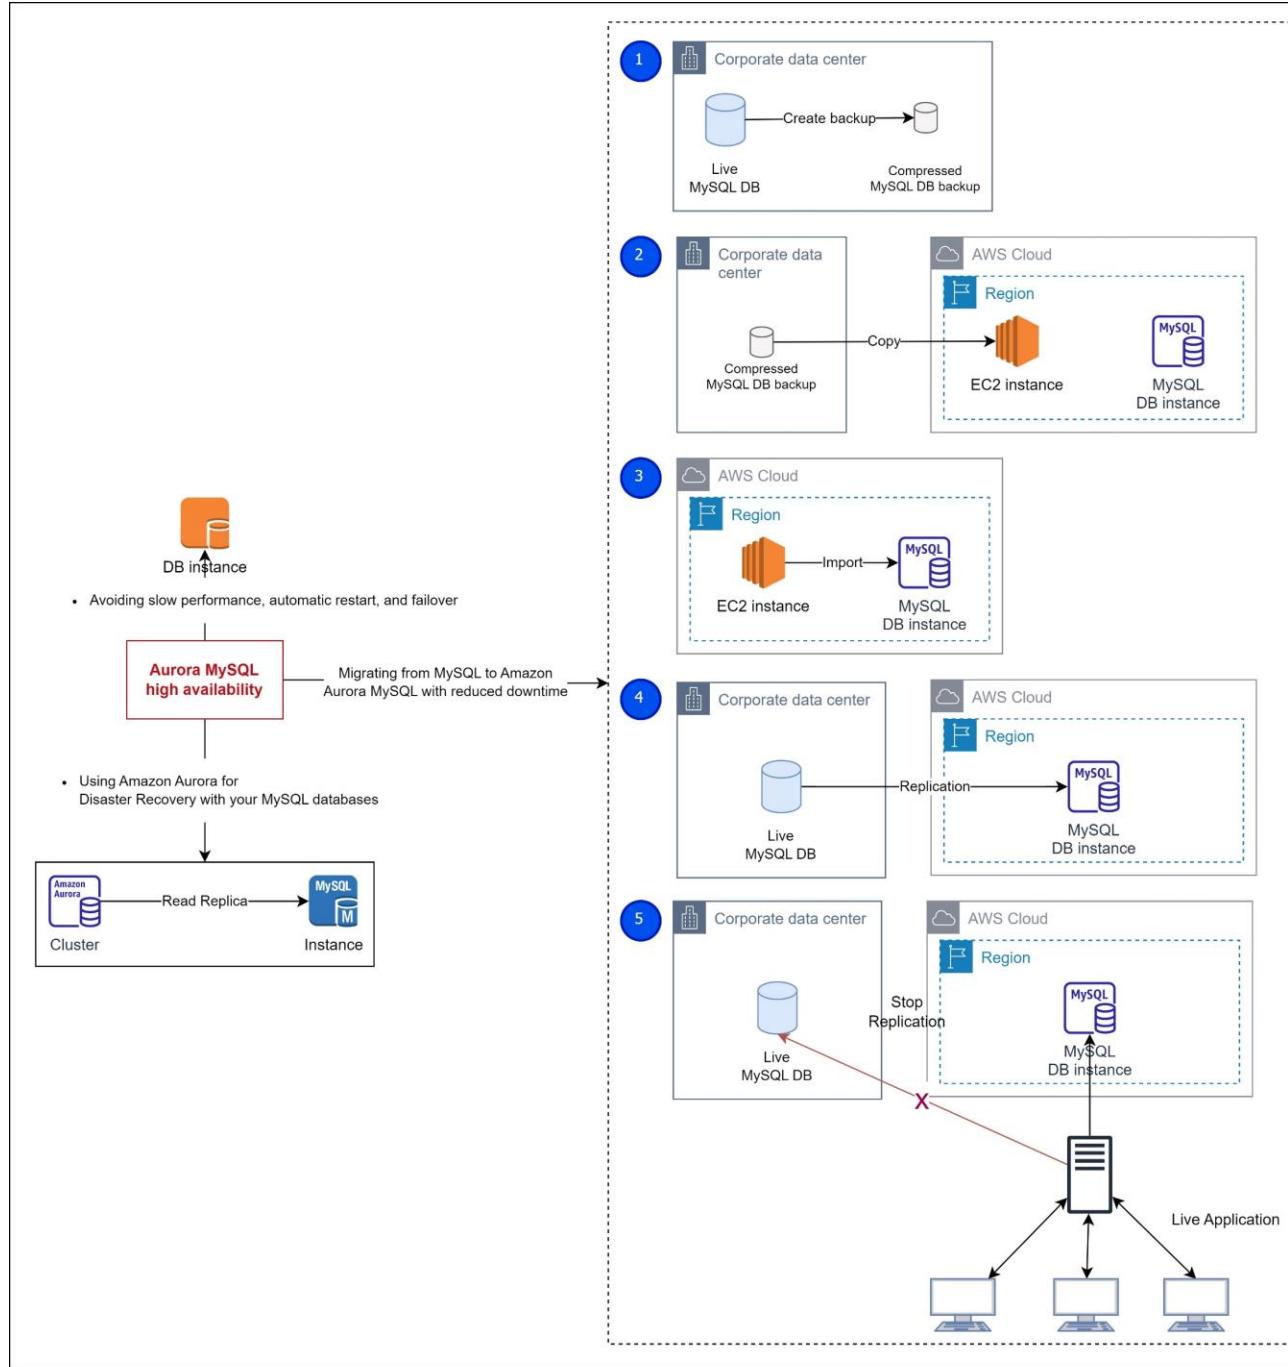
Determining which DB instance

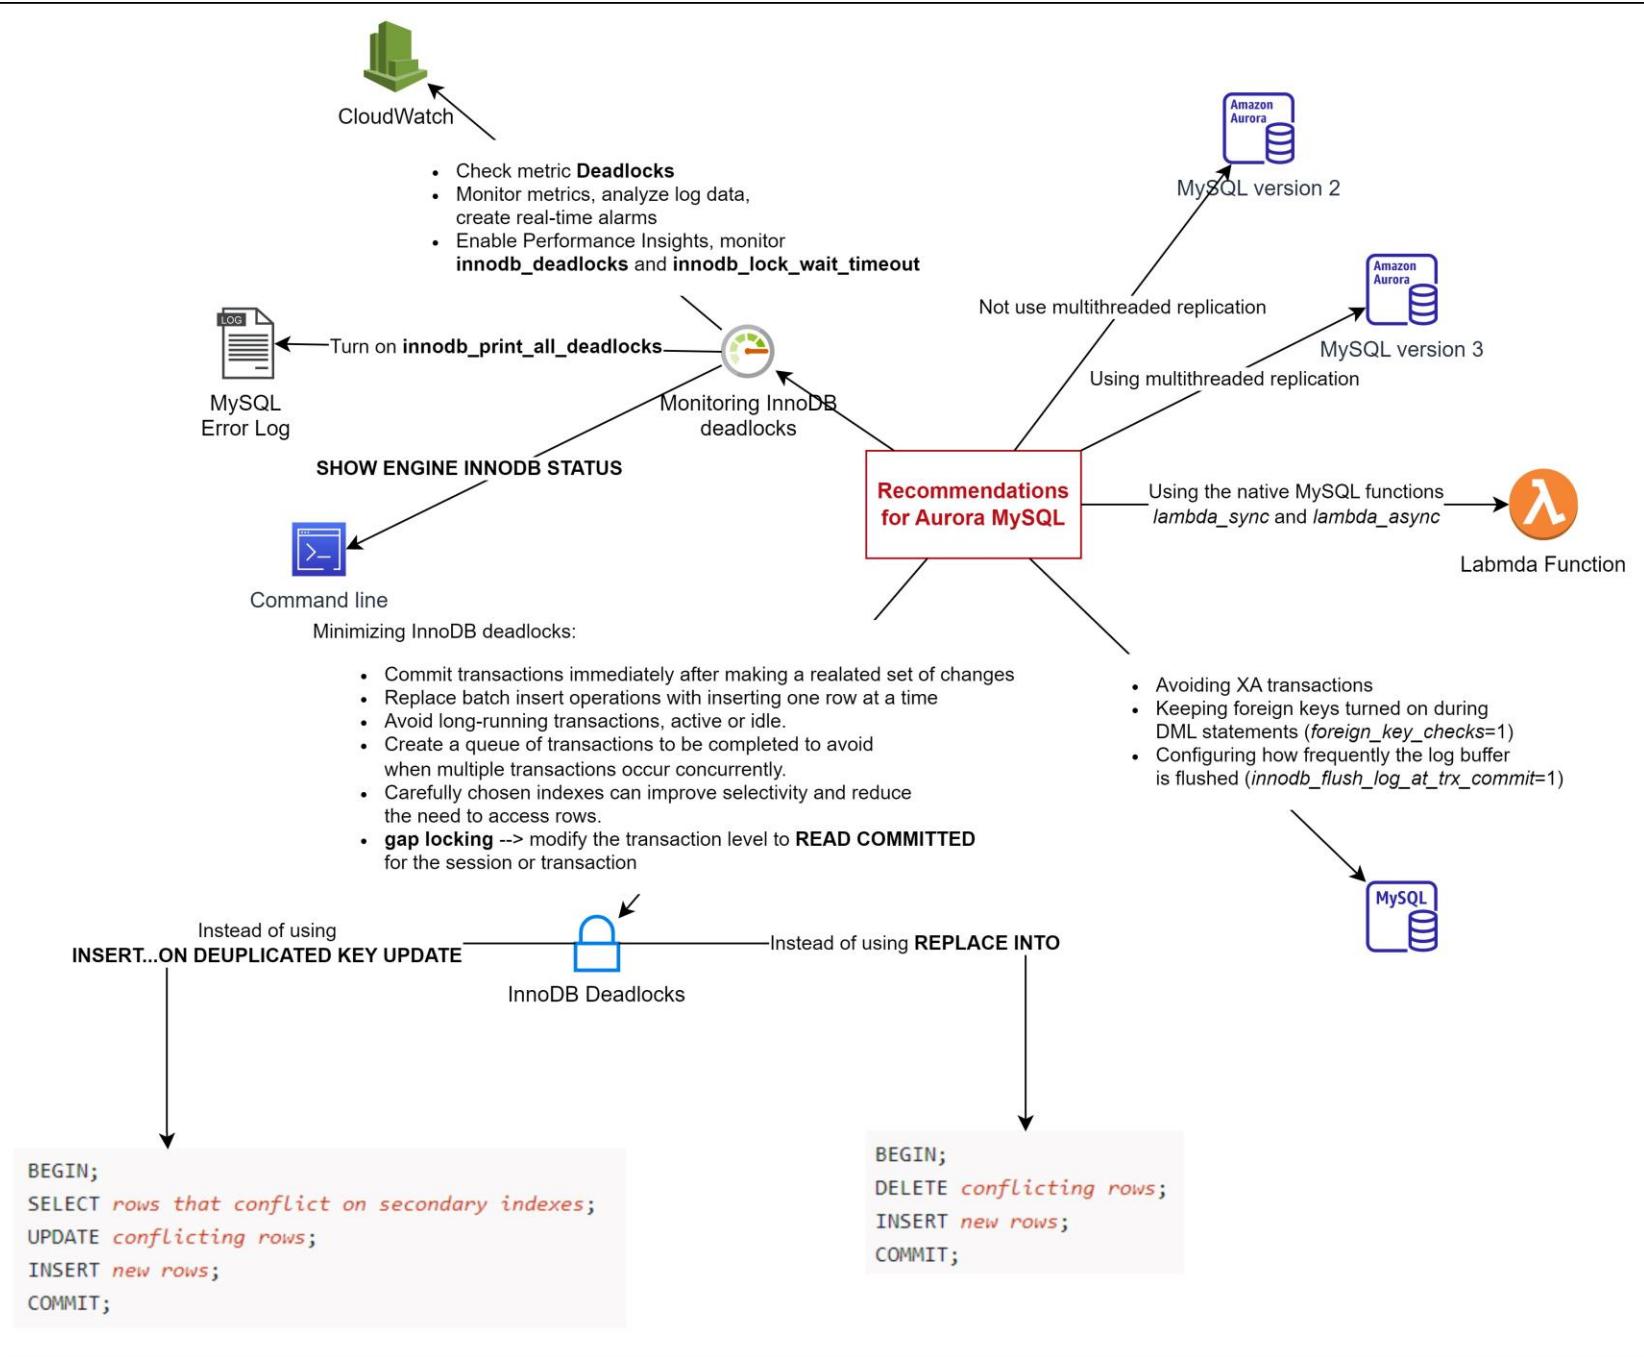
- To determine which DB instance in an Aurora MySQL DB cluster a connection is connected to, check the `innodb_read_only` global variable, as shown in the following example.

```
SHOW GLOBAL VARIABLES LIKE 'innodb_read_only';
```

- The `innodb_read_only` variable is set to ON if you are connected to a reader DB instance. This setting is OFF if you are connected to a writer DB instance, such as a primary instance in a provisioned cluster.
- This approach can be helpful if you want to add logic to your application code to balance the workload or ensure that a write operation uses the correct connection.

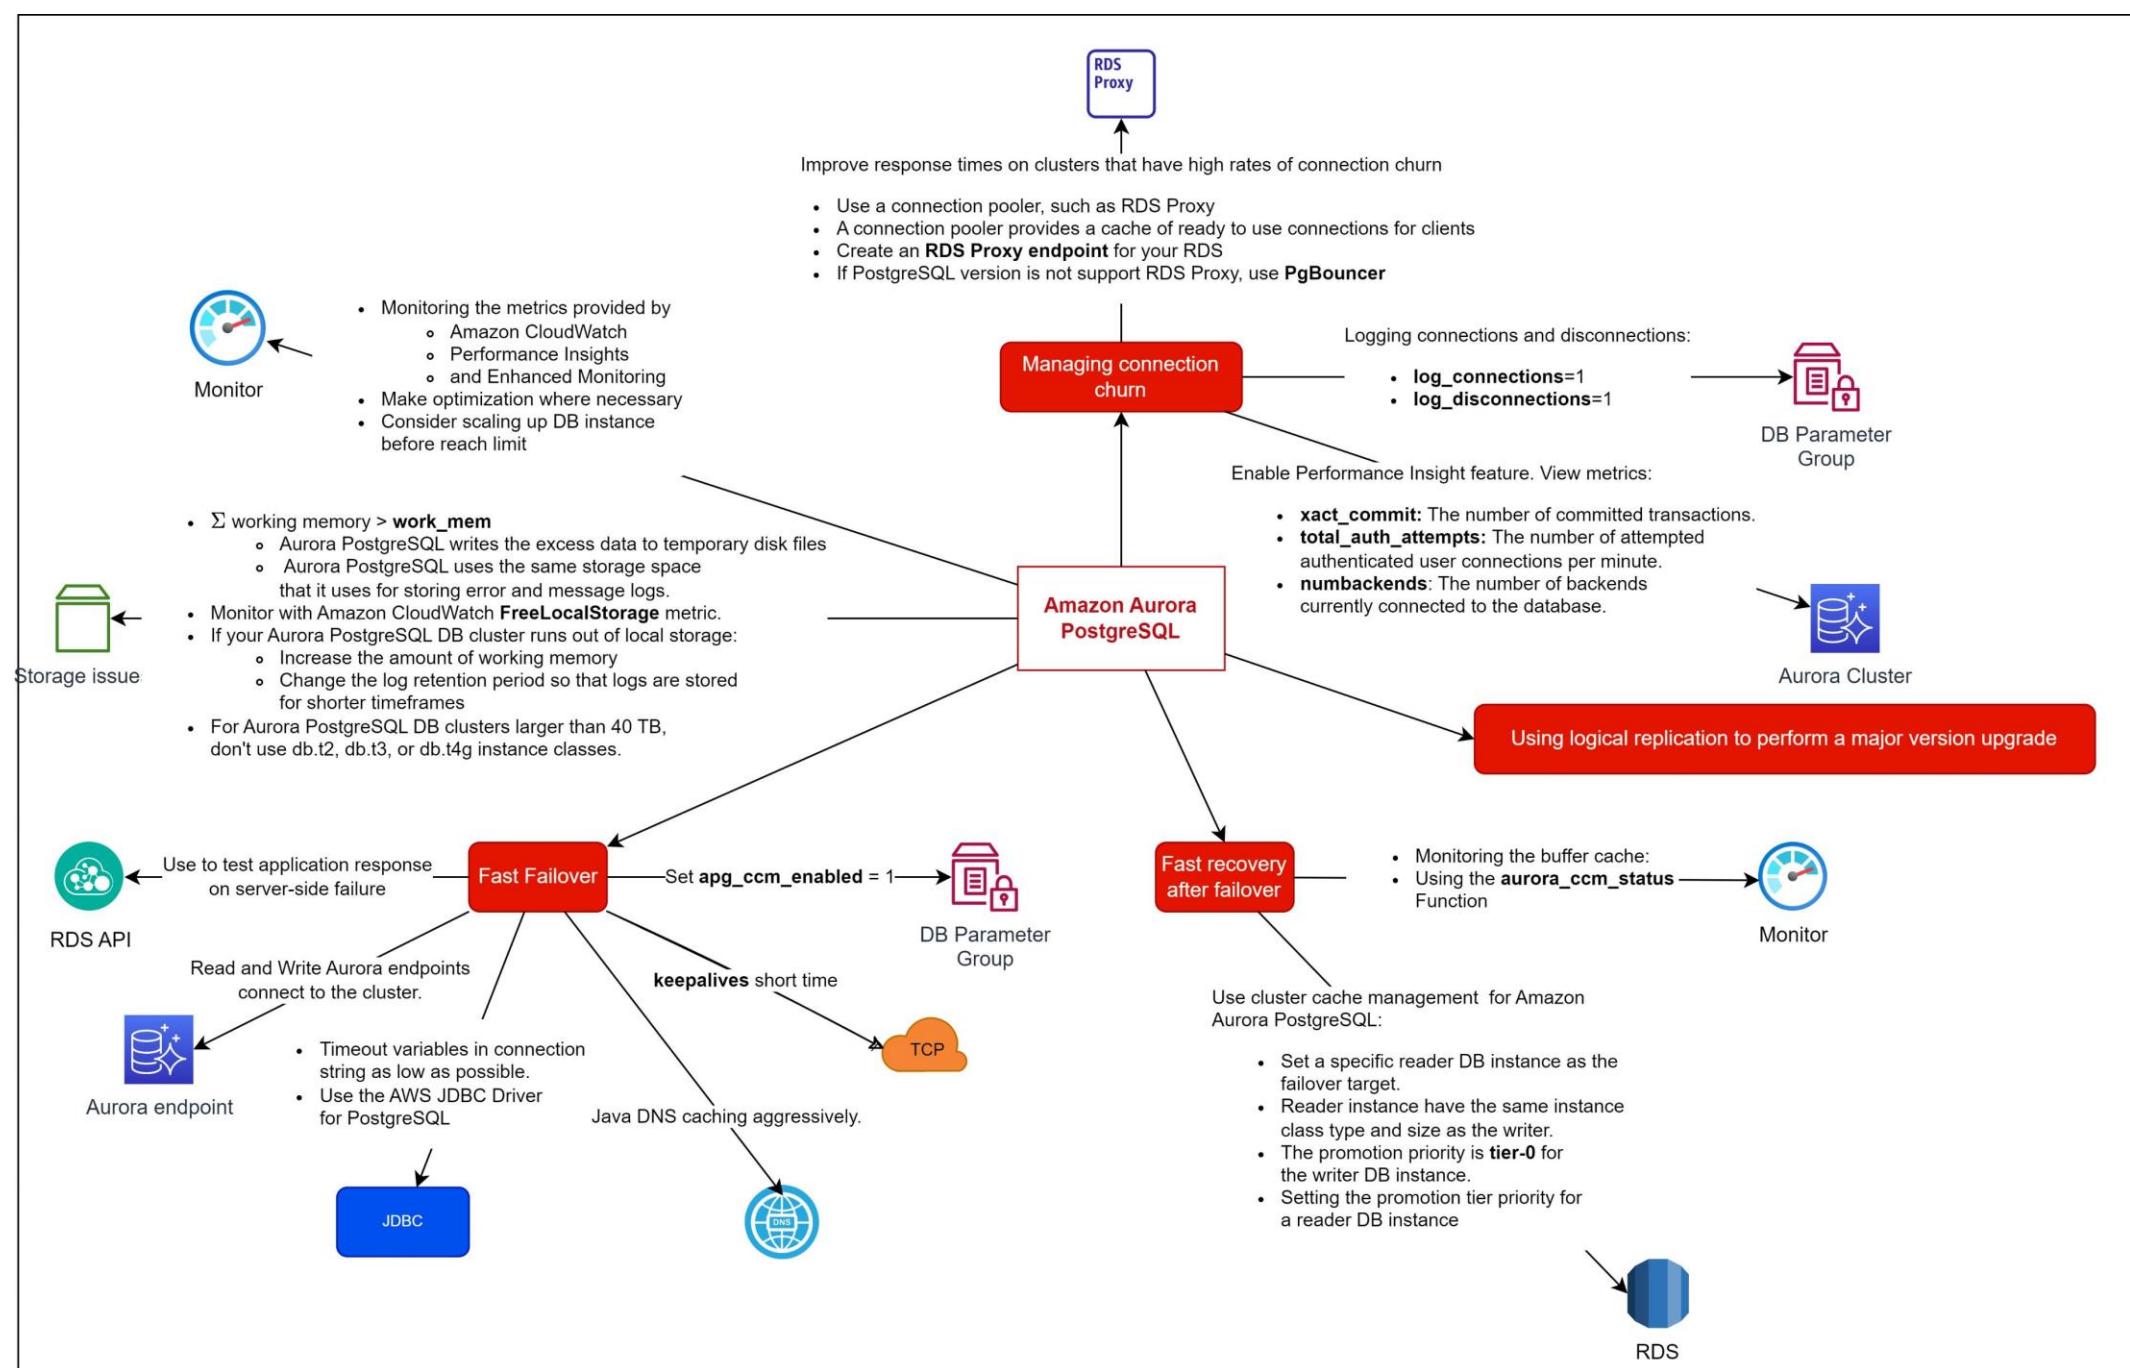






Best practices with Amazon Aurora PostgreSQL

- [Troubleshooting storage issues](#)
- [Avoiding slow performance, automatic restart, and failover for Aurora PostgreSQL DB instances](#)
- [Fast failover with Amazon Aurora PostgreSQL](#)
- [Fast recovery after failover with cluster cache management for Aurora PostgreSQL](#)
- [Managing Aurora PostgreSQL connection churn with pooling](#)
- [Using logical replication to perform a major version upgrade for Aurora PostgreSQL](#)
- [Tuning memory parameters for Aurora PostgreSQL](#)
- [Using Amazon CloudWatch metrics to analyze resource usage for Aurora PostgreSQL](#)
- [Diagnosing table and index bloat](#)

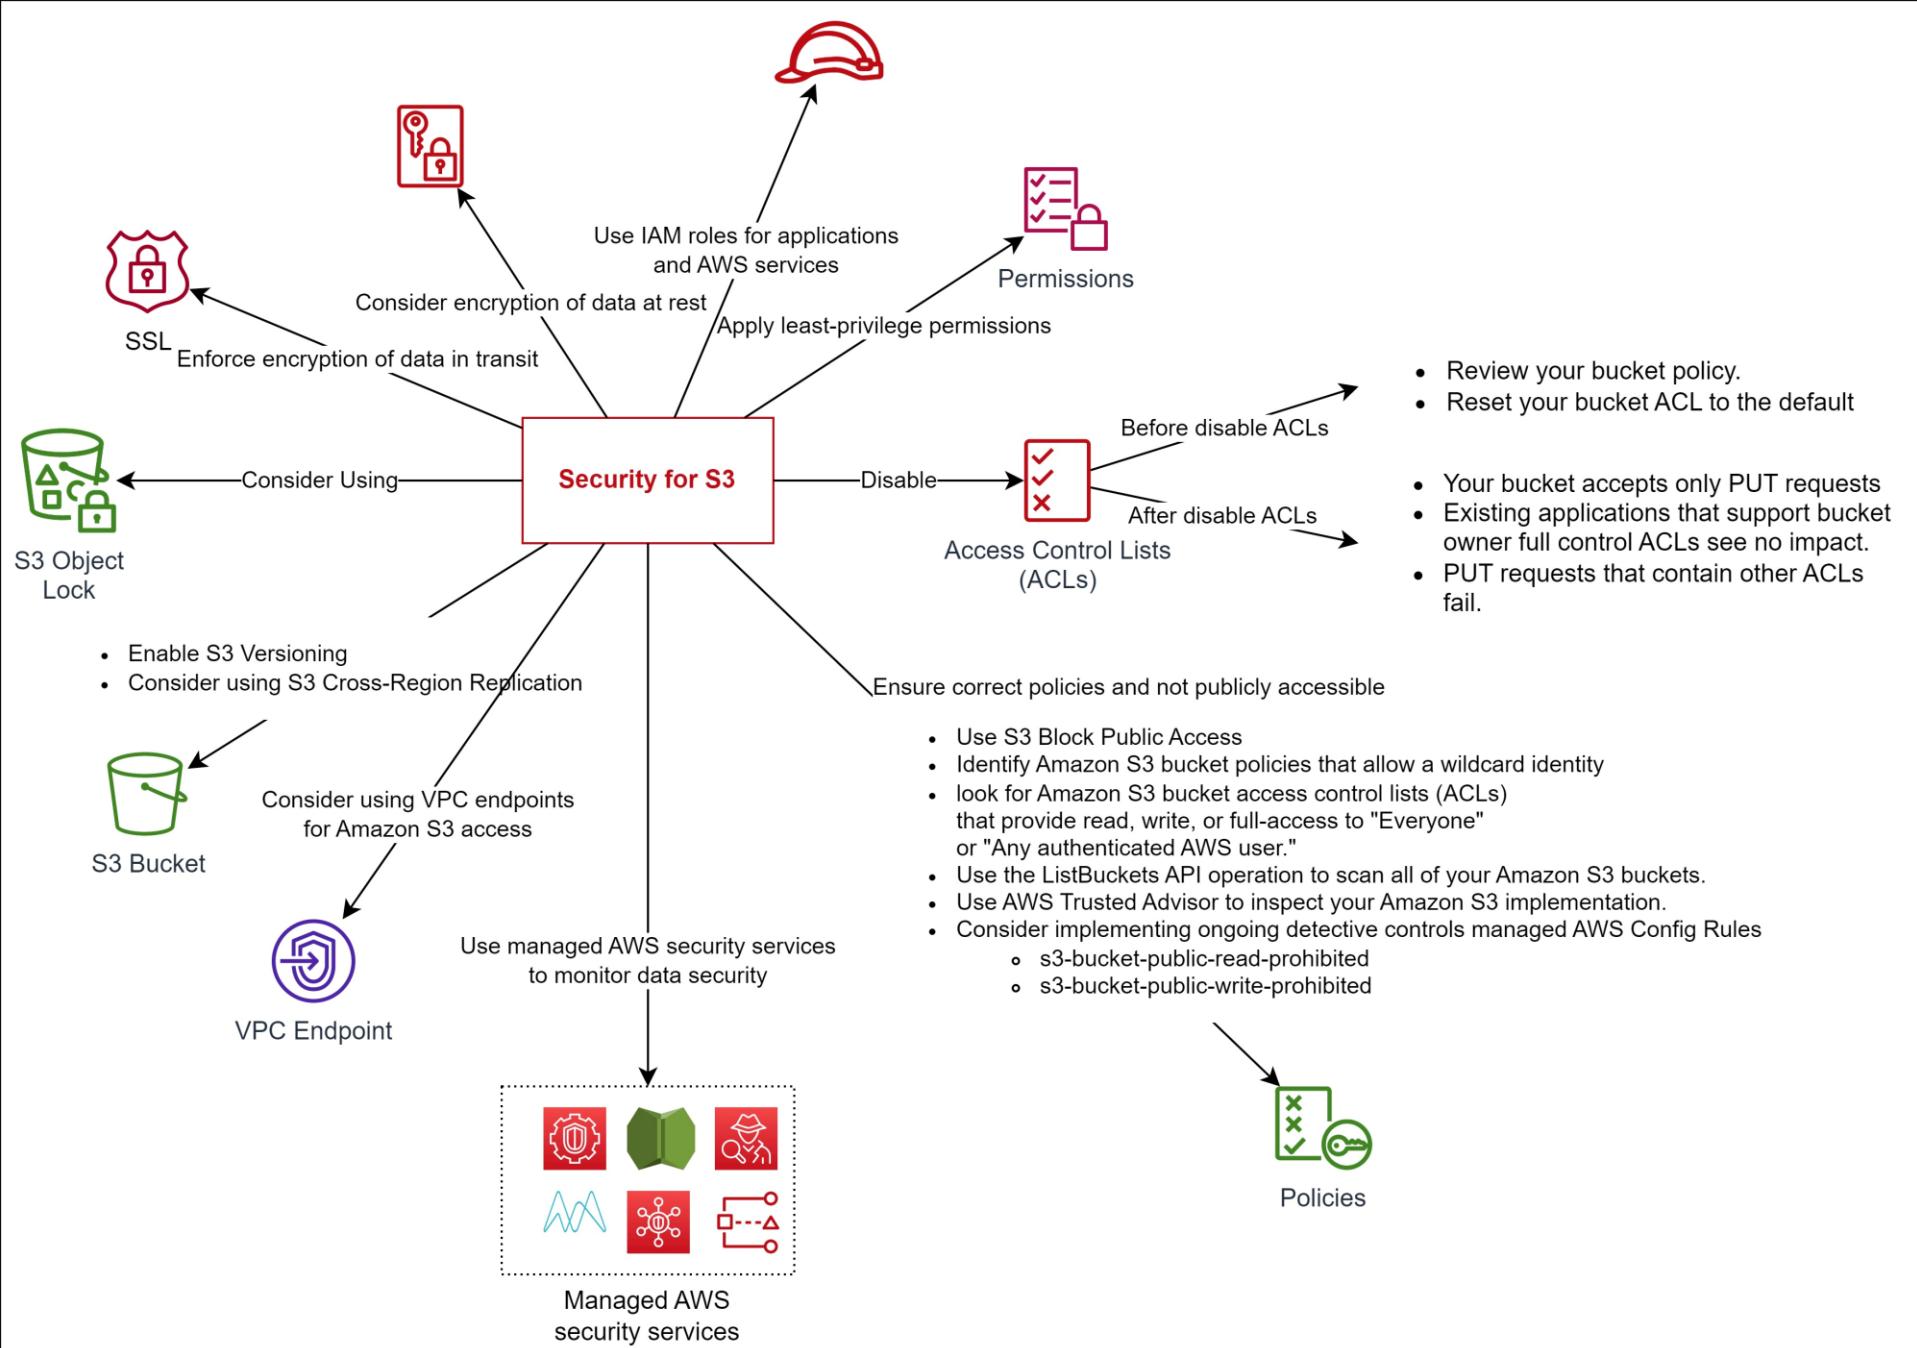


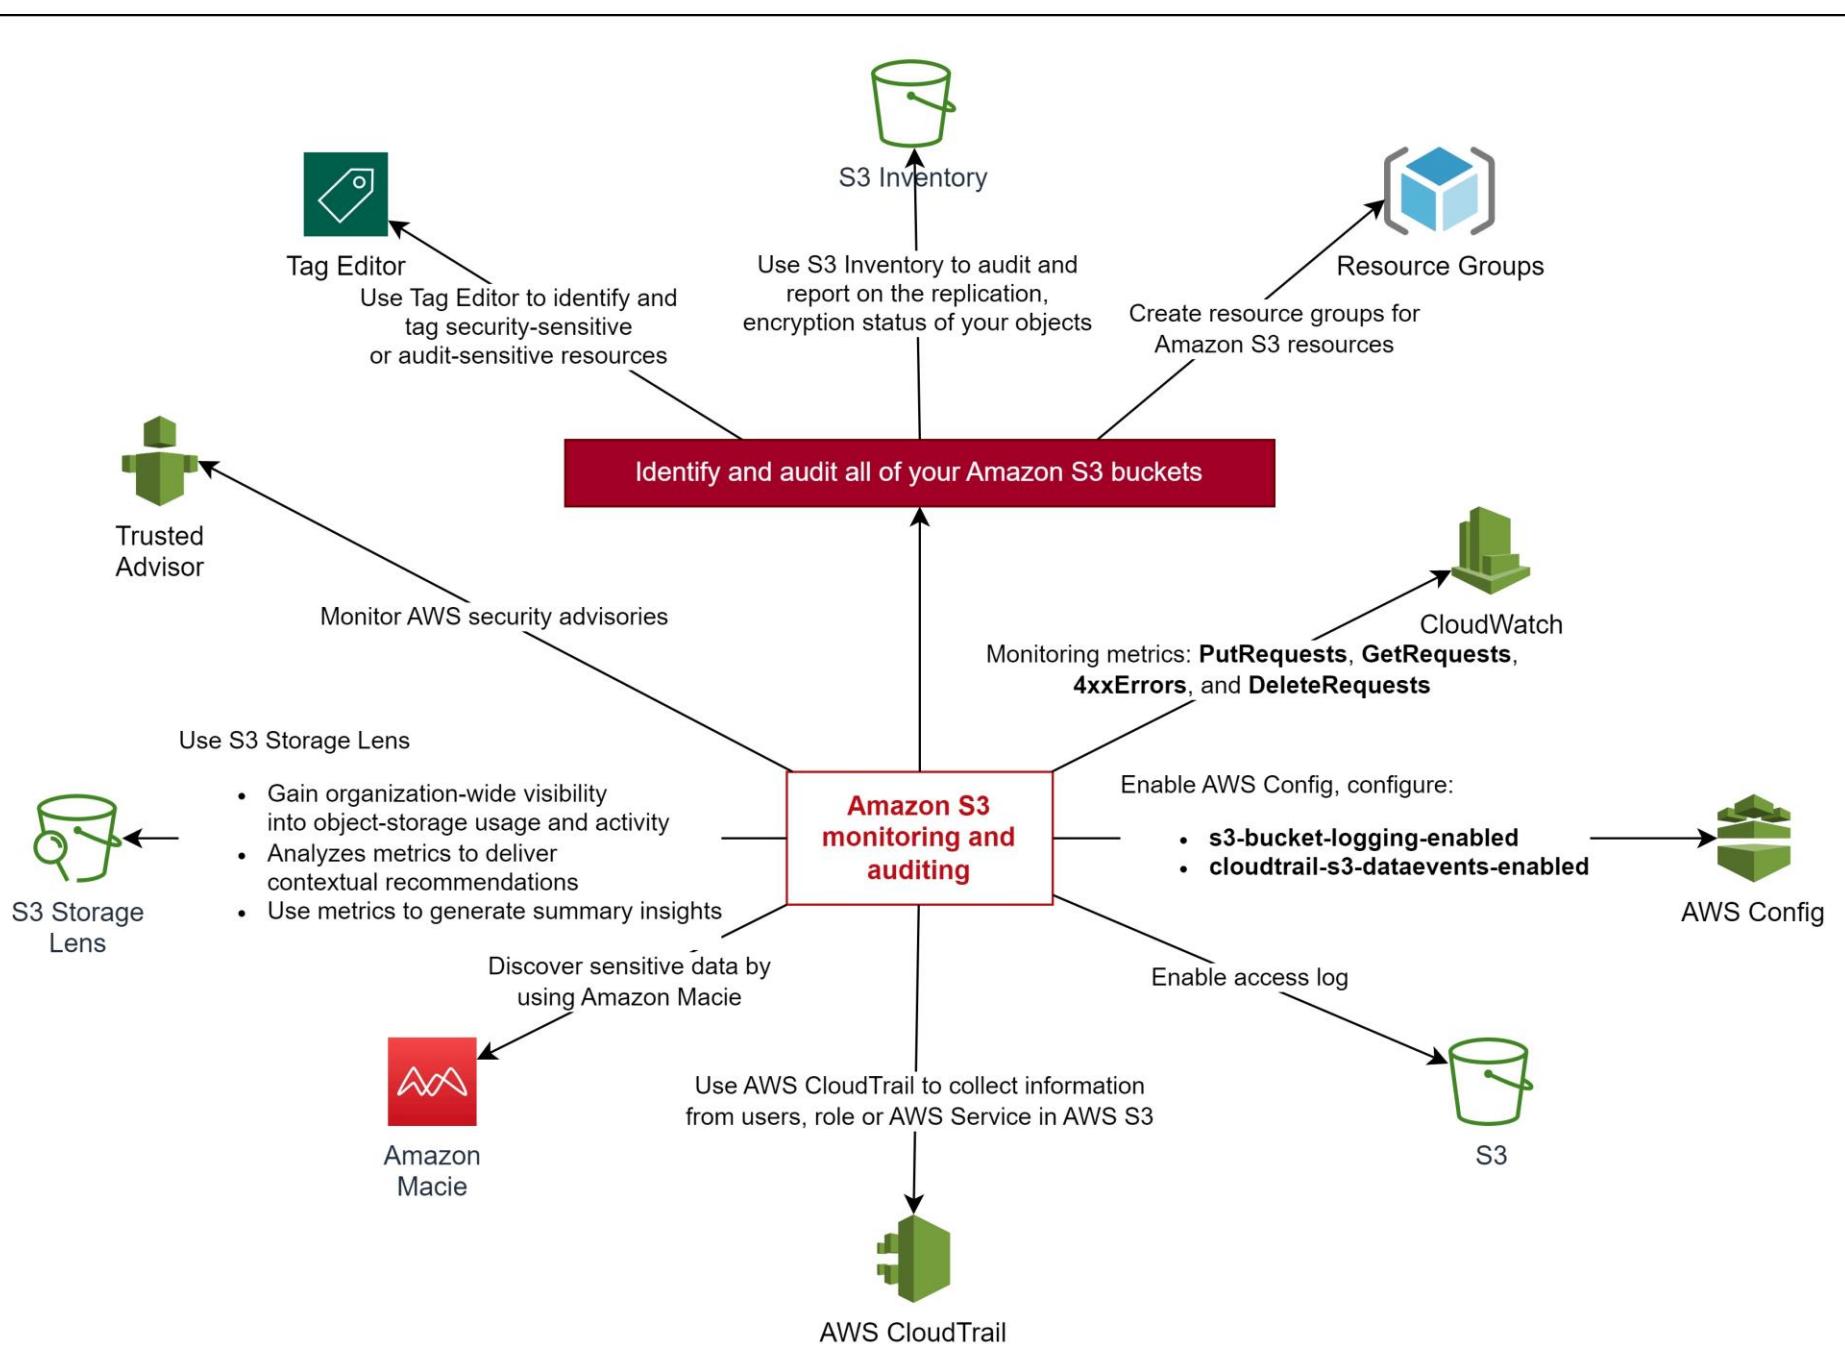
Amazon S3 (Simple Storage Service)

Nash
Tech.

Security best practice for Amazon S3

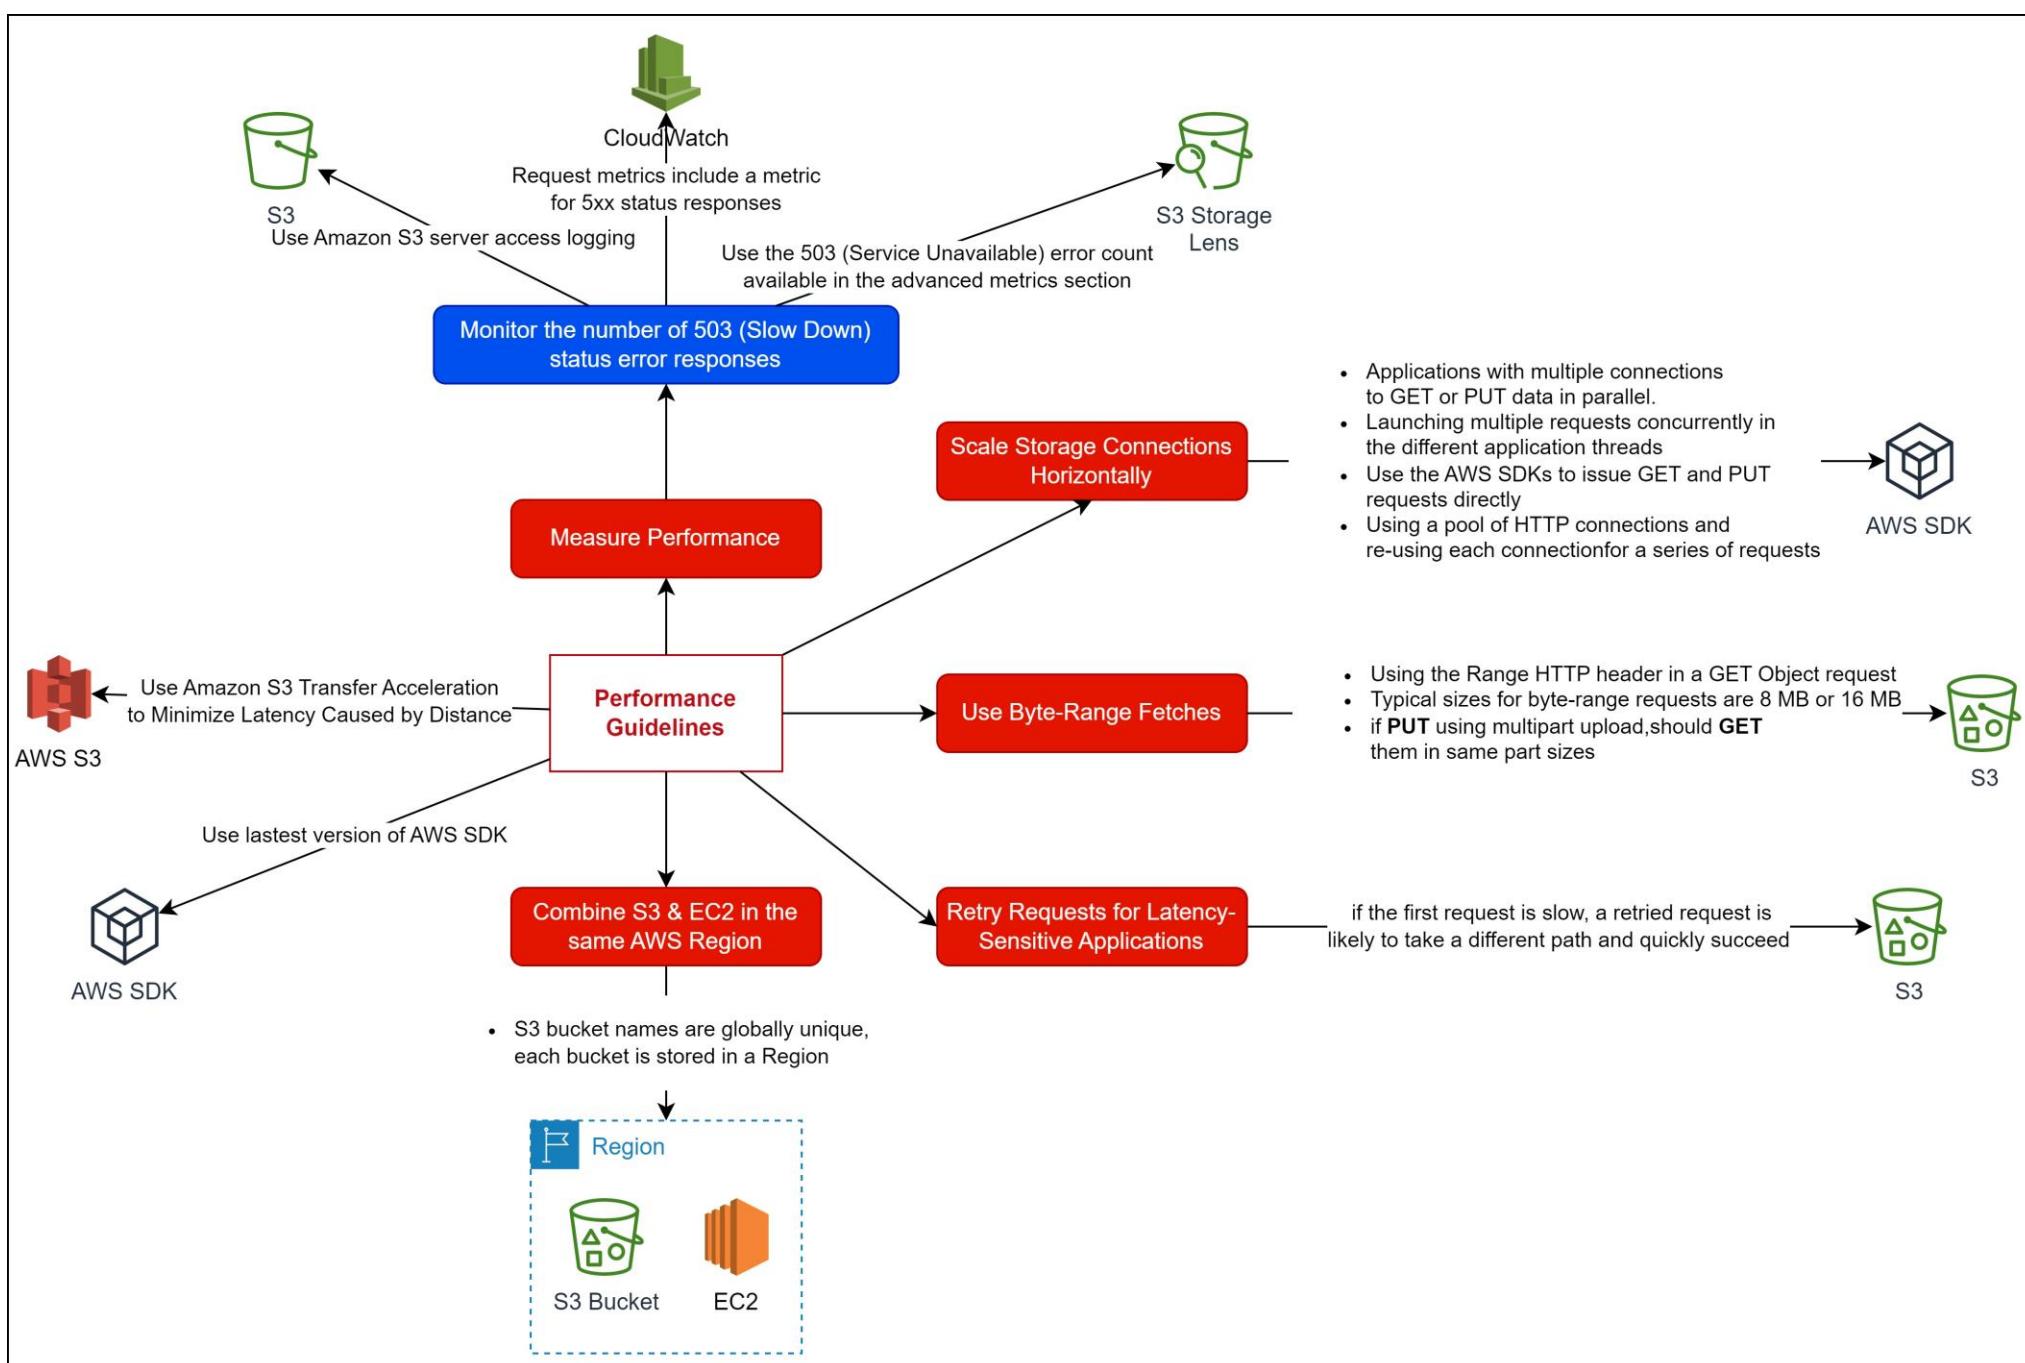
- [Amazon S3 security best practices](#)
- [Amazon S3 monitoring and auditing best practices](#)

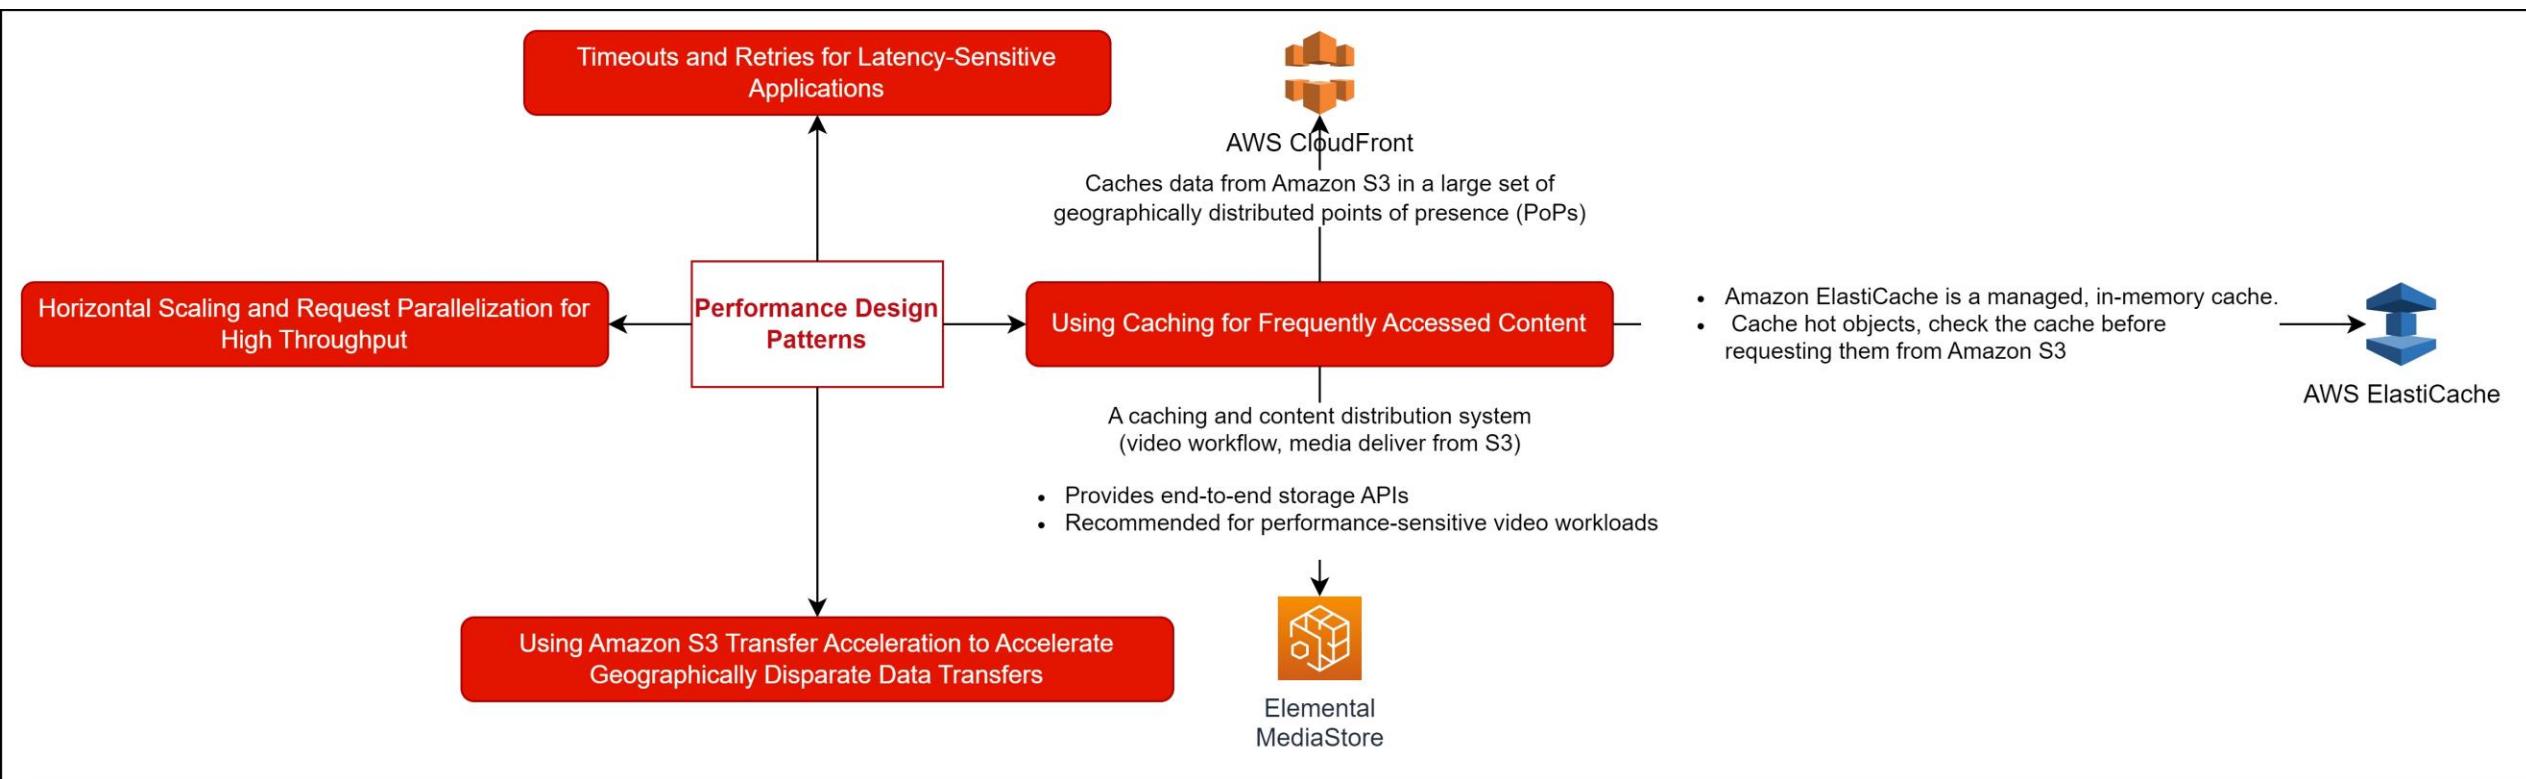




Optimizing Amazon S3 performance

- [Performance Guidelines](#)
- [Performance Design Patterns](#)



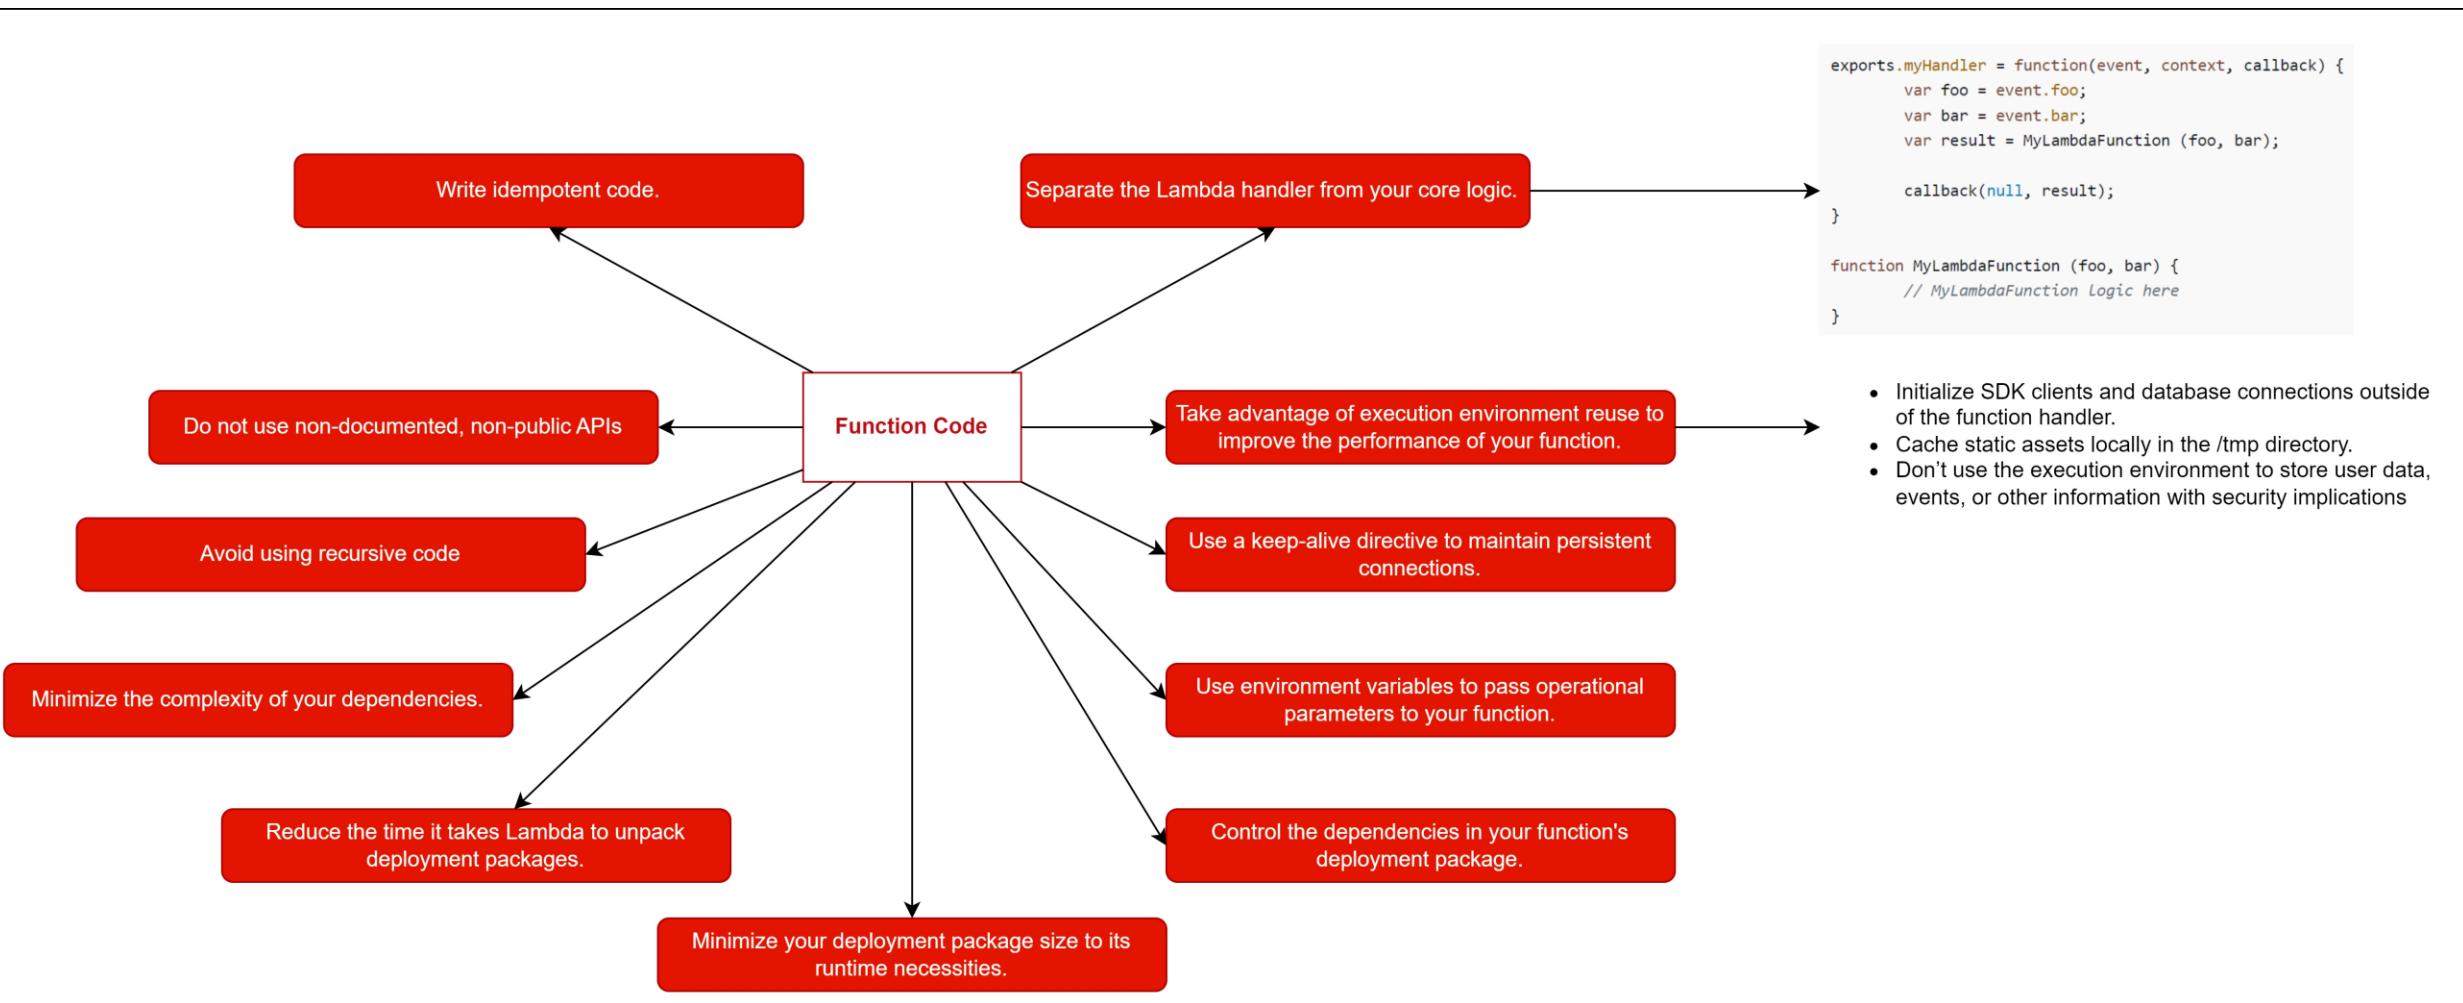


Amazon Lambda

Nash
Tech.

Best practices for working with AWS Lambda functions

- [Function code](#)
- [Function configuration](#)
- [Metrics and alarms](#)
- [Working with streams](#)
- [Security best practices](#)





AWS CloudWatch

Check **Max Memory Used**

Performance testing your Lambda function

Performance testing your Lambda function

Function configuration

Delete Lambda functions that you are no longer using

Use most-restrictive permissions when setting IAM policies.

Be familiar with Lambda quotas.

If you are using Amazon Simple Queue Service as an event source

function's expected invocation time < Visibility Timeout



AWS CloudWatch

- Use Working with Lambda function metrics and CloudWatch Alarms
- Leverage your logging library and AWSLambda Metrics and Dimensions to catch app errors

Metrics and alarms

Use **Cost Anomaly Detection** feature
to detect unusual activity on your account



Cost Explorer

Working with streams

Test with different batch and record sizes

Increase Kinesis stream processing throughput by adding shards

Use Amazon CloudWatch on IteratorAge to determine if your Kinesis stream is being processed

Security

Monitor your usage of AWS Lambda as it relates to security best practices

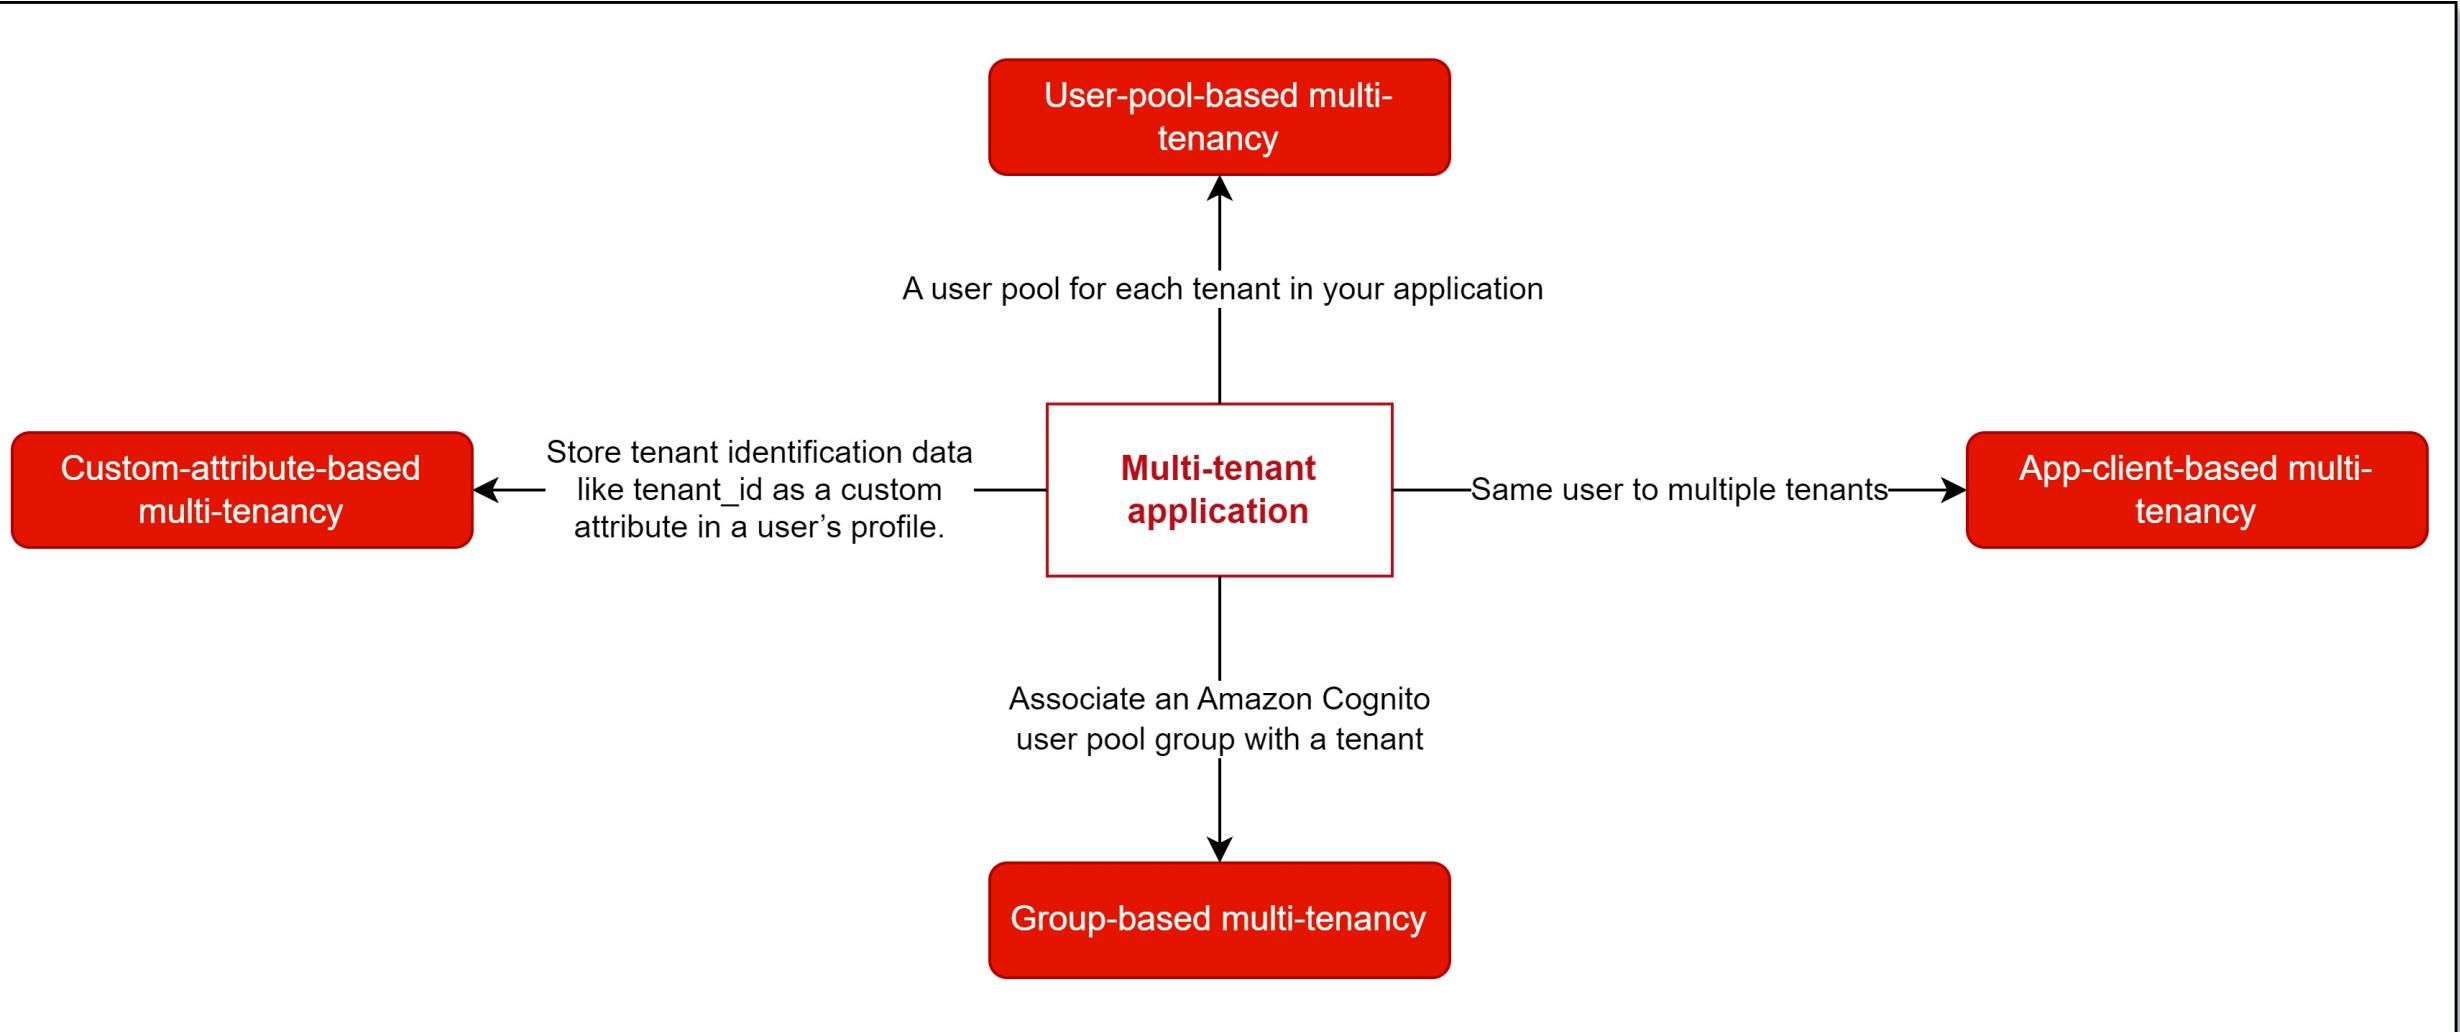


Amazon Cognito

Nash
Tech.

Multi-tenant application best practice

- [User-pool-based multi-tenancy](#)
- [App-client-based multi-tenancy](#)
- [Group-based multi-tenancy](#)
- [Custom-attribute-based multi-tenancy](#)
- [Multi-tenancy security recommendations](#)



Multi-tenancy security recommendations

Use only a verified email address to authorize user access to a tenant based on domain match.

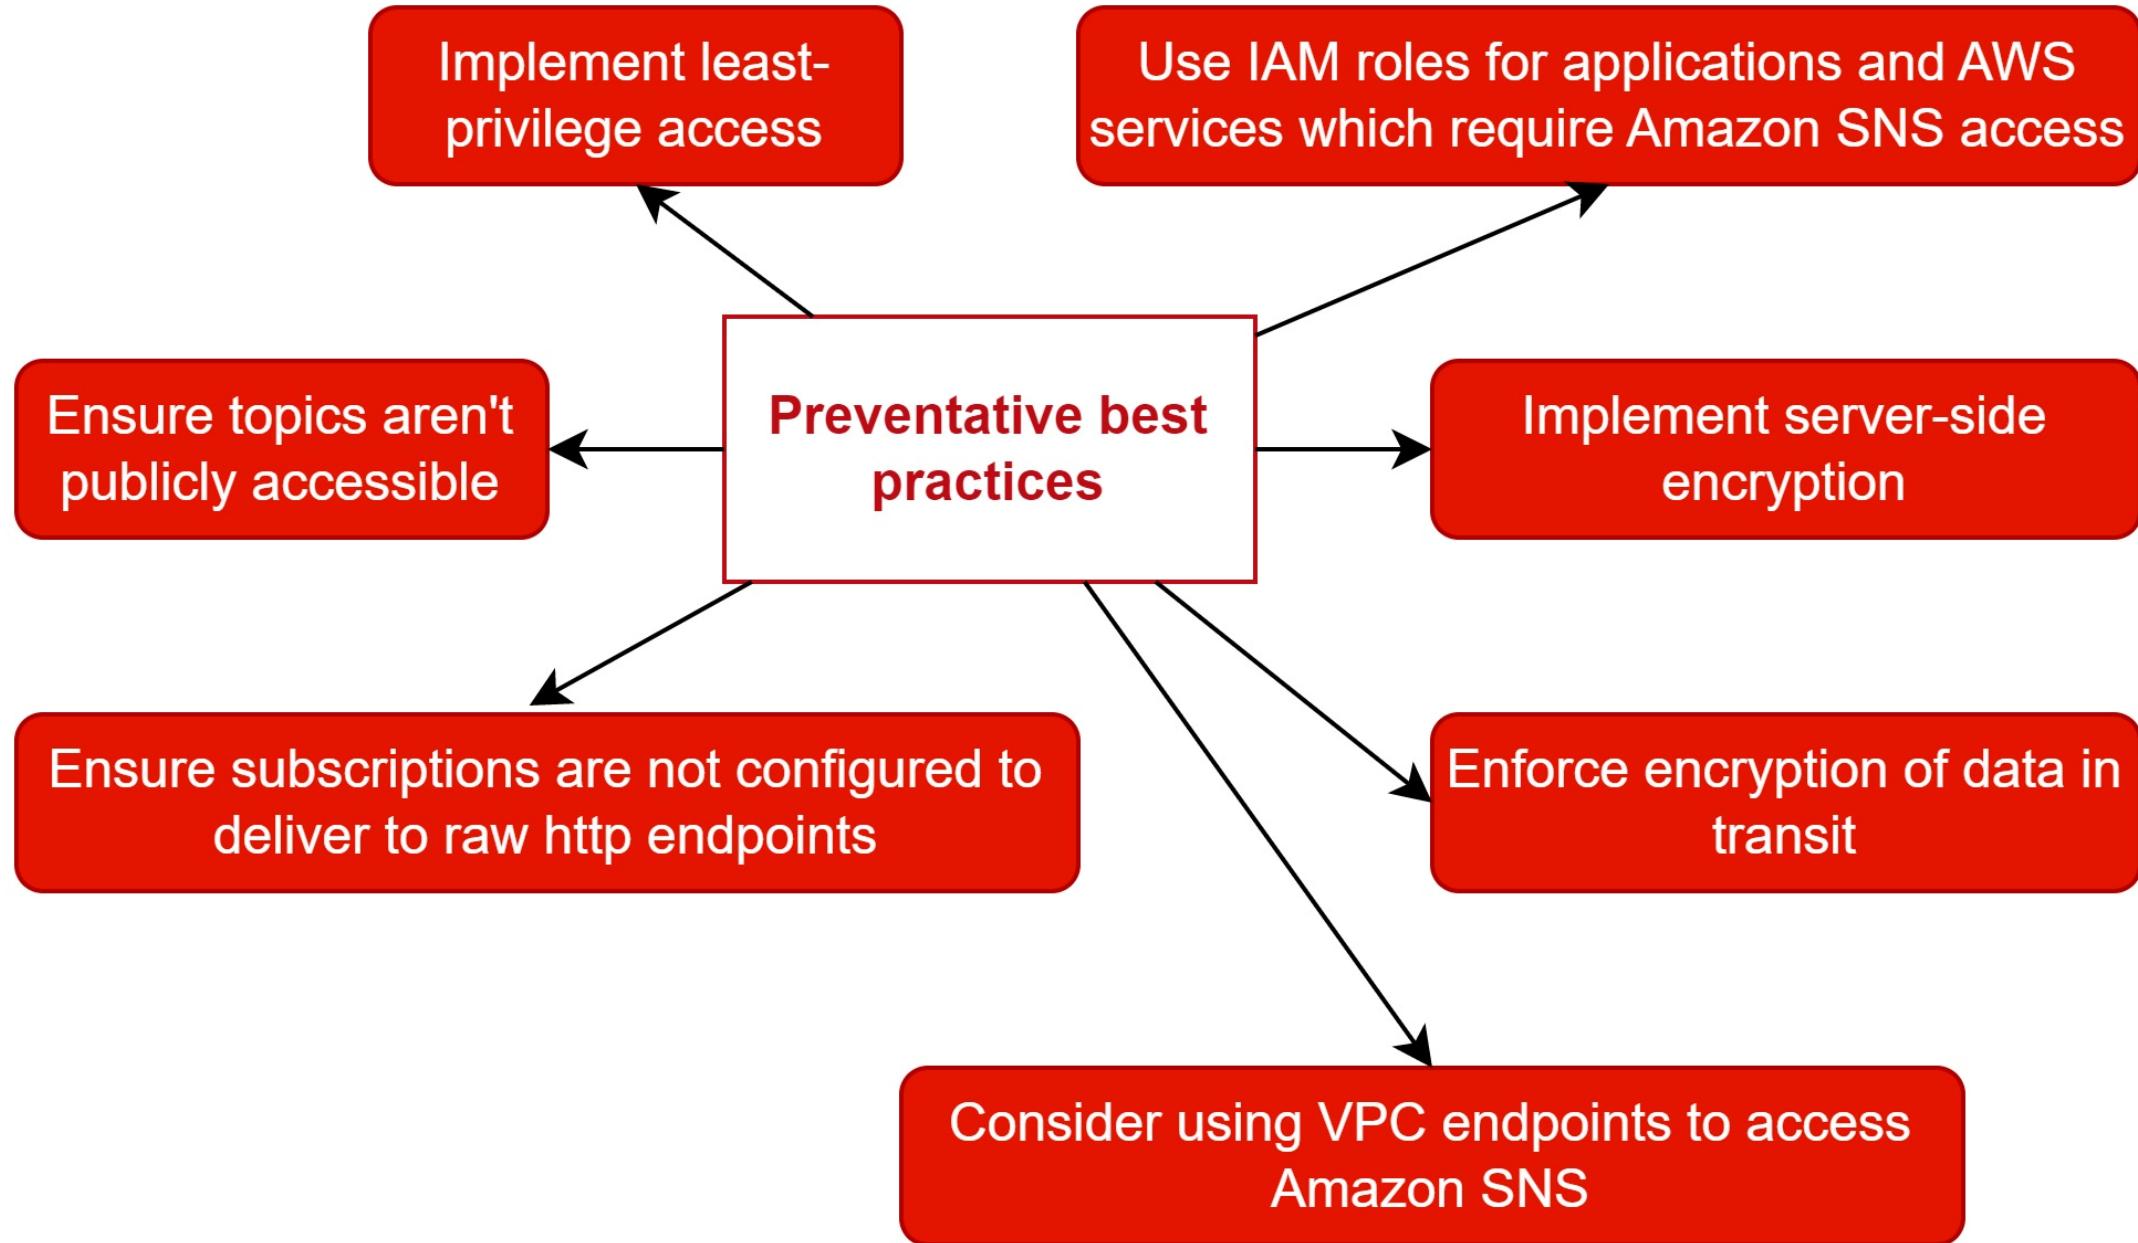
Use immutable, or read-only, attributes for the user profile attributes that identify tenants.

Use 1:1 mapping between external IdP and application client to prevent unauthorized cross-tenant access.

When you implement tenant-matching and authorization logic in your application, restrict users so that they can't modify the criteria that authorize user access to the tenants.

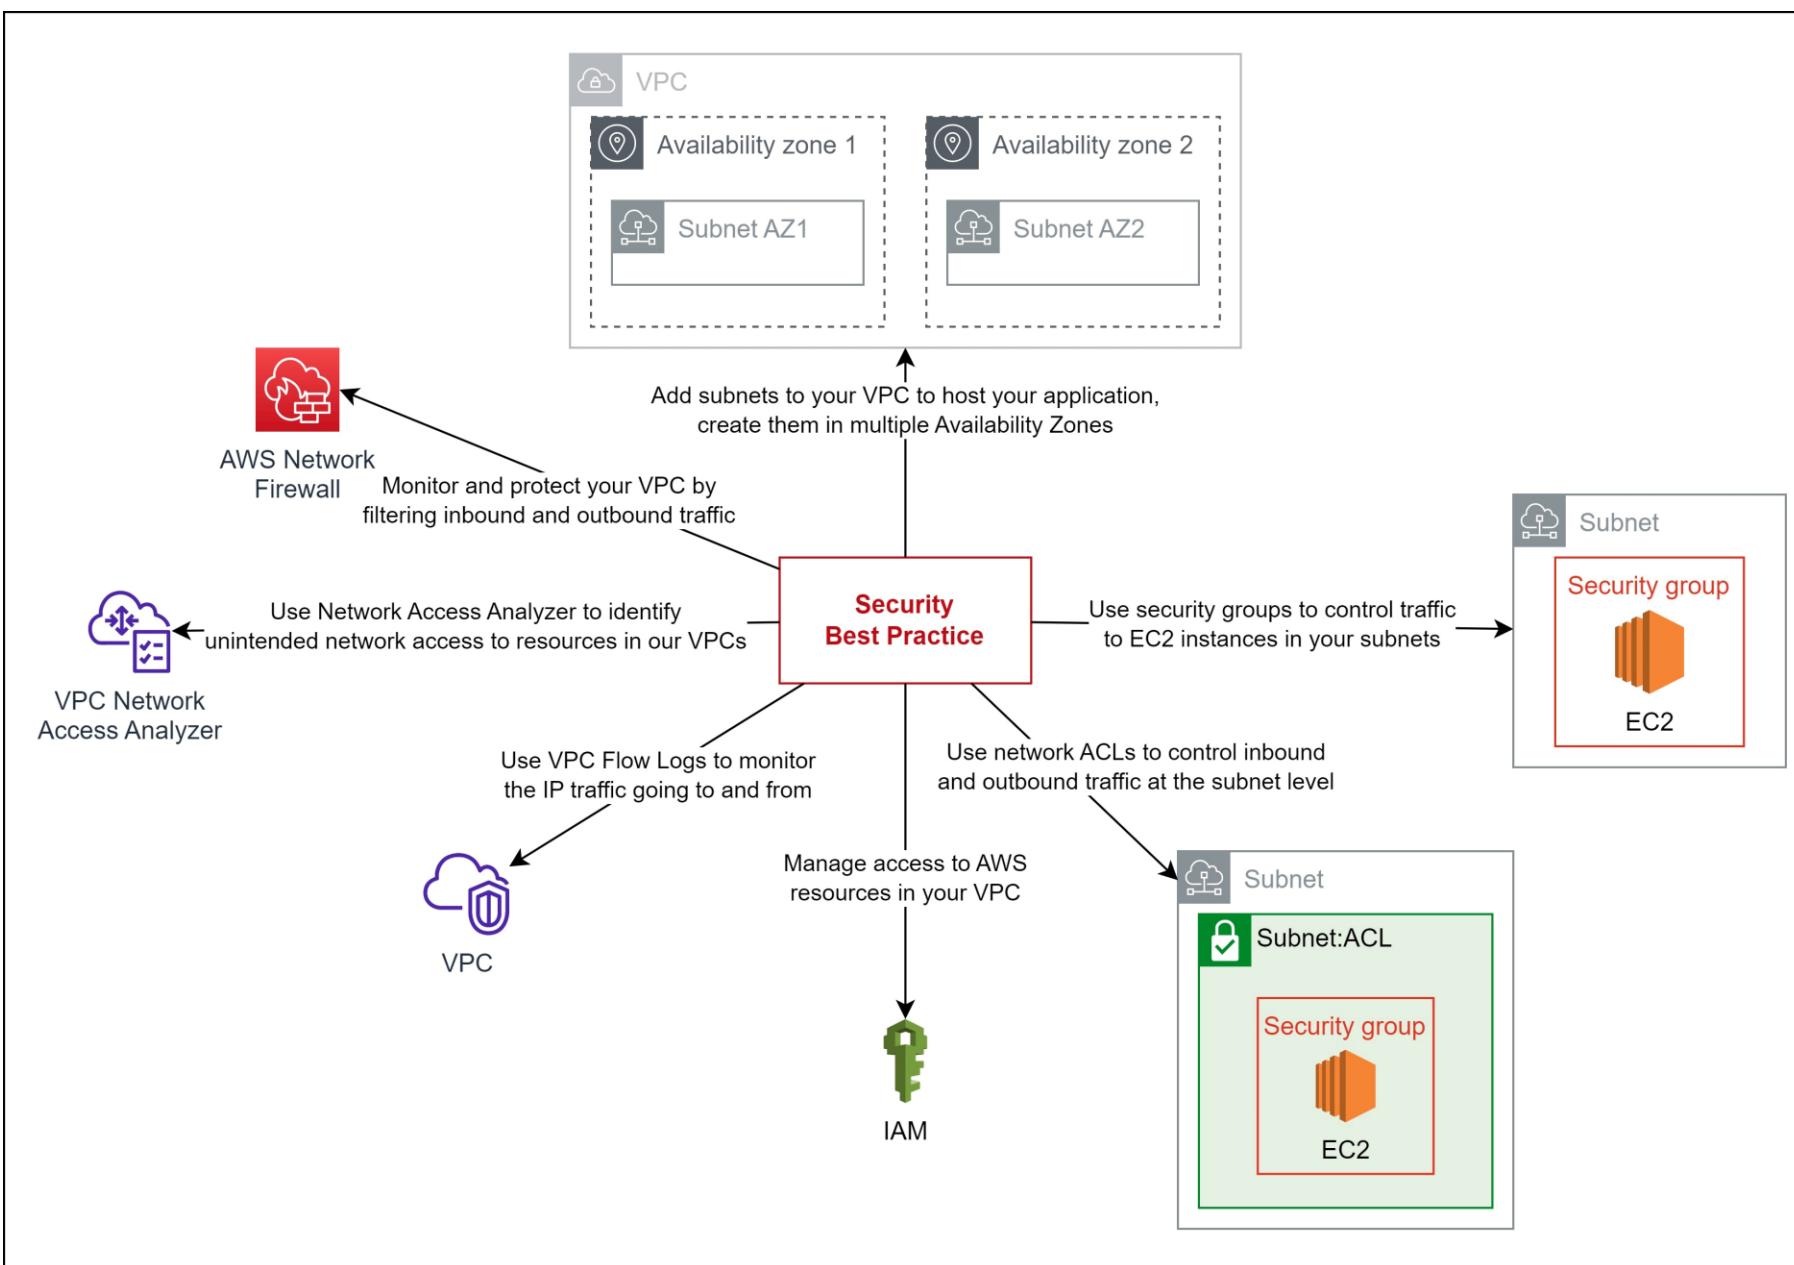
Amazon SNS (Simple Notification Service)

Nash
Tech.



Amazon VPC (Virtual Private Cloud)

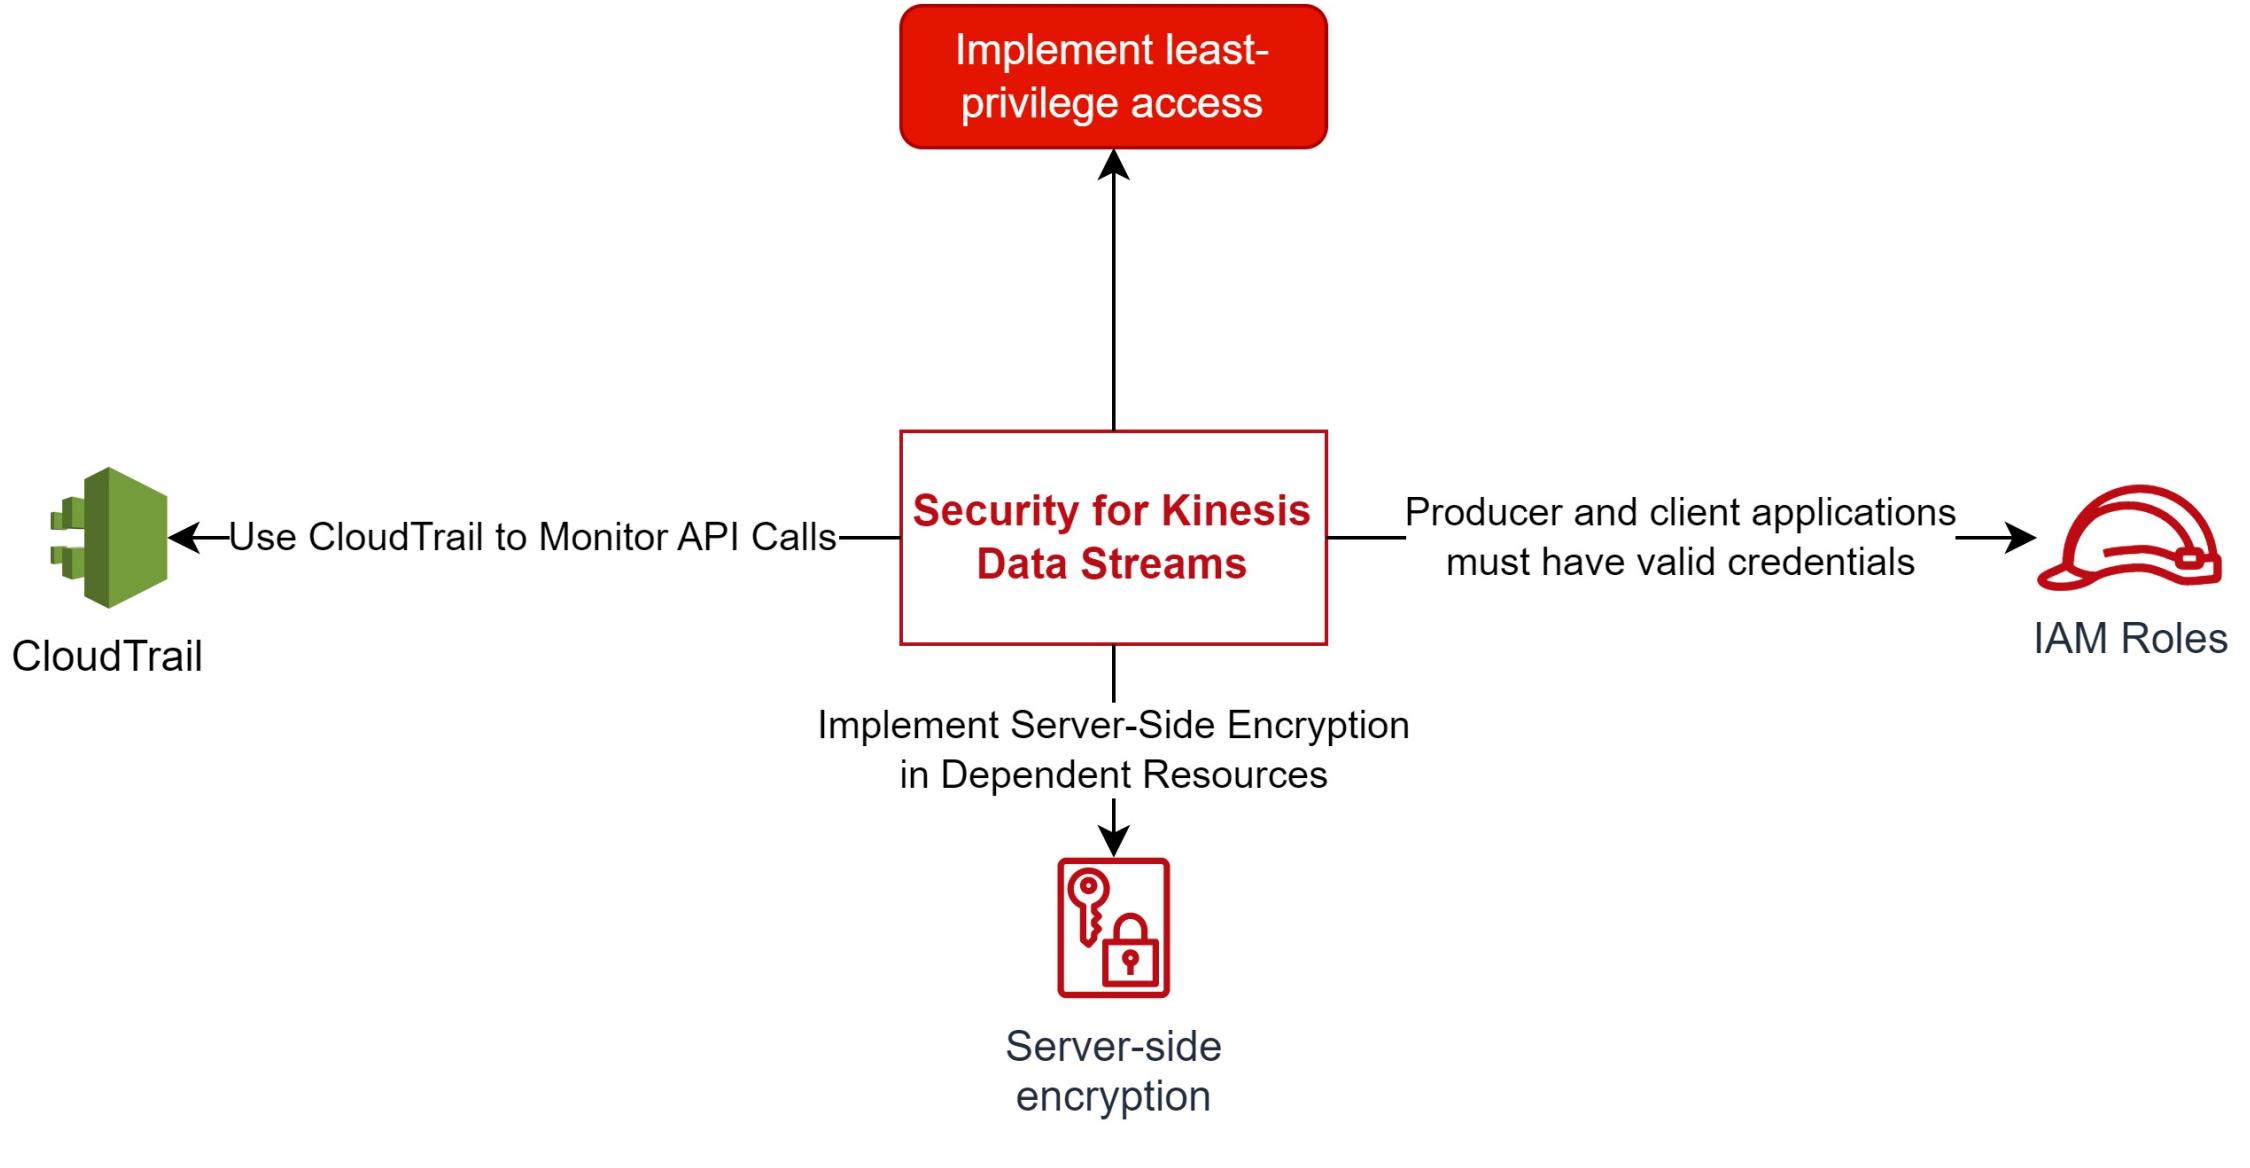
Nash
Tech.

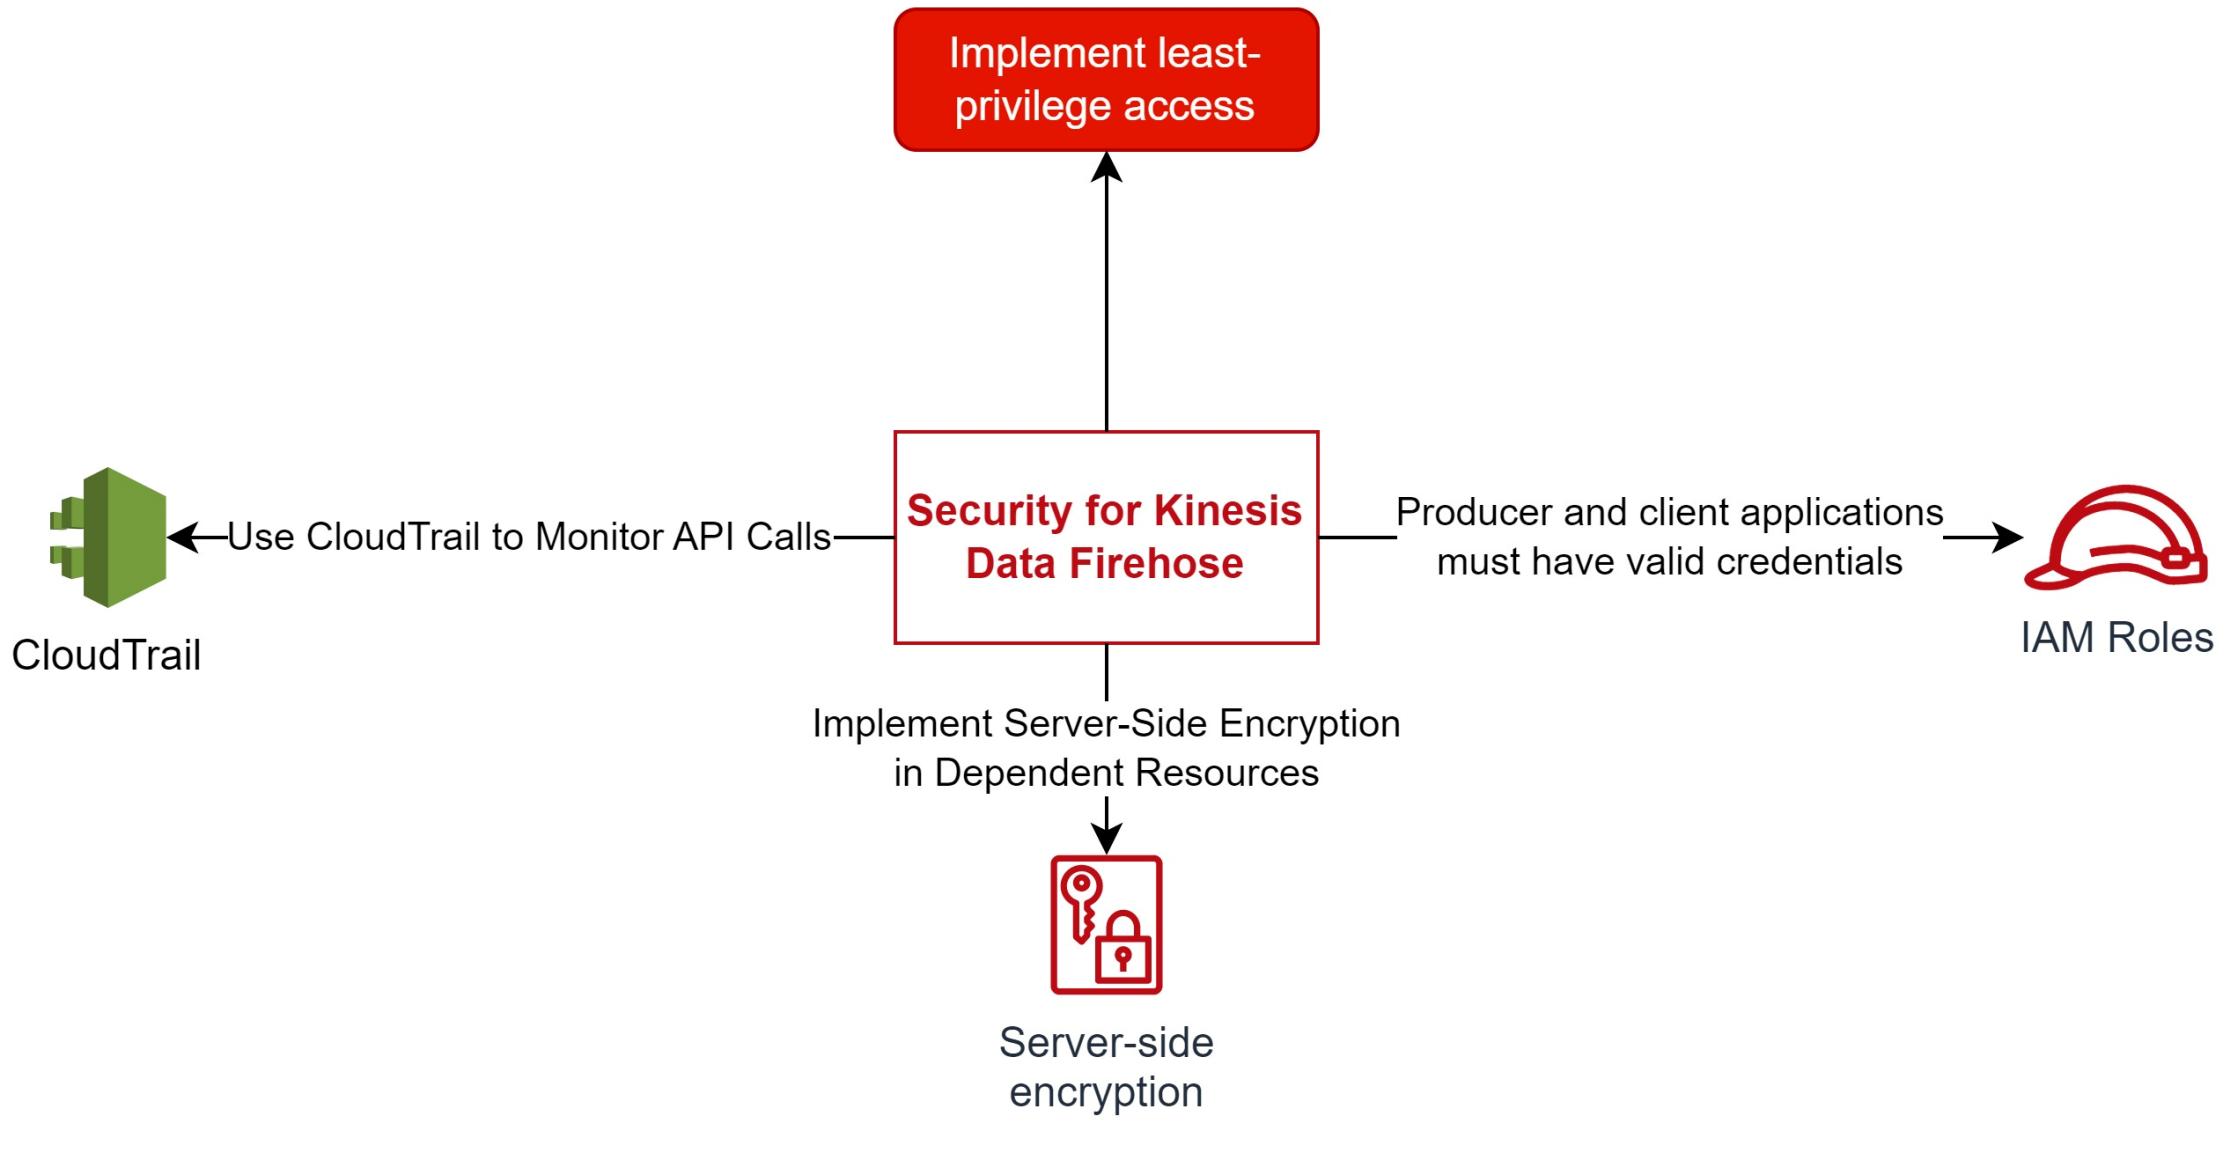


Security best practices for your VPC - Amazon Virtual Private Cloud

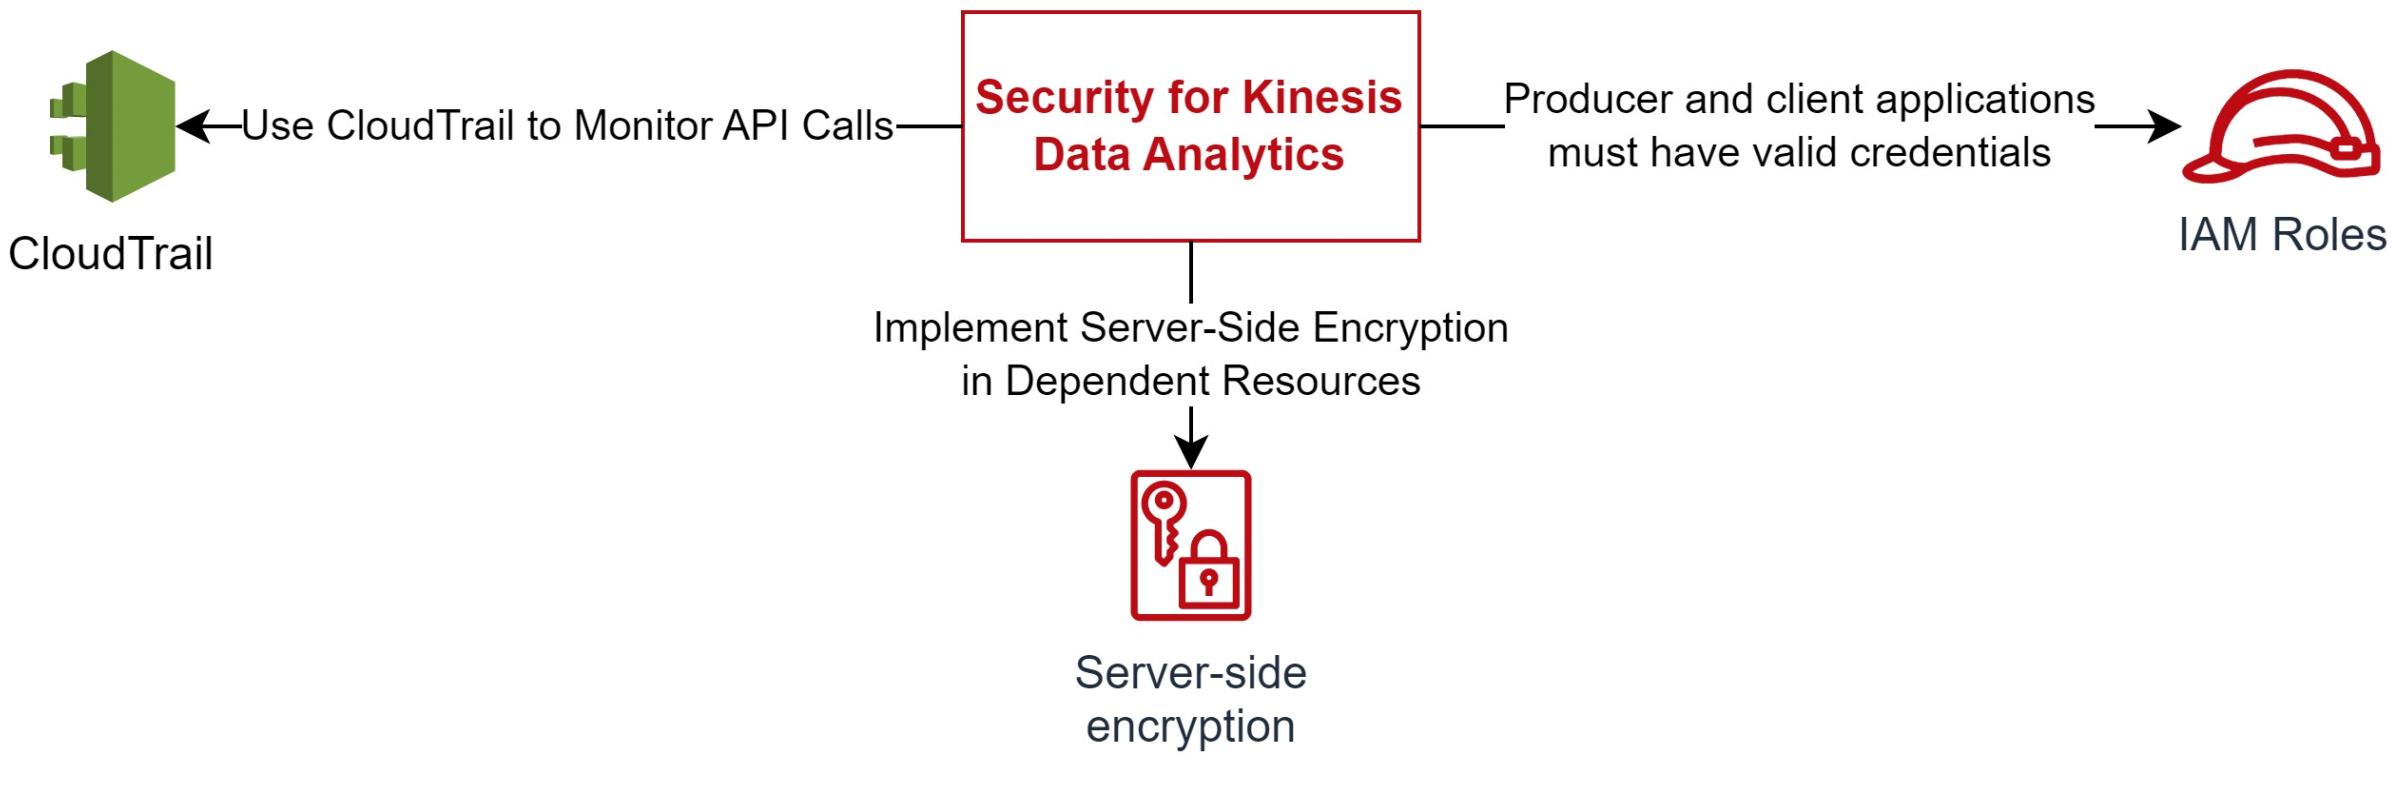
Amazon Kinesis

Nash
Tech.





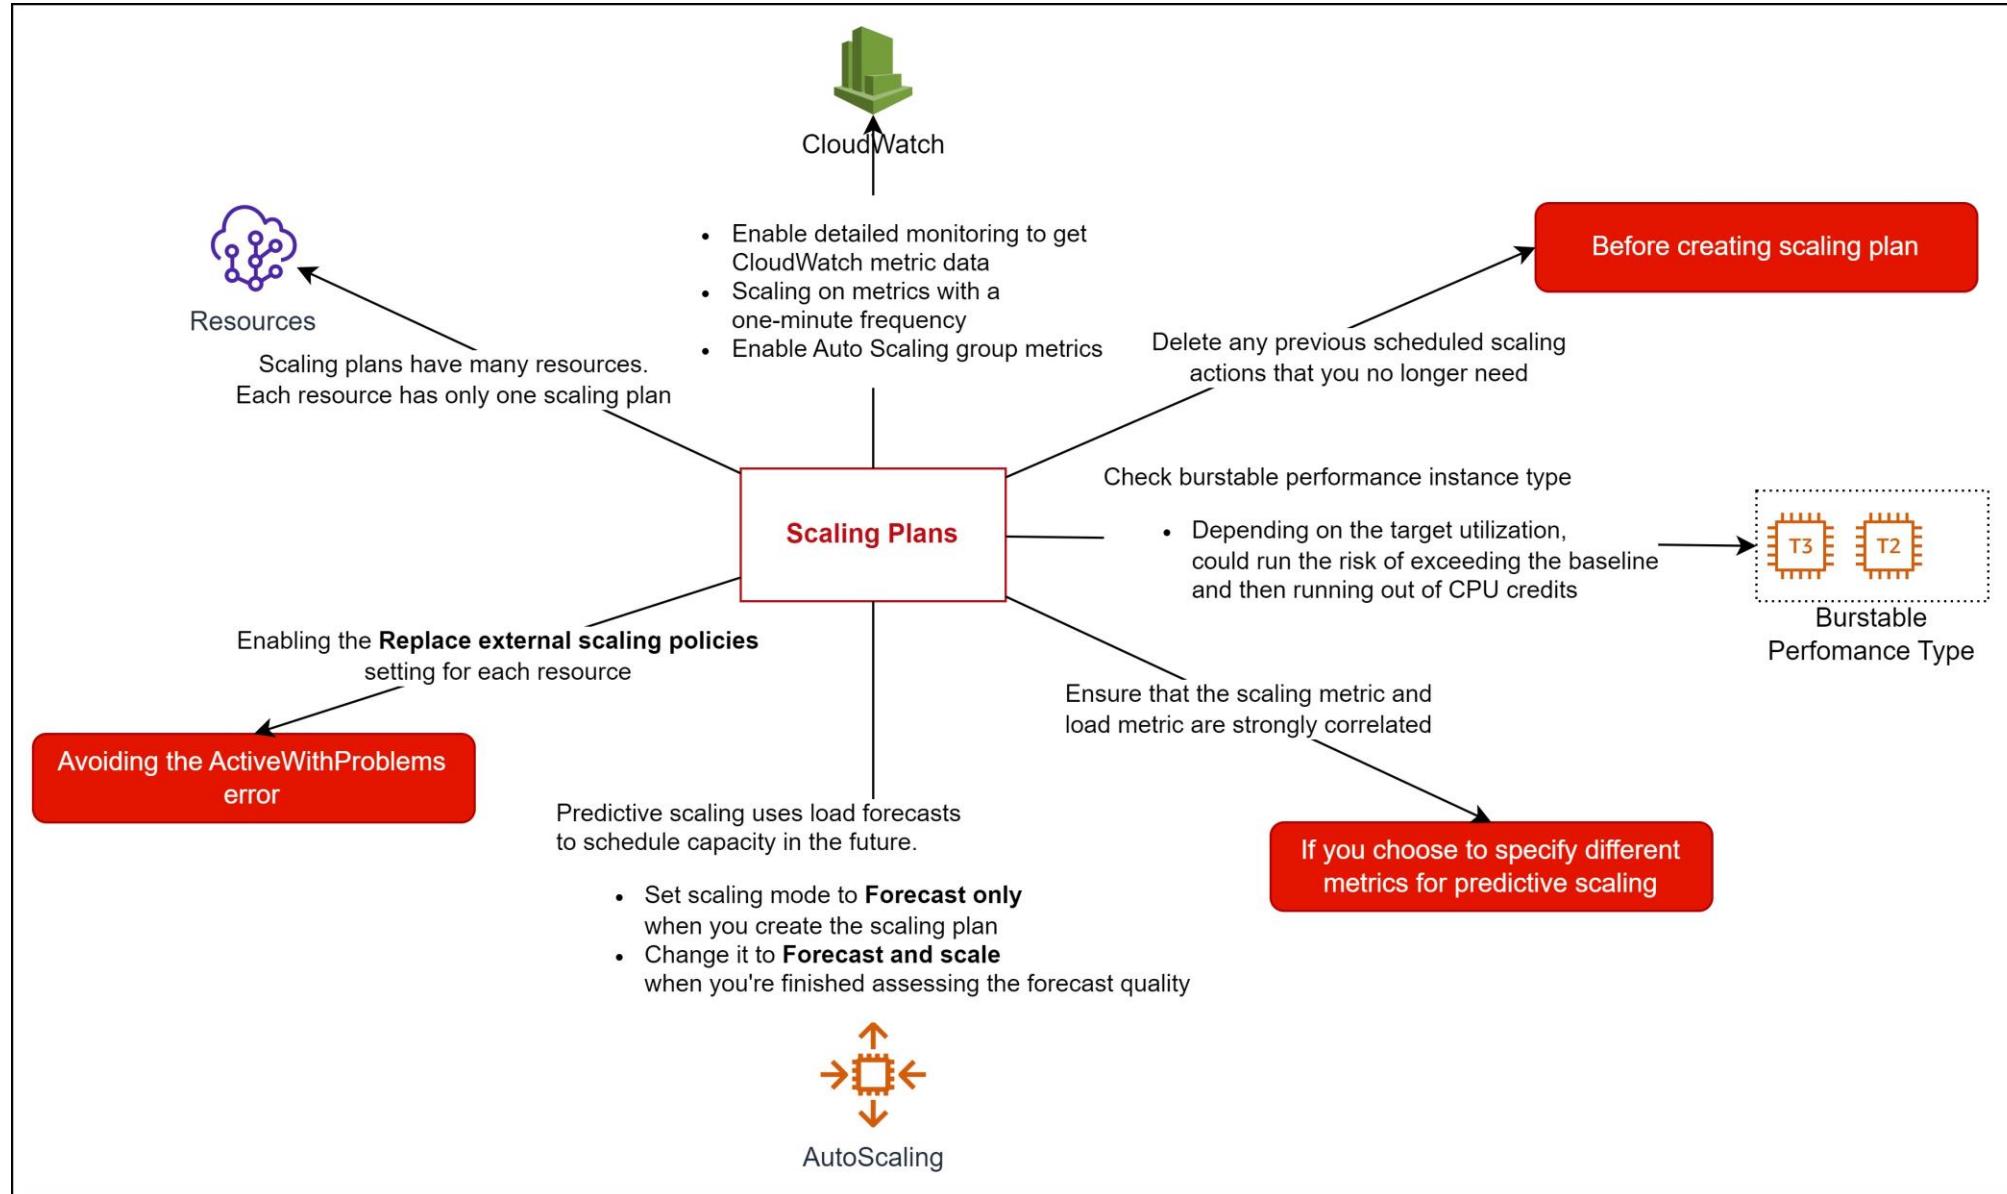
[Security Best Practices for Kinesis Data Firehose - Amazon Kinesis Data Firehose](#)



[Security Best Practices for Kinesis Data Analytics - Amazon Kinesis Data Analytics for SQL Applications Developer Guide](#)

Amazon Auto-scaling

Nash
Tech.



<https://docs.aws.amazon.com/autoscaling/plans/userguide/best-practices-for-scaling-plans.html>

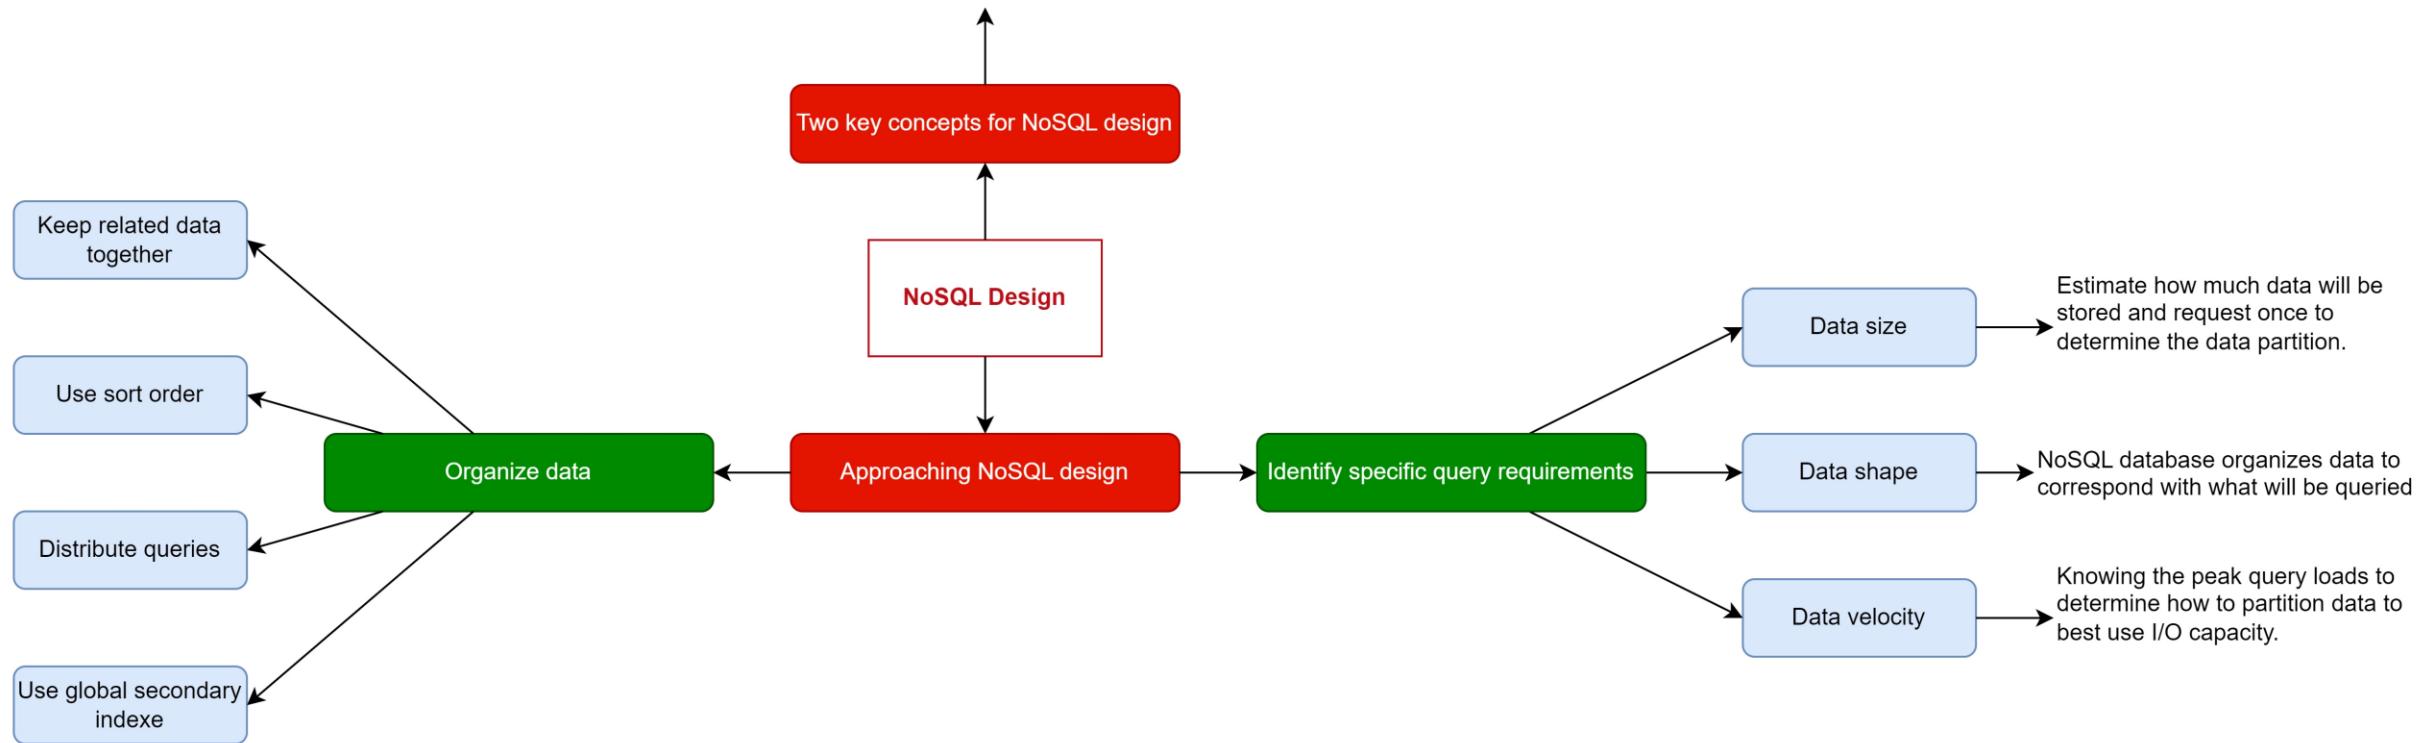
Dynamo DB

Nash
Tech.

DynamoDB best practices

- [NoSQL design for DynamoDB](#)
- [Using deletion protection to protect your table](#)
- [Using the DynamoDB Well-Architected Lens to optimize your DynamoDB workload](#)
- [Best practices for designing and using partition keys effectively](#)
- [Best practices for using sort keys to organize data](#)
- [Best practices for using secondary indexes in DynamoDB](#)
- [Best practices for storing large items and attributes](#)
- [Best practices for handling time series data in DynamoDB](#)
- [Best practices for managing many-to-many relationships](#)
- [Best practices for implementing a hybrid database system](#)
- [Best practices for modeling relational data in DynamoDB](#)
- [Best practices for querying and scanning data](#)
- [Best practices for DynamoDB table design](#)
- [Best practices for DynamoDB global table design](#)
- [Best practices for managing the control plane in DynamoDB](#)

- Understanding the business problems and the application use case before design schema
- You should maintain as few tables as possible in a DynamoDB application

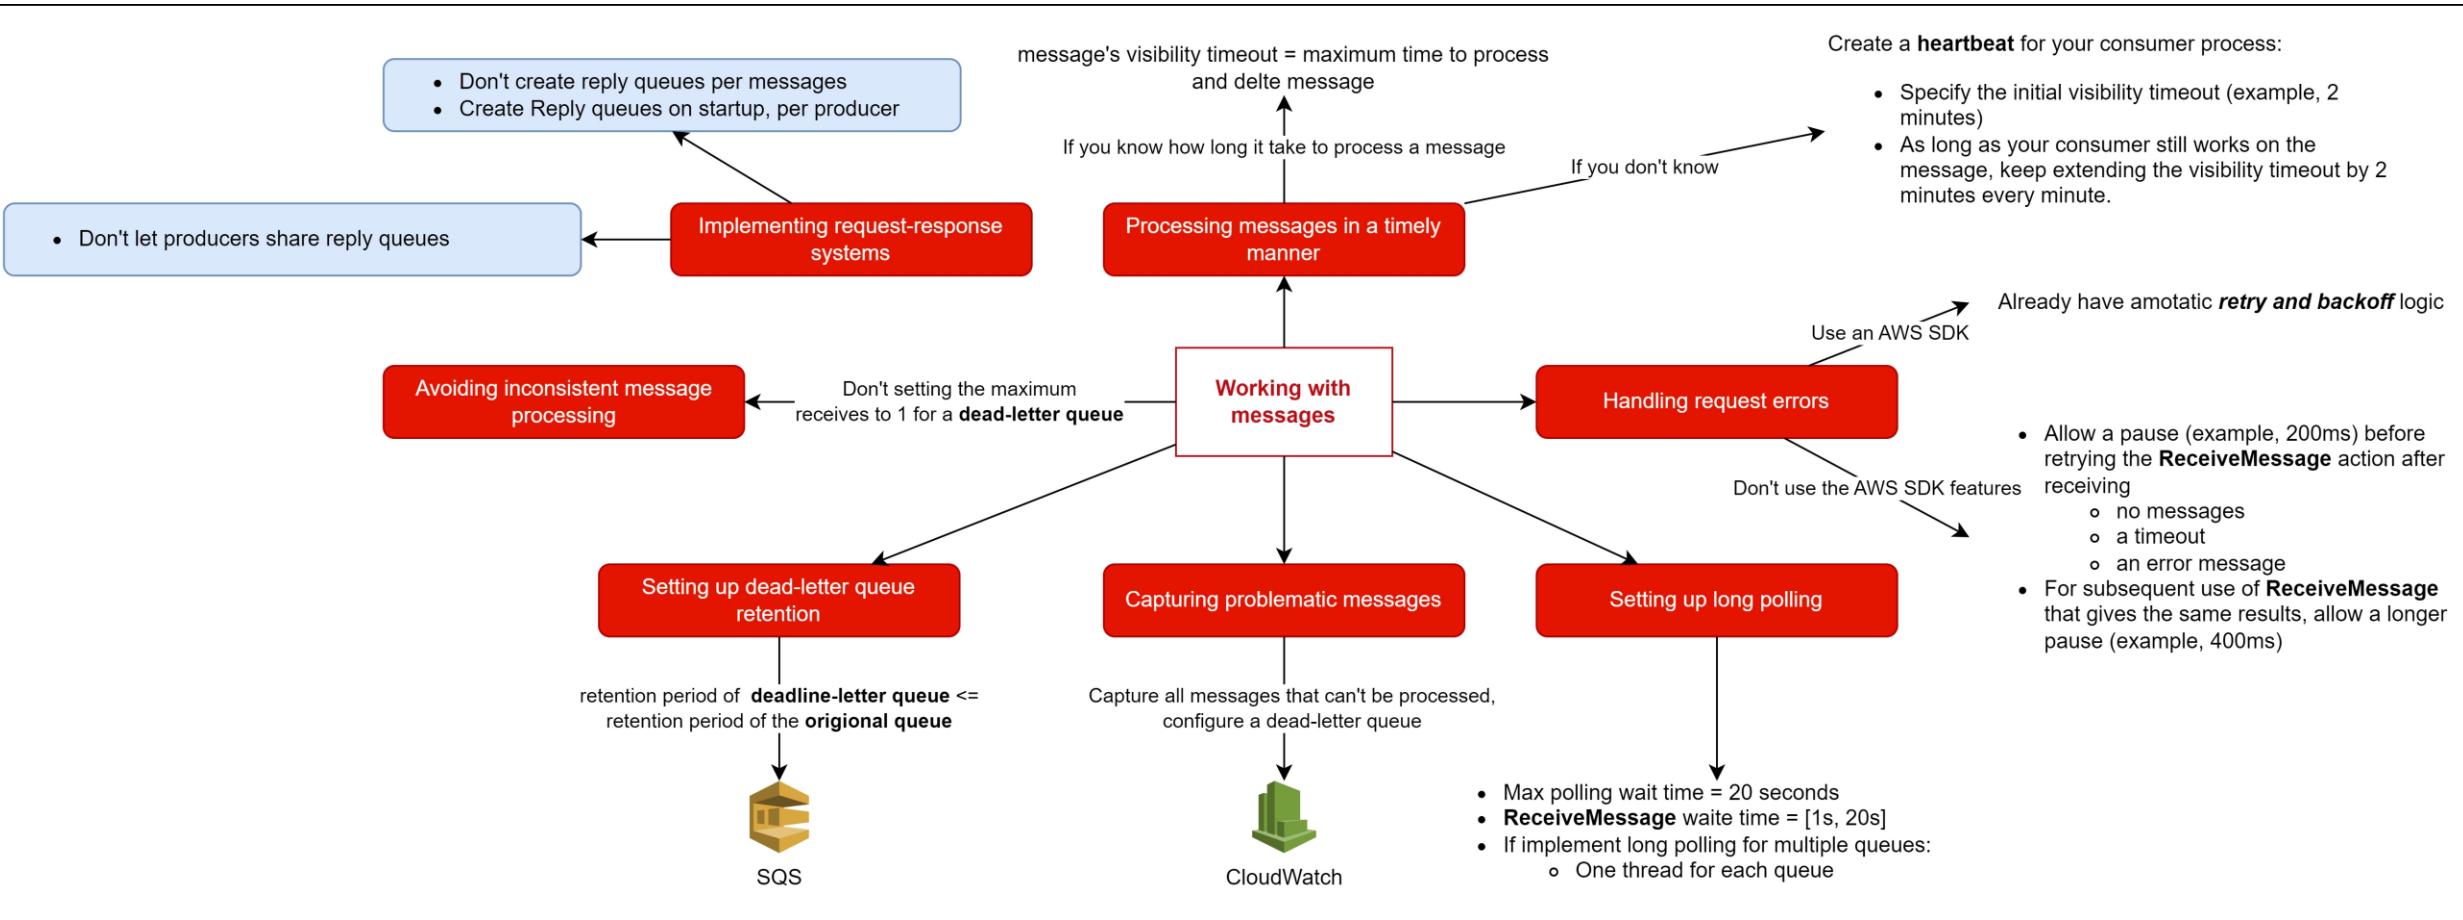


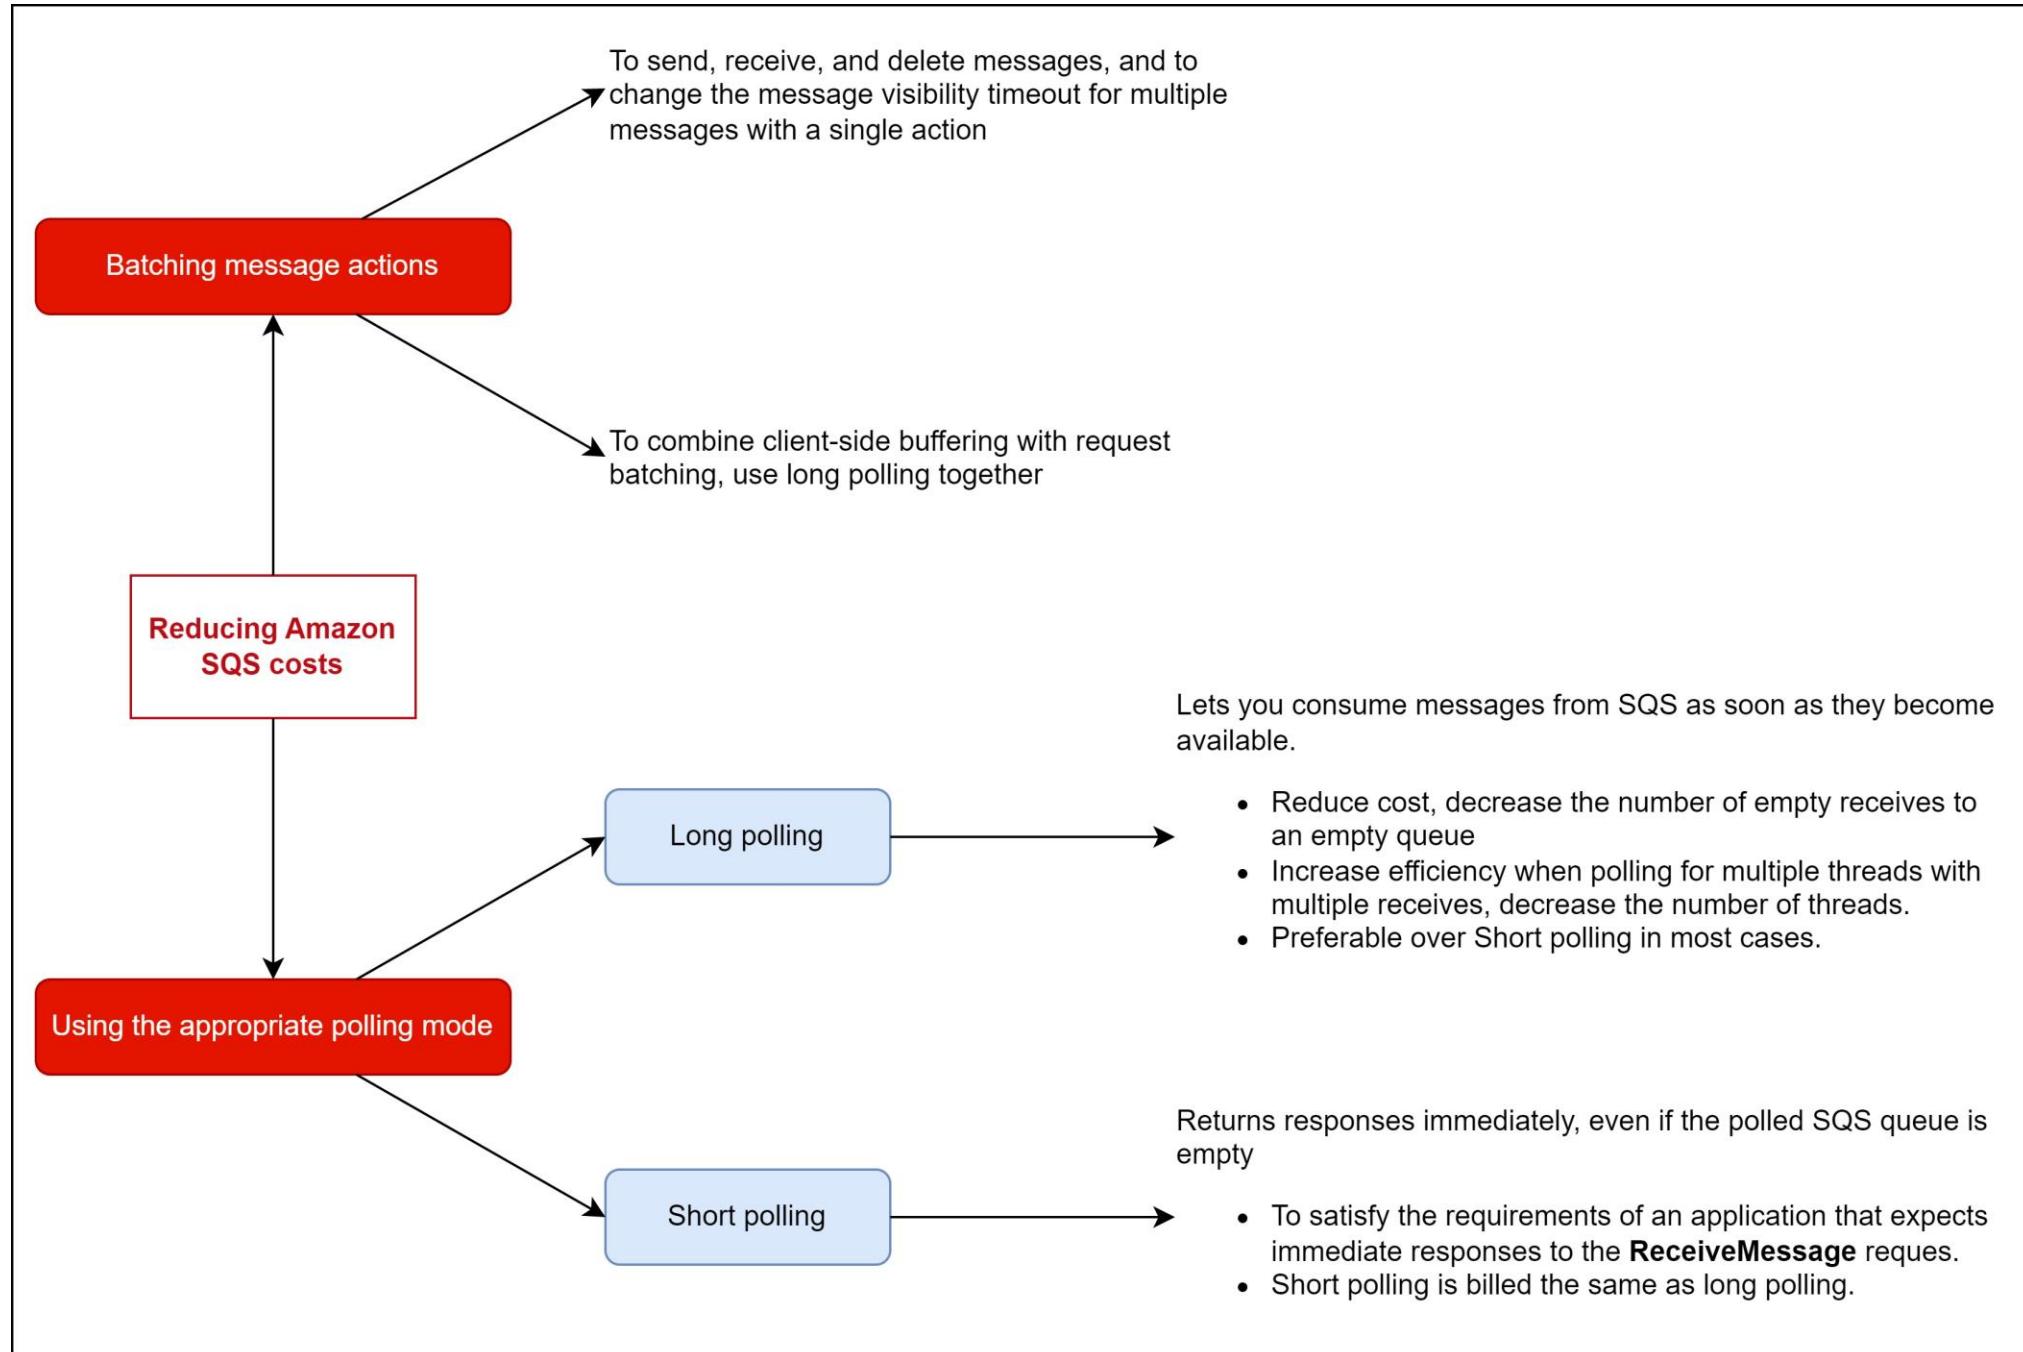
Amazon SQS (Simple Queue Service)

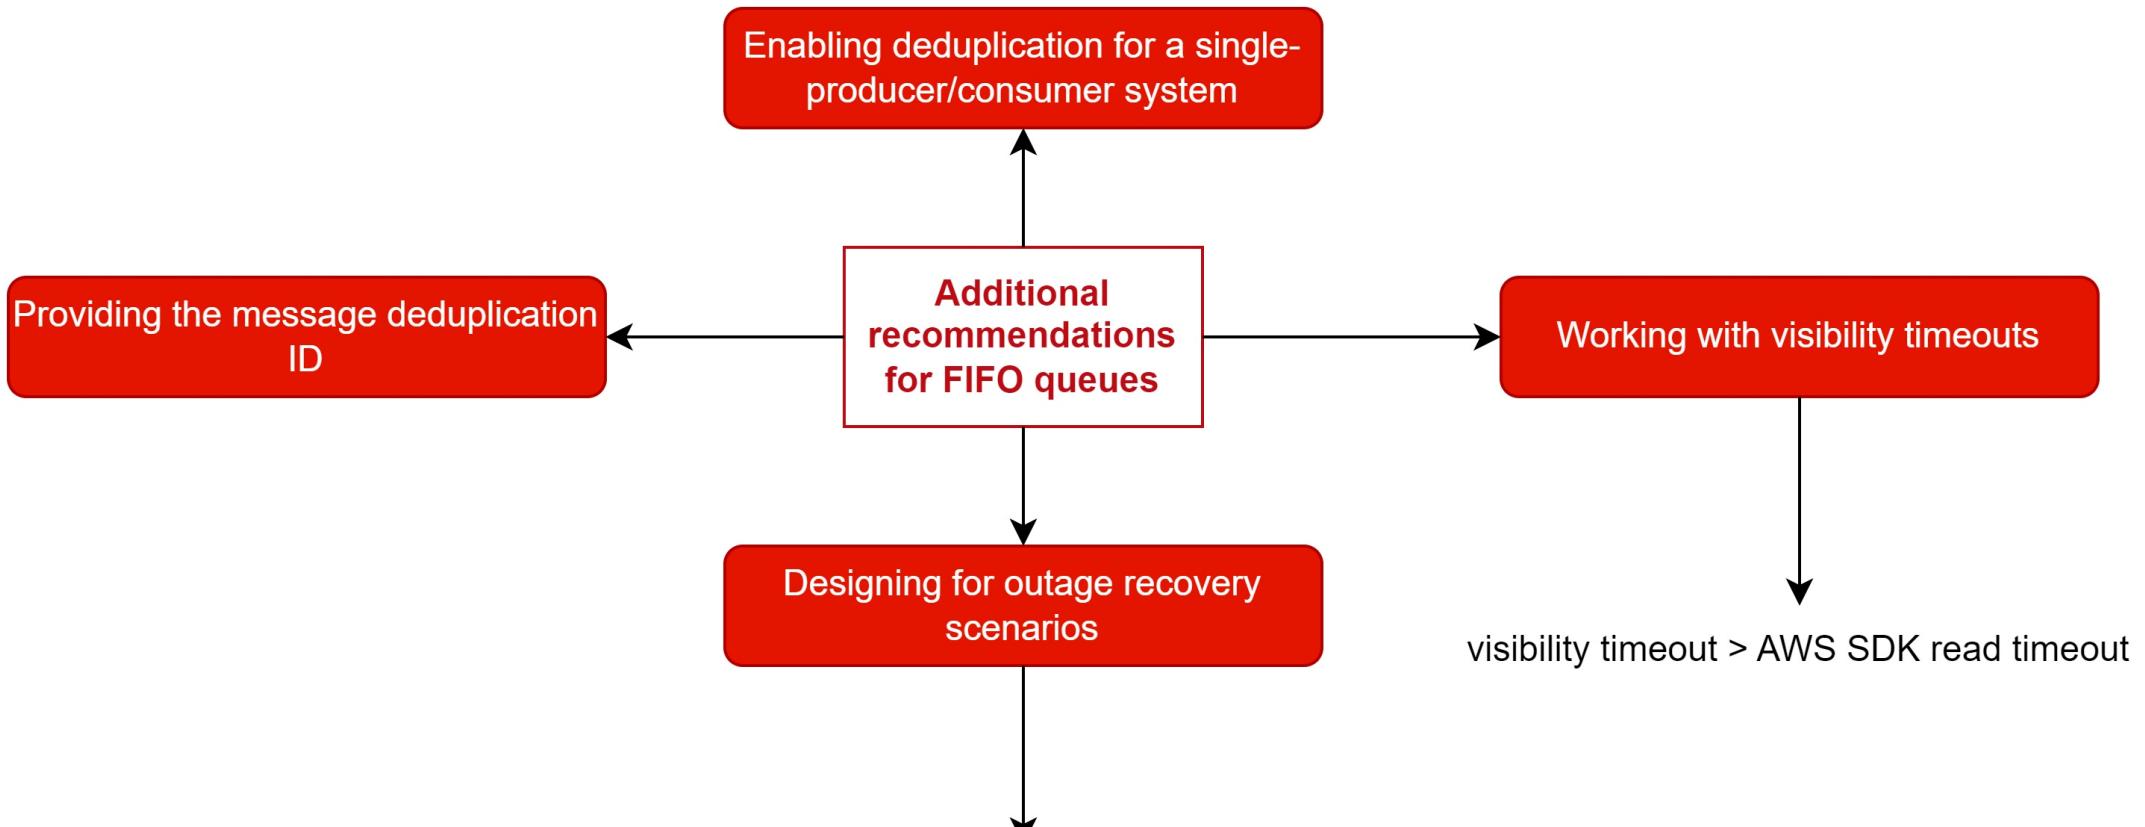
Nash
Tech.

Best practices for Amazon SQS

- [Recommendations for standard and FIFO queues](#)
- [Additional recommendations for FIFO queues](#)







The deduplication process in FIFO queues is time-sensitive.

- The producer must aware of the deduplication interval of the queue.
 - Deduplication interval = 5 mins
 - **Don't Retry** SendMessage requests after the deduplication interval expires
- The consumer must have a visibility timeout that minimizes the risk of being unable to process messages before the visibility timeout expires.

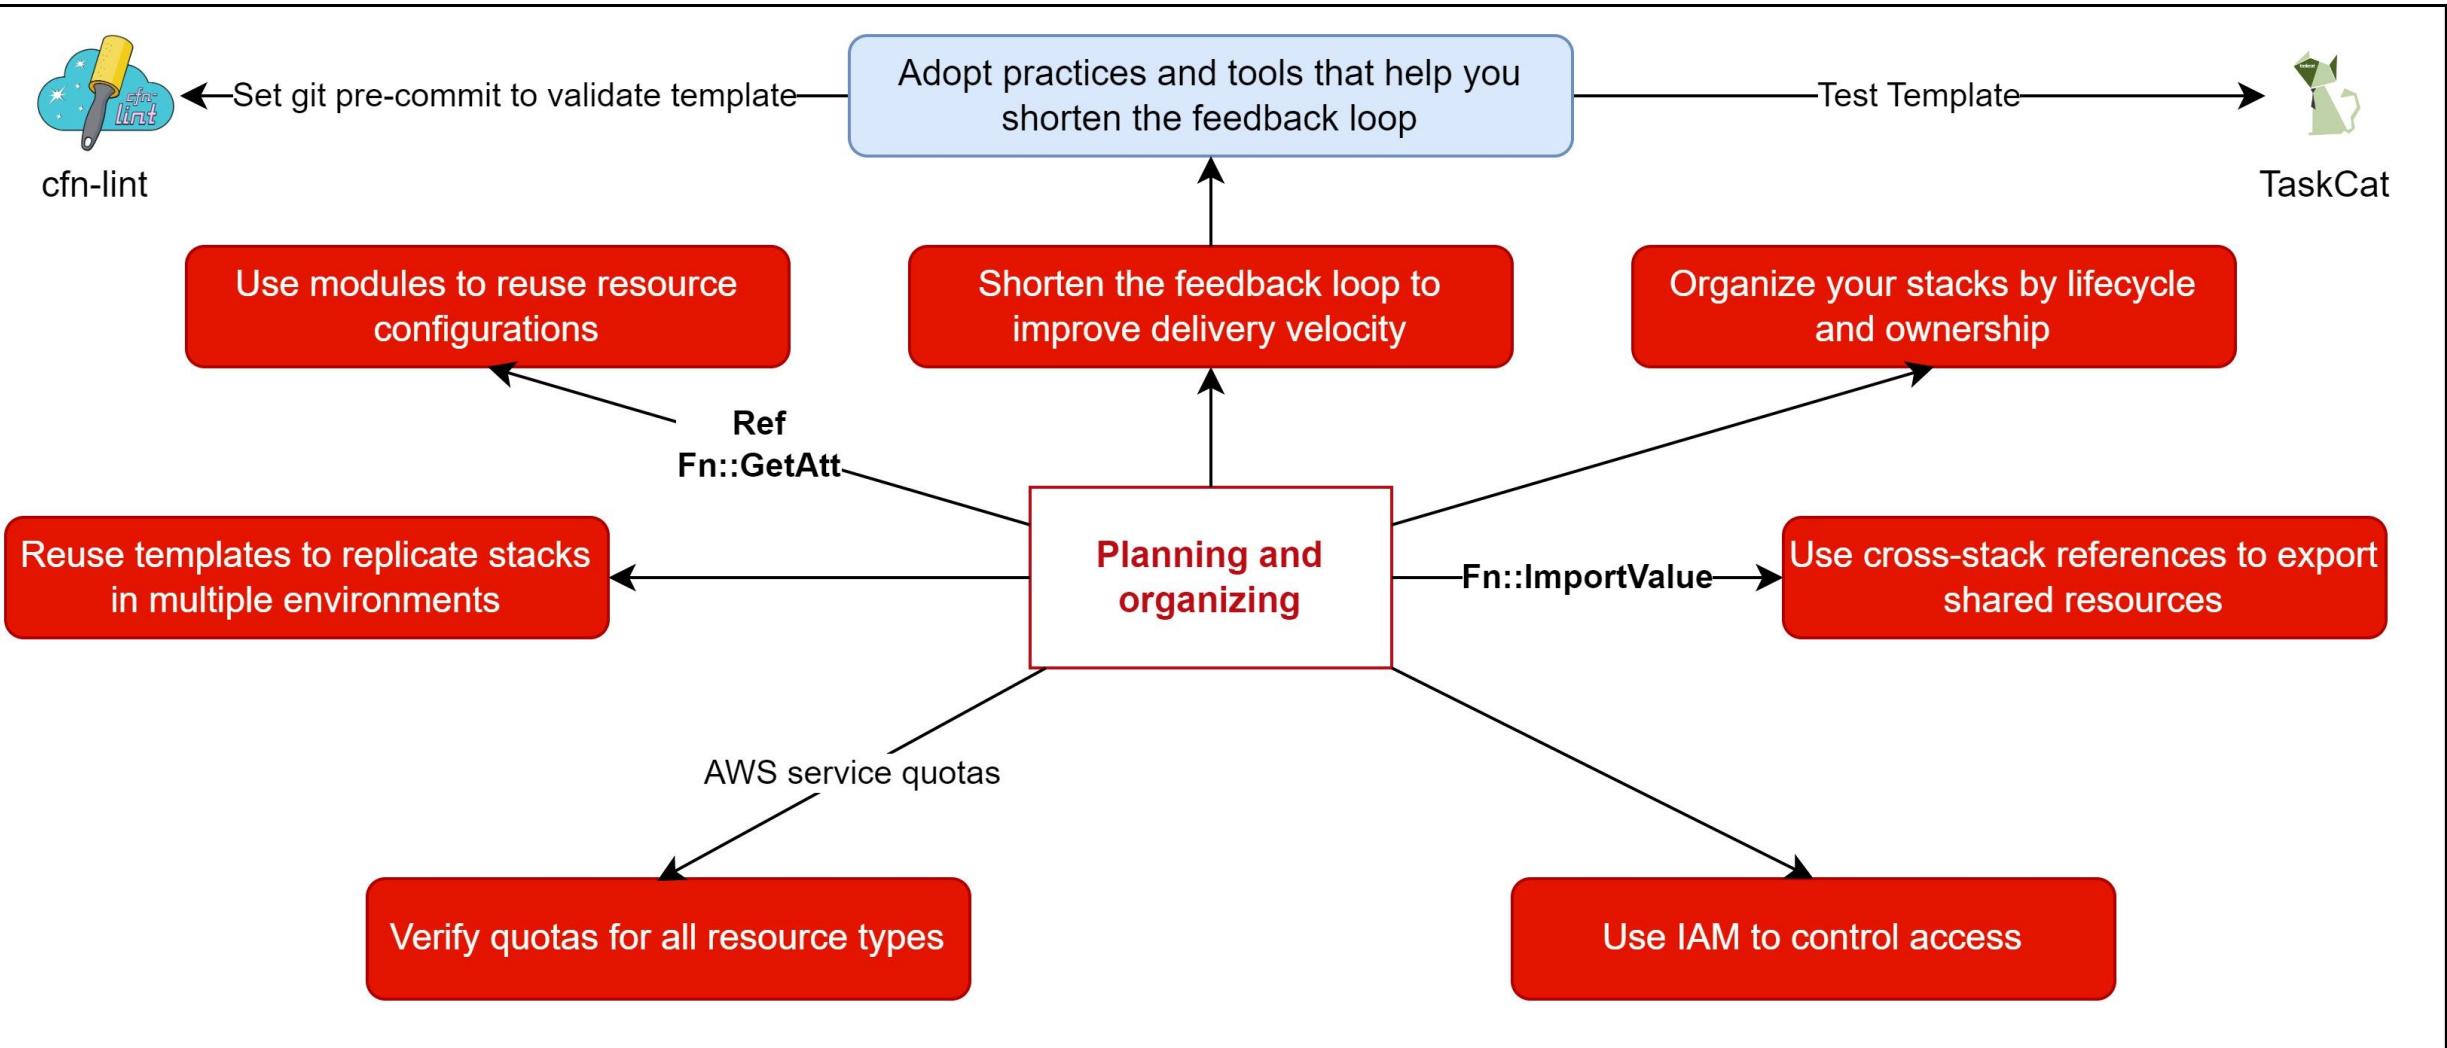
AWS CloudFormation

Nash
Tech.

Best practices

Planning and organizing

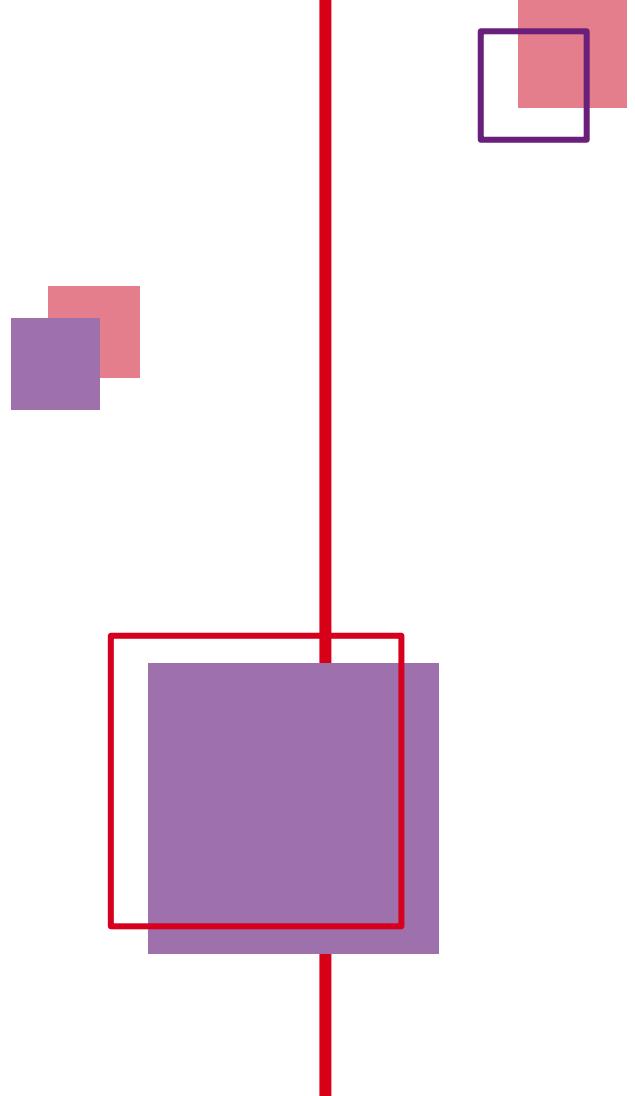
- Shorten the feedback loop to improve delivery velocity
- Organize your stacks by lifecycle and ownership
- Use cross-stack references to export shared resources
- Use IAM to control access
- Reuse templates to replicate stacks in multiple environments
- Verify quotas for all resource types
- Use modules to reuse resource configurations

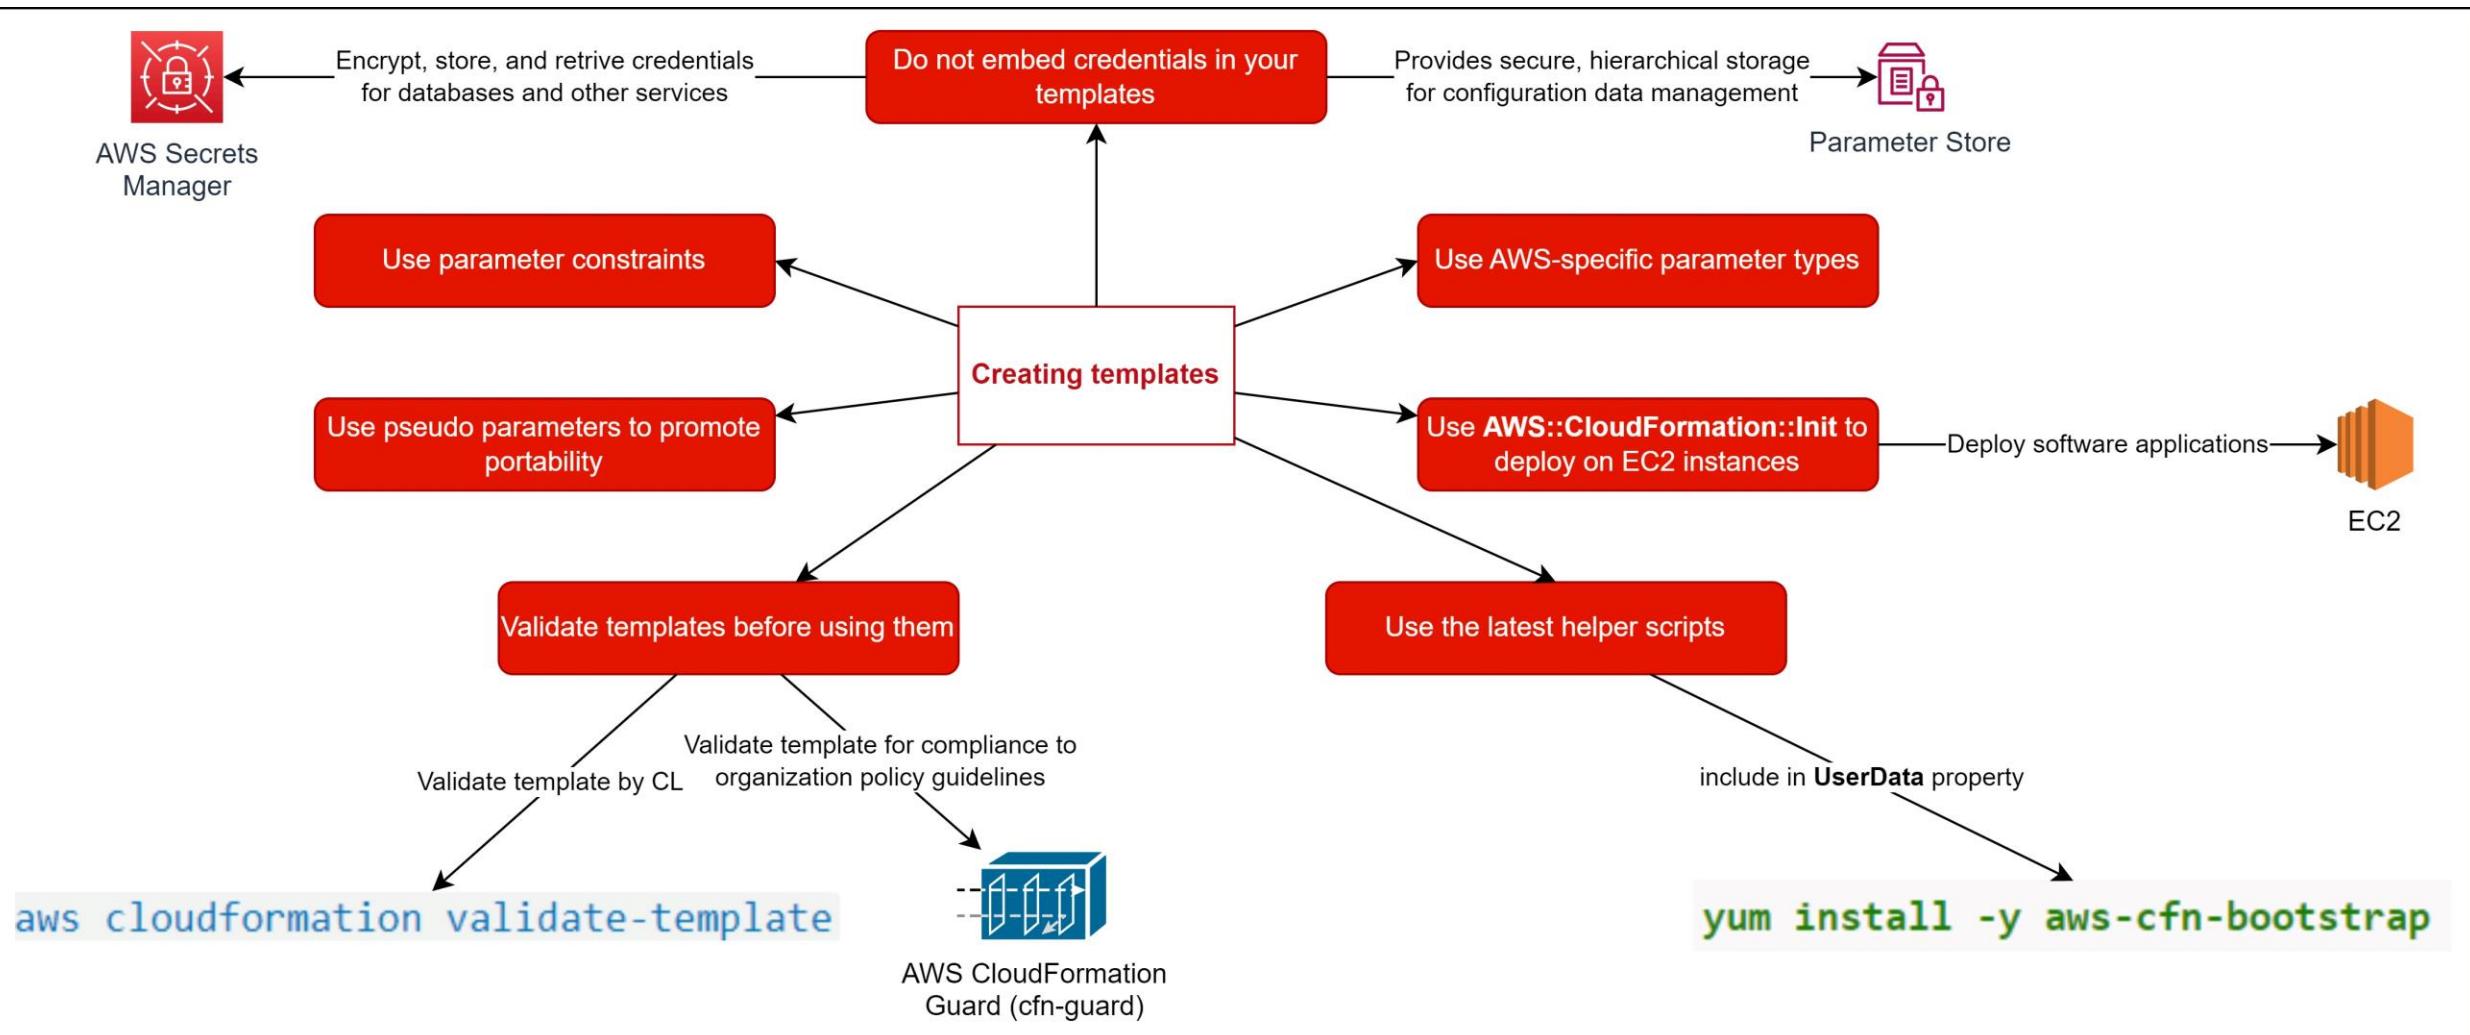


Best practices (2)

Creating templates

- [Do not embed credentials in your templates](#)
- [Use AWS-specific parameter types](#)
- [Use parameter constraints](#)
- [Use pseudo parameters to promote portability](#)
- [Use AWS::CloudFormation::Init to deploy software applications on Amazon EC2 instances](#)
- [Use the latest helper scripts](#)
- [Validate templates before using them](#)

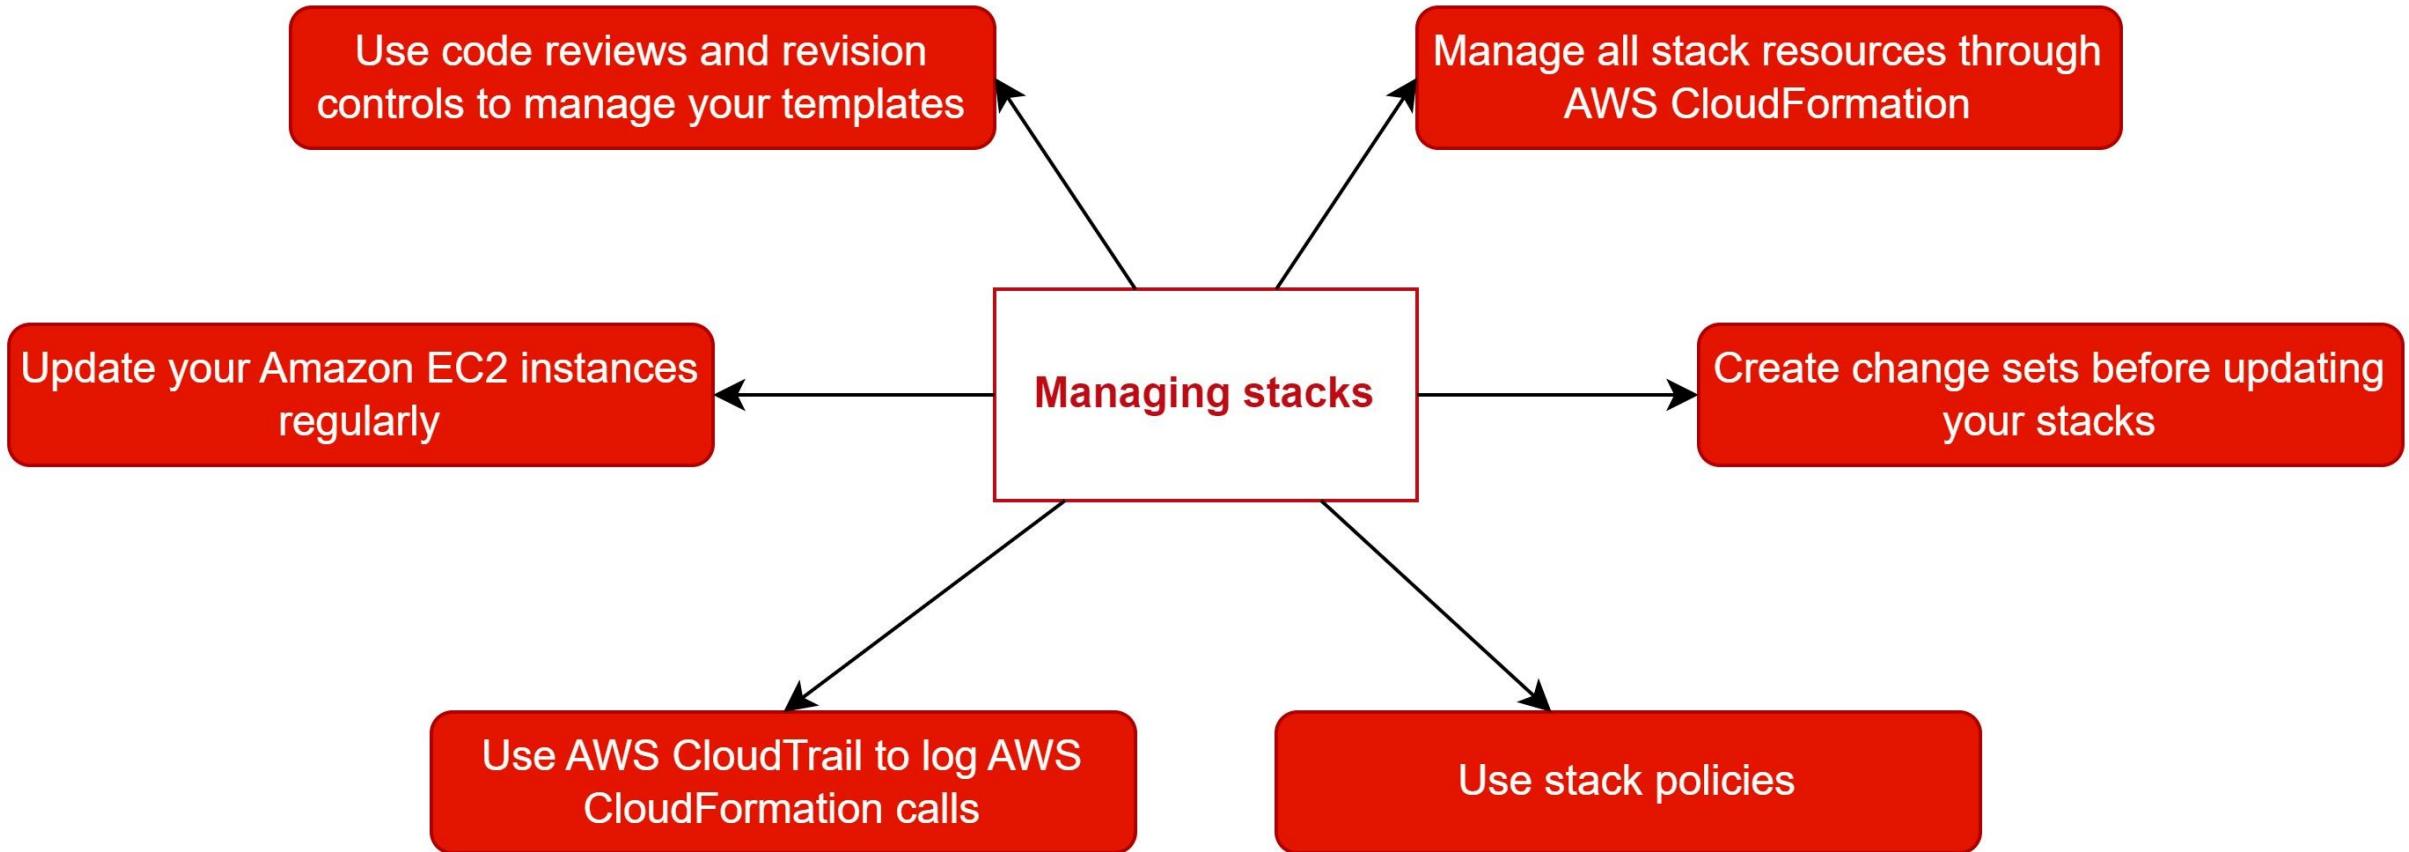




Best practices (3)

Managing stacks

- [Manage all stack resources through AWS CloudFormation](#)
- [Create change sets before updating your stacks](#)
- [Use stack policies](#)
- [Use AWS CloudTrail to log AWS CloudFormation calls](#)
- [Use code reviews and revision controls to manage your templates](#)
- [Update your Amazon EC2 instances regularly](#)



Thank you